



# RTM-Aware Quantum Computing

A Multiscale, Slope-First Framework for Coherence, Scheduling, and Design

Álvaro Quiceno

# Abstract

We introduce a **slope-first** methodology for quantum computing based on **Multiscale Temporal Relativity (RTM)**. Inside a fixed operational regime, RTM posits that a characteristic time  $T$  scales with a size/scale proxy  $L$  by a power law,

$$\log T = \alpha \log L + c,$$

where the **coherence exponent**  $\alpha$  is the **clock-invariant** structural signal and  $c$  encodes clock/units. We adapt RTM to quantum stacks—**physical**, **QEC**, **compiler/runtime**, and **I/O-cryo**—by defining layer-specific  $(L, T)$  pairs (e.g., number of active qubits vs. stable calibration time; code distance vs. logical-failure time; multiplexing degree vs. readout latency; circuit width vs. makespan), and estimating binwise slopes under errors-in-variables (ODR/TLS, Theil-Sen, SIMEX). A **collapse test** validates scaling and guards against regime mixing; clean family-wise slopes are fused into a real-time  $ECI_{QC}(t)$  with uncertainty and QA gates.

We formulate **falsifiable** hypotheses: **(H1)** higher pre-shock  $\alpha$  predicts longer stability margins (fewer forced recalibrations, lower logical error at fixed  $d$ ); **(H2) decoherence events**—significant QA-clean drops in  $ECI_{QC}$ —lead spikes in logical error, queueing, or makespan; **(H3)** micro→meso→macro **tempo cascades** exhibit non-decreasing  $\alpha$  within stable regimes. We demonstrate how **RTM-aware scheduling** (batching, staggered resets, low-variance routing), **QEC cadence design** (desynchronization of syndrome cycles), and **modular sizing** (sweet spots for interconnect) can improve throughput and reliability without changing physical fidelities. The framework is reproducible, gauge-robust (unit/clock changes do not affect  $\alpha$ ), and designed to fail gracefully (no-collapse and high heterogeneity become scope boundaries, not post-hoc fixes).

## 1. Introduction

### 1.1 Motivation: beyond fidelities and error rates

Quantum performance is usually summarized by **point metrics**—single- and two-qubit fidelities,  $T_1/T_2$ , logical error rates, or benchmark figures (QED-C, QV). Yet practical reliability and throughput hinge on something orthogonal: **how timing stretches across scale** in a multistage stack—qubits and resonators, code cycles, compilers, cryogenic I/O. When small subsystems respond quickly and larger ones respond more slowly in a disciplined, layered fashion, shocks are **dissipated**; when timings **flatten**, disturbances percolate across layers and synchronize failures (stalling readout, spiking logical error, or forcing global recalibrations).

**Multiscale Temporal Relativity (RTM)** provides a compact language for this phenomenon. Inside a fixed regime, RTM expects a power-law relation between a **characteristic time**  $T$  and a **scale proxy**  $L$ : the **slope**  $\alpha$  in  $\log T = \alpha \log L + c$  is structural (invariant to time units), while the intercept  $c$  is a **clock** (gauge). We bring this principle to quantum computing and show that measuring, validating, and **engineering**  $\alpha$  yields actionable levers—**independent of nominal units**—to improve stability and throughput.

## 1.2 RTM in one line

**Structure lives in the slope; clocks live in the gauge.**

A change of clock or units shifts  $c$  but leaves  $\alpha$  unchanged. Thus  $\alpha$  can be compared across devices, stacks, and labs, while  $c$  cannot.

## 1.3 Contributions

This paper makes five contributions:

1. **Operationalization of RTM for QC.** We define layer-specific ( $L, T$ ) pairs for **physical**, **QEC**, **compiler/runtime**, and **I/O-cryo** layers (e.g.,  $L$  = active qubits,  $T$  = stable calibration time;  $L = d$ ,  $T$  = cycles to logical failure;  $L$  = multiplexing degree,  $T$  = readout latency;  $L$  = circuit width,  $T$  = makespan).
2. **Validation & estimation.** We provide a **collapse test** (residual independence of  $\log T - \alpha \log L$  from  $\log L$ ) to detect regime mixing and non-power curvature, and adopt **errors-in-variables** estimation (ODR/TLS, Theil-Sen, SIMEX) with bootstrap uncertainty and changepoint guards.
3. **A single real-time indicator.** We fuse family-wise slopes into  $ECI_{QC}(t)$  via random-effects meta-analysis with heterogeneity controls ( $Q, I^2, \hat{\tau}^2$ ); we publish QA flags and withhold fusion when proxies disagree.
4. **Design levers.** We formalize **RTM-aware scheduling** (batching, staggered resets, low-variance routing), **QEC cadence design** (desynchronization to avoid phase lock between physical errors and syndrome extraction), and **modular sizing** (choosing module/interconnect scales that elevate  $\alpha$  without throttling throughput).
5. **Falsifiable hypotheses & protocols.** We pre-register **H1-H3** with A/B protocols on superconducting and trapped-ion platforms, metrics (throughput, makespan, logical error, uptime, p95/p50 ratios), and decision thresholds for adoption.

## 1.4 What $\alpha$ is—and is not

- **Is:** a **binwise slope** linking a time  $T$  to a scale  $L$  inside a **fixed environment** (same temperature/firmware/topology/syndrome schedule). It captures the **geometry of tempo across scale**.

- **Is not:** a causal parameter by default; level changes in  $T$  (units, clocks, offsets) do **not** change  $\alpha$ . When collapse fails,  $\alpha$  is **undefined** for that bin and should not be fused.

## 1.5 Layer-specific ( $L, T$ )exemplars (preview)

- **Physical:**  $L$  = active qubits / coupler degree / cluster size;  $T$  = stable calibration interval, gate/RO latency, mean time to drift.
- **QEC:**  $L = d$  (code distance) or number of logical qubits;  $T$  = cycles to logical failure; cadence of syndrome extraction.
- **Compiler/runtime:**  $L$  = circuit width or depth after mapping;  $T$  = makespan; queueing delay and rescheduling latency.
- **I/O-cryo:**  $L$  =multiplexing degree or channels;  $T$  = readout latency/BER recovery; p95 queue length.

Each layer yields a slope  $\hat{\alpha}_f$ ; after QA and collapse, we fuse them into  $ECI_{QC}(t)$  with uncertainty bands. Clean **decoherence events** are significant drops in  $ECI_{QC}$  over pre-registered horizons.

## 1.6 Hypotheses (falsifiable)

- **H1 (Resilience):** Higher pre-shock  $\alpha$  associates with smaller logical-error spikes at fixed  $d$  and longer stable calibration intervals.
- **H2 (Anticipation):** QA-clean  $ECI_{QC}$  drops lead increases in makespan, queueing, or logical error by weeks to months, adding predictive value over baselines (fidelity, utilization, temperature).
- **H3 (Cascade):** Within stable regimes,  $\alpha_{\text{physical}} \leq \alpha_{\text{QEC}} \leq \alpha_{\text{runtime/I/O}}$ ; directionality tests favor micro→meso→macro timing flow.

## 1.7 RTM-aware design (intuitions we will test)

- **Scheduling:** Avoid patterns that **flatten**  $\alpha$  (long, tightly coupled operations in parallel); favor **batching** readouts and **staggered** resets to prevent synchronization cascades.
- **QEC cadence:** Introduce slight **desynchronization** (phase offsets) between syndrome cycles and known noise rhythms to raise  $\alpha_{\text{QEC}}$ .
- **Modularity:** Choose module size and interconnect density where  $\alpha$  is high enough to damp inter-module cascades but not so high that throughput is throttled.

## 1.8 Relation to prior work

Our framework complements fidelity-centric and error-model approaches by adding a **scale–tempo geometry**. It is compatible with (not a replacement for) surface/LDPC code theory, compilation/routing heuristics, and queueing models; it contributes a **gauge-invariant** statistic  $\alpha$  and a **collapse** specification test to separate **structure** from **clock** effects. In the language of stochastic processes, our dynamics section (later) connects RTM to **time-changed diffusions**; in meta-analysis terms, our fusion mimics **random-effects** with explicit **heterogeneity gates**.

## 1.9 Paper roadmap

- **Sec. 2–3:** RTM foundations adapted to QC; geometry of collapse and variable exponents.
- **Sec. 4:** Operational  $(L, T)$  definitions and binning.
- **Sec. 5:** Estimation (ODR/TLS, Theil–Sen, SIMEX), collapse thresholds, uncertainty.
- **Sec. 6:** ECI<sub>QC</sub>(t) construction, fusion, and QA.
- **Sec. 7:** RTM-aware design: scheduling, QEC cadence, modular sizing.
- **Sec. 8:** Experimental protocols (superconducting, ion traps), metrics, and evaluation.
- **Sec. 9:** Results templates and robustness panels.
- **Sec. 10:** Discussion; **Sec. 11:** Limitations; **Sec. 12:** Methods & Reproducibility; **Sec. 13:** Conclusion.

**Takeaway.** By measuring and engineering  $\alpha$ —the multiscale **coherence exponent**—we obtain a reproducible, falsifiable handle on timing structure across the quantum stack. Used with discipline (collapse/QA, heterogeneity gates), RTM delivers a third axis for quantum system design: not only *how accurate* and *how many*, but *how timing stretches with scale*.

## 2. RTM Foundations Adapted to Quantum Computing

This section states the RTM axioms, derives the **power-law** form  $T = \kappa L^\alpha$ , and tailors **clock/gauge** and **collapse** notions to quantum stacks. Throughout,  $L > 0$  is a **scale proxy** (layer-specific) and  $T > 0$  is a **characteristic time** measured in a **fixed environment/bin** (same temperature, firmware, topology, syndrome schedule, utilization band).

### 2.1 Axioms (binwise)

**A1 — Scale semigroup.** For any dilation  $b > 0$ ,

$$T(bL) = f(b) T(L),$$

with  $f(1) = 1$  and  $f(b_1 b_2) = f(b_1)f(b_2)$ .

**A2 — Mild regularity.**  $f$  is measurable (or continuous at  $b = 1$ ).

**A3 — Clock invariance in-bin (multiplicative gauge).** Allowed clock changes multiply  $T$  by a factor  $c > 0$  independent of  $L$  inside the bin (unit changes, constant gain factors, uniform timebase rescaling).

Additive latencies are not gauges in log–log: if observed times satisfy  $T_{\text{obs}} = c \cdot T + b$  with  $b$  constant, then log–log slopes can be biased when  $T$  is not  $\gg b/c$ .

Latency policy (must be stated): (i) calibrate and subtract  $b$  before log-fitting (fit  $\log(T_{\text{obs}} - b)$ ); (ii) model explicitly via  $T_{\text{obs}} = c \cdot \kappa \cdot L^\alpha + b$  (nonlinear fit / profile likelihood); or (iii) restrict fits to regimes where  $T_{\text{obs}} \gg b$  and report sensitivity to plausible  $b$ .

Pure timestamp re-referencing (changing the time origin) does not affect durations once  $T$  is computed as a difference.

**A4 — Binning.** Comparisons are made within bins where the operational environment is stable (e.g., temperature, firmware, topology, QEC cadence, utilization band). If a changepoint is detected or heterogeneity gates fail, the bin must be split; otherwise RTM slope estimates are not comparable.

---

## 2.2 Functional-equation solution → power law

Let  $u = \log L, v = \log T$ . From A1–A2, the multiplicative Cauchy equation gives  $f(b) = b^\alpha$  for some  $\alpha \in \mathbb{R}$ . Hence

$$T(L) = \kappa L^\alpha, v(u) = \alpha u + \log \kappa.$$

**Interpretation.**  $\alpha$  is the **coherence exponent** (slope);  $\kappa$  is a **clock** (intercept).

---

## 2.3 Clocks (gauges) and what is identifiable

A multiplicative clock transform  $g(T) = c \cdot T$  gives  $\log T' = \log T + \log c$ . Thus:

$\alpha$  is invariant to multiplicative clocks (unit changes, uniform gain factors) within a bin.

The intercept encodes clock/units and should not be compared across bins or labs without explicit calibration.

Additive latencies/overheads are not log–log gauges: for  $T_{\text{obs}} = c \cdot T + b$ , slope invariance holds only when  $b=0$ , after correcting  $b$ , or when  $T_{\text{obs}} \gg b$  over the fitted range.

Estimation is therefore slope-first with an explicit latency policy (subtract / model / restrict) stated up front.

QC-specific examples of “clocks”.

Multiplicative clocks (safe): unit conversions; uniform tick-rate rescaling; uniform firmware timing gains applied across widths.

Additive latencies (not gauge in log–log unless handled): fixed readout preamble; fixed scheduling baseline offset that affects measured durations; constant cryo I/O/context-switch baseline.

If any overhead depends on  $L$  inside the bin (e.g., per-channel overhead that grows with multiplexing degree), it is not a valid clock transform; the bin violates RTM and should fail collapse or be split.

### **Remark 2.3a (Additive timestamp offsets are not log–log gauges).**

The clock transform in Definition 2.3 is **multiplicative**:  $T' = cT$  with  $c > 0$ . This implies  $\log T' = \log T + \log c$ , so the **log–log slope** is invariant and only the intercept shifts.

In practice, instrumentation may report an **affine** time:  $\tilde{T} = cT + b$  (constant latency/dead time  $b$ ). Then

$$\log \tilde{T} = \log T + \log c + \log \left(1 + \frac{b}{cT}\right),$$

so slope invariance holds **exactly only when  $b = 0$** , and holds **approximately** only when  $T \gg b/c$  over the fitted range. Otherwise, either (i) estimate/subtract  $b$  before logging (use  $T_{\text{eff}} = \tilde{T} - b$  with  $\tilde{T} > b$ ), or (ii) explicitly model the offset (e.g., fit  $\log(\tilde{T} - b)$  after validating  $b$ ).

## **2.4 Collapse as a binwise specification test**

Given observations  $\{(L_i, T_i)\}_i$  in a bin, define  $x_i = \log L_i$ ,  $y_i = \log T_i$ . Fit a binwise slope  $\hat{\alpha}$  (Section 5) and examine **residuals**

$$\tilde{y}_i := y_i - \hat{\alpha}x_i.$$

**Collapse test.** In a valid RTM bin,  $\tilde{y}$  should be **independent of  $x$**  (up to noise). We operationalize with:

- A regression  $\tilde{y} \sim x$  and require  $R_{\text{collapse}}^2 < \tau$  (default  $\tau = 0.05$ ).
- A **clock placebo**: multiply all  $T_i$  by a constant;  $\hat{\alpha}$  and  $R_{\text{collapse}}^2$  must be unchanged.
- A **smooth check** (LOESS or spline) for visible trend; if present, reject the bin.

**Meaning.** Collapse establishes that, after removing  $\hat{\alpha} \log L$ , only a **gauge** remains (intercept noise), not a trend vs. scale.

---

## 2.5 Variable exponents and finite-window bias

In practice,  $\alpha$  can drift slowly with environment or scale (e.g., across utilization bands or multiplexing factors). Write

$$v(u) = \int_{u_0}^u \alpha(s) ds + \log \kappa(u),$$

with  $|\alpha'(u)| \leq \varepsilon$  small on the window and  $\kappa$  **slowly varying**. For any symmetric window of width  $h$  in  $u$ ,

$$\hat{\alpha}(u; h) = \alpha(u) + O(\varepsilon h) + O(\text{slow-variation}),$$

and

$$R_{\text{collapse}}^2 = O((\varepsilon h)^2).$$

**Rule.** Choose bins/windows small enough that curvature is negligible; otherwise split the bin.

---

## 2.6 Failure modes (should fail)

RTM is designed to **predict its own failure**:

1. **Regime mixing (kinks).** Example: changing the readout chain or syndrome scheduler mid-bin. The log-log plot shows a slope change at  $L^*$ ; collapse fails.
2. **Curvature (non-power).** Example: a multiplexing-dependent overhead that grows nonlinearly with  $L$ . Residuals trend with  $x$ ; collapse fails even after rebinning.

- 
3. **Scale-dependent clocks.** Any “clock” factor  $c(L)$  that depends on  $L$  is not a gauge; it injects  $du$ -components into the 1-form and must be modeled explicitly (or the bin rejected).

---

## 2.7 QC layer mapping (notation and exemplars)

We will use these **canonical**  $(L, T)$  pairs in later sections (others may be added if they pass collapse):

- **Physical:**  
 $L$  =number of **active qubits** (or cluster/coupler degree);  
 $T$  =**stable calibration interval**, **gate latency**, **readout latency**, or **mean time to drift**.
- **QEC:**  
 $L$  =**code distance  $d$**  (or logical-qubit count);  
 $T$  =**cycles to logical failure** at fixed target error.
- **Compiler/Runtime:**  
 $L$  =**circuit width** or **post-mapping depth**;  
 $T$  =**makespan** or **queueing delay**.
- **I/O-Cryo:**  
 $L$  =**multiplexing degree** or readout-channel count;  
 $T$  =**effective readout latency / BER-recovery half-life / p95 queue length (in time)**.

Each family produces a binwise  $\hat{\alpha}_f$ . Only families that **pass collapse** and QA contribute to the fused indicator  $ECI_{QC}(\mathbf{t})$  (Section 6).

---

## 2.8 Why $\alpha$ matters operationally

- **Comparability:**  $\alpha$  is invariant to unit changes and constant overheads, enabling **cross-lab** and **cross-generation** comparison.
  - **Early warning:** significant **drops** in  $\alpha$  (per family or fused) signal **decoherence events** likely to precede spikes in logical error, makespan, or forced recalibrations.
  - **Design lever:** raising  $\alpha$  (without over-layering) via **scheduling**, **QEC cadence**, or **module sizing** improves damping of cross-scale cascades.
-

## 2.9 Summary

RTM in QC reduces to three binwise statements: (i) **power-law** scaling  $T = \kappa L^\alpha$ , (ii) **gauge invariance** (only the slope  $\alpha$  is structural), and (iii) **collapse** as a falsifiable specification test. With careful binning and EIV-aware estimation,  $\alpha$  becomes a reproducible, unit-robust **coherence exponent** that guides both **diagnostics** and **design** across the quantum stack.

## 3. Scale–Clock Geometry for QC (Collapse as Exactness)

We recast RTM for quantum stacks in geometric form. The key object is the **RTM 1-form**

$$\omega = d(\log T) - \alpha(x) d(\log L),$$

defined on a bin  $E$  with **environment** coordinates  $x$  (temperature, firmware, topology, syndrome schedule, utilization band) and **scale**  $u = \log L$ . In this language, **collapse** is equivalent to **exactness/flatness** of  $\omega$ ; regime seams and non-power curvature appear as **holonomy/curvature**. This section states the results and instantiates them with QC failure modes.

---

### 3.1 Spaces, bins, and the RTM 1-form

- **State space.**  $M = X \times \mathbb{R}$  with coordinates  $(x, u)$ , where  $u = \log L$ .
- **Clock potential.**  $v(x, u) = \log T(x, L)$ .
- **RTM 1-form.**  $\omega = dv - \alpha(x) du$  (constant- $\alpha$  case) or  $\omega = dv - \alpha(x, u) du$  (slow drift allowed).

A **clock change** (unit/baseline shift independent of  $L$  inside a bin) is  $v \mapsto v^\# = v + \phi(x)$ . Then

$$\omega \mapsto \omega^\# = \omega + d\phi(x)$$

— a **gauge transformation** by an exact 1-form pulled back from  $X$ . Hence  $\alpha$  is **gauge-invariant**.

---

### 3.2 Collapse $\Leftrightarrow$ exactness/flatness

**Theorem 3.1 (Collapse  $\Leftrightarrow$  exactness).**

On a simply connected bin  $E$ , the following are equivalent:

1. (RTM chart)  $v(x, u) = \alpha(x) u + \log \kappa(x)$  (or  $v = \int \alpha(x, s) ds + \log \kappa(x)$  for slow drift).
2. (**Collapse**) Residual  $\tilde{v} := v - \alpha u$  is independent of  $u$  in  $E$ .
3. (**Exactness**)  $\omega = d\psi$  on  $E$  for some  $\psi(x)$  (no  $u$ -dependence).

**Corollary 3.2 (Flatness test).**

$d\omega = 0$  is necessary and (on simply connected  $E$ ) sufficient for collapse. With  $\alpha = \alpha(x, u)$ ,

$$d\omega = -d\alpha \wedge du.$$

Thus curvature (non-power behavior) or regime mixing gives  $d\alpha / du \neq 0$  and **breaks collapse**.

---

### 3.3 Holonomy and regime seams (QC failure modes)

Define **holonomy** around a closed loop  $\gamma \subset E$ :  $\mathcal{H}(\gamma) = \oint_{\gamma} \omega$ . If  $\mathcal{H}(\gamma) \neq 0$ , collapse cannot hold globally.

#### QC instances.

- **Scheduler seam.** Changing the syndrome-extraction cadence mid-bin (new FPGA image) produces a kink in  $v(u)$ ; loops that cross the seam pick up nonzero holonomy  $\rightarrow$  **rebin**.
  - **Readout chain swap.** A per-channel overhead that *depends on multiplexing* behaves like a scale-dependent clock  $c(L)$ ; this is **not gauge** and injects  $du$ -components  $\rightarrow$  collapse fails (and should).
  - **Thermal drift window.** A slow utilization ramp changes  $\alpha$  across  $u$ ; if  $\partial_u \alpha$  is not small on the window,  $d\omega \neq 0 \rightarrow$  split the bin or shrink the window.
- 

### 3.4 Adiabatic collapse (slowly varying $\alpha$ )

If  $|\partial_u \alpha| \leq \varepsilon$  on a window of width  $h$ ,

$$\tilde{v}(x, u) = v - \alpha(u_0, x) u = \log \kappa(x) + O(\varepsilon h),$$

and the empirical collapse statistic obeys

$$R_{\text{collapse}}^2 = O((\varepsilon h)^2).$$

**Practice.** Choose  $h$  so that  $\varepsilon h \ll 1$ ; otherwise, reduce the bin or model the drift explicitly.

---

### 3.5 Morphisms (reparametrizations) and gauge

Let  $\Phi = (\varphi, \psi)$  map  $(X_A, L_A, v_A) \rightarrow (X_B, L_B, v_B)$ , where  $\varphi: X_A \rightarrow X_B$  reparametrizes environment and  $\psi: X_B \rightarrow \mathbb{R}$  is a clock change. Then

$$\Phi^* \omega_B = \omega_A + d(\psi \circ \varphi).$$

Interpretation: transporting the structure from  $B$  to  $A$  preserves **slope** and alters only the **clock** by an exact form. This formalizes cross-lab/device comparisons when units/baselines differ.

---

### 3.6 Diagnostics and acceptance gates (QC checklist)

1. **Collapse test.** Fit  $\hat{\alpha}$  (Section 5), compute residuals  $\tilde{y} = y - \hat{\alpha}x$ ; require  $R_{\text{collapse}}^2 < 0.05$  and no trend in a nonparametric smooth.
  2. **Clock placebo.** Multiply all  $T$  by a constant;  $\hat{\alpha}$  and  $R_{\text{collapse}}^2$  must be unchanged.
  3. **Changepoints.** Run detectors on  $(x, y)$  and on  $\tilde{y}$ ; any kink  $\Rightarrow$  rebin.
  4. **Window control.** Ensure  $|\partial_u \alpha|$  is small (adiabatic regime).
  5. **Publish/withhold.** Only bins passing 1–4 contribute to  $\text{ECI}_{\text{QC}}(t)$ ; otherwise label `NO_COLLAPSE` or `REGIME_MIX`.
- 

### 3.7 What this buys us operationally

- A **proof-obligation:** show flatness/exactness (collapse) before trusting a slope.
  - A **debugger:** nonzero holonomy localizes seams (scheduler swaps, readout changes).
  - A **tuning rule:** reduce  $h$  or rebin until  $d\omega \approx 0$ ; if impossible, the domain is **non-power**—treat  $\alpha$  as undefined there.
- 

### 3.8 Summary

The scale–clock geometry makes two RTM statements precise for QC:

1.  **$\alpha$  is a gauge-invariant structural quantity**, unaffected by unit/baseline changes;
  2. **Collapse equals exactness/flatness of  $\omega$** , and its failure is informative (curvature or seams).
- We will now leverage this to define **operational**  $(L, T)$  (Sec. 4) and to estimate  $\hat{\alpha}$  robustly under measurement error (Sec. 5).

## 4. Operational $(L, T)$ Definitions and Binning Protocol

This section turns RTM into **measurable practice** for quantum stacks. We define layer-specific  $(L, T)$  pairs, specify **sampling**, **units**, and **guards**, and give a binning protocol that avoids regime mixing. Throughout,  $u = \log L$ ,  $v = \log T$ .

---

### 4.1 Design principles for $(L, T)$

- **One mechanism per family.** Each  $(L, T)$  pair should reflect a single dominant mechanism (e.g., readout pipeline, not a mixture of readout + routing).
- **Monotone  $L$ .**  $L$  should increase with “problem size” at that layer (width, distance, channels, cluster size).
- **Clock independence.** Within a bin, all **constant** overheads in  $T$  are allowed (gauges). Anything **scale-dependent** in  $T$  must be modeled or excluded.
- **Steady sampling.** Use **fixed cadence** collection; record raw timestamps to allow reslicing.

---

### 4.2 Physical layer

#### Candidates for $L$ :

- $L$  = number of **active qubits** in the workload window;
- $L$  = **cluster size** (connected qubits participating simultaneously);
- $L$  = **coupler degree** (average fanout).

#### Candidates for $T$ :

- **Stable calibration interval** (time until any qubit in the cluster exits tolerance);

- **Gate latency** (median single/two-qubit gate duration across the active set);
- **Readout latency** (median per-shot time to valid symbol under fixed thresholds);
- **Mean time to drift** (MTTD) for frequency/phase.

### Instrumentation.

- Log per-shot timestamps; a calibration watchdog recording when thresholds are breached; attach environment tags: temperature band, firmware hash, bias point.

### Non-examples.

- Mixing *both* gate latency and readout latency in the same  $T$ .
  - Letting  $L$  be “qubits defined on chip” (not necessarily active).
- 

## 4.3 Error Correction (QEC)

$L$ : code **distance  $d$**  (primary), or number of **logical qubits** at fixed  $d$ .

$T$ :

- **Cycles to logical failure** at a fixed target error (median or survival quantile);
- **Syndrome-cycle latency** (mean time per cycle under fixed schedule).

### Scheduling notes.

- Freeze a **syndrome schedule** (FPGA image + cadence). Any change  $\Rightarrow$  new bin.
- Record bias (X/Z) and leakage mitigation settings.

### Edge cases.

- If  $T$  is dominated by **rare catastrophic events** (e.g., resonator latch-ups), prefer **conditional medians** (exclude known catastrophic flags) and report a sensitivity panel.
- 

## 4.4 Compiler / Runtime

$L$ : circuit **width** (max concurrent qubits) or **post-mapping depth**; optionally **active layers** after routing.

$T$ :

- **Makespan** (submission  $\rightarrow$  completion);
- **Queueing delay** (submission  $\rightarrow$  start);

- **Rescheduling latency** after a calibration event.

### Controls.

- Fix **routing policy** and **placement heuristic** inside a bin.
  - Stratify by utilization band (e.g., 0–30%, 30–60%, >60%). If utilization drifts, split the bin.
- 

## 4.5 I/O – Cryo / Readout

$L$ : multiplexing degree (channels per line) or number of concurrent readout channels.

$T$ :

- **Readout latency** (median p50 and tail p95);
- **BER recovery half-life** after a controlled burst;
- **Queue p95** expressed in time.

### Instrumentation.

- Timestamp every DMA/ADC burst; log per-channel buffers; annotate firmware versions of DSP.

### Caveat.

- Per-channel overheads that **grow with  $L$**  are *not* gauges; they are genuine scale effects—permissible for RTM—but if the overhead itself changes mid-bin, collapse should fail and trigger a split.
- 

## 4.6 Binning protocol (environment fixing)

A **bin** is a maximal interval where the environment is effectively constant.

### Bin key (example):

BIN

= {platform, temperature band, firmware hash, topology ID, routing policy, syndrome cadence, utilization band}.

### Procedure.

1. **Slice** data by BIN; discard slices with  $< N_{\min}$  distinct  $L$  values (default 6).
2. **Changepoint scan** on  $y = \log T$  vs.  $x = \log L$  (and on residuals if available). If a changepoint is detected (BIC/AIC/PELT), **split**.

3. **Windowing**: for slowly drifting regimes, use sliding windows in  $x$  of width  $h$  such that  $|\partial_u \alpha| h \ll 1$  (from Sec. 3.4).
  4. **Clock placebo**: multiply  $T$  by a constant; the slope  $\hat{\alpha}$  must not change.
- 

## 4.7 Estimation-ready dataset

Create a tidy table per bin with columns:

$$x = \log L, y = \log T, \text{family}, \text{BIN tags}, \text{replicate ID}, \text{timestamp}, \text{weights} ].$$

- **Replicates**. If multiple runs at same  $L$ , aggregate to robust summaries (median  $y$ , MAD-based SE) or pass all and let ODR handle them with replicate weights.
  - **Weights**. Prefer inverse-variance weights from bootstrap over simple counts.
  - **Outliers**. Tag catastrophic events (hardware flags); report both **with** and **without** them.
- 

## 4.8 Acceptance gates (per bin, per family)

A family contributes a slope  $\hat{\alpha}_f$  **only if** all hold:

1. **Coverage**: at least 6 distinct  $L$  points and span  $\geq 0.6$  in  $\log L$ .
2. **Collapse**: regress  $\tilde{y} = y - \hat{\alpha}x$  on  $x$ ; require  $R^2_{\text{collapse}} < 0.05$  and no visible trend (smooth check).
3. **Clock placebo**:  $\hat{\alpha}$  unchanged under  $T \mapsto cT$ .
4. **Changepoints**: none within bin (else split and re-estimate).
5. **EIV fit quality**: ODR/TLS converged; residual diagnostics acceptable (no single leverage point dominates).

Bins or families failing any gate are flagged (`NO_COLLAPSE`, `REGIME_MIX`, `THIN_COVERAGE`, `EIV_FAIL`) and **excluded from fusion**.

---

## 4.9 Examples vs. non-examples (QC-flavored)

- **Good physical family**:  $L$  = active-qubit cluster size;  $T$  = stable calibration interval. Single firmware, stable temperature, no routing change. Collapses cleanly  $\rightarrow$  accept.

- **Bad physical family:** Same, but mid-bin the PLL loop parameters change. Changepoint triggers; split required.
  - **Good QEC family:**  $L = d$ ,  $T = \text{cycles}$  to logical failure, fixed syndrome cadence. Residuals flat  $\rightarrow$  accept.
  - **Bad QEC family:** Mix of two cadences (fast and slow) inside one bin  $\rightarrow$  kink in log-log  $\rightarrow$  reject until split.
  - **Good I/O family:**  $L = \text{multiplexing degree}$ ;  $T = \text{readout latency}$  p95. Firmware constant; latency rises as  $L^\alpha$ , collapse holds  $\rightarrow$  accept.
  - **Bad I/O family:** Switch of DSP firmware that changes per-channel overhead nonlinearly mid-bin  $\rightarrow$  curvature; reject or rebin around the switch.
- 

## 4.10 Summary

- We fixed **operational** ( $L, T$ ) per layer and specified **instrumentation** to make them measurable.
- We defined a **binning protocol** that enforces environment constancy and guards against regime mixing.
- We set **acceptance gates** (coverage, collapse, placebo, changepoints, EIV fit) that determine whether a family's slope enters downstream fusion ( $ECI_{QC}(t)$ ).

## 5. Estimation Under Errors-in-Variables (EIV) and Collapse Thresholds

We now specify **how** to estimate the binwise slope  $\alpha$  robustly when both axes are noisy, and how to decide—via a **collapse threshold**—whether a family's data are RTM-consistent. Throughout,  $x = \log L$ ,  $y = \log T$ . Observations are  $x^{obs} = x + \xi$ ,  $y^{obs} = y + \zeta$  with mean-zero errors.

---

### 5.1 Estimation targets and models

Inside a **fixed bin**, the target is the **local slope**  $\alpha$  in

$$y = \alpha x + c + r(x),$$

with  $r \equiv 0$  under exact RTM or  $|r'(x)| \leq \varepsilon$  under slow drift on a window. Because  $x$  is noisy, **OLS is attenuated**; we use EIV-aware estimators.

**Default target:** point slope  $\alpha$  for the bin; intercept  $c$  is a **gauge** (not compared across bins).

---

## 5.2 Orthogonal Distance Regression (Total Least Squares)

**Definition.** ODR minimizes orthogonal residuals to a line:

$$\min_{\alpha, c} \sum_i \frac{(y_i^{obs} - \alpha x_i^{obs} - c)^2}{\sigma_y^2 + \alpha^2 \sigma_x^2}$$

with effective (possibly heterogeneous)  $(\sigma_x, \sigma_y)$  from replicate variance or bootstrap.

**Practice.**

- Initialize by Theil–Sen (Sec. 5.4) to avoid poor local minima.
- Use **cluster/bootstrap** (replicate or job-level) for CIs.
- If per-point SEs are available, weight them; else use robust Huber weights on orthogonal residuals.

**Convergence gates.**

- Condition number of the centered covariance matrix  $< 10^4$ .
  - Jackknife leverage check: no single point contributes  $> 25\%$  of slope influence.
- 

## 5.3 SIMEX (when $\text{Var}(\xi)$ is known/estimated)

If you can estimate  $\sigma_\xi^2 = \text{Var}(\xi)$  (e.g., repeated  $L$  at the same setting), apply **SIMEX**:

1. For  $\lambda \in \Lambda = \{0.5, 1.0, 1.5, 2.0\}$ , generate pseudo-samples  
 $x_i^{(\lambda)} = x_i^{obs} + \sqrt{\lambda} \tilde{\xi}_i, \tilde{\xi}_i \sim \mathcal{N}(0, \sigma_\xi^2)$ .
2. Fit a naive slope  $\hat{\alpha}(\lambda)$  by ODR or OLS.
3. Fit a quadratic  $\hat{\alpha}(\lambda) = a + b\lambda + c\lambda^2$  and **extrapolate to**  $\lambda = -1$ :  
 $\hat{\alpha}_{\text{SIMEX}} = a - b + c$ .

**Use.** Prefer ODR as the base fitting routine; report SIMEX as a **sensitivity** estimate next to ODR. If  $\sigma_\xi^2$  is uncertain, give a band (low/med/high) for  $\hat{\alpha}_{\text{SIMEX}}$ .

---

## 5.4 Theil-Sen (robust median slope)

The **Theil-Sen** slope is the median of all pairwise slopes

$$\alpha_{ij} = \frac{y_j^{obs} - y_i^{obs}}{x_j^{obs} - x_i^{obs}} \quad (i < j),$$

with a robust intercept from the median of  $y_i^{obs} - \hat{\alpha}x_i^{obs}$ .

**Role.**

- Initialization for ODR.
  - **Outlier-robust** cross-check reported alongside ODR.
  - When EIV is severe and  $\sigma_\xi^2$  is unknown, Theil-Sen may still be stable (expect mild attenuation).
- 

## 5.5 Windowing and finite-window bias

If slow drift is suspected, estimate slopes on **symmetric windows** in  $x$  of width  $h$ . From the adiabatic bias bound,

$$\hat{\alpha}(u; h) = \alpha(u) + O(\varepsilon h),$$

choose  $h$  so that  $\varepsilon h \ll 1$ . Practically: start with  $h \approx 0.8$  in log  $L$  span if coverage allows; shrink until collapse passes (Sec. 5.7) without exploding variance.

---

## 5.6 Uncertainty and diagnostics

- **Bootstrap** (pairs within bin or block/cluster if natural replicates exist) for 50/95% CIs.
  - **Jackknife-after-bootstrap** to detect leverage points.
  - **Residual plots**: orthogonal residual vs.  $x$ ; LOESS smooth must be flat within bands.
  - **EIV adequacy**: if OLS and ODR differ by  $\geq 0.2$  absolute slope **and** ODR CI excludes OLS, report EIV as material.
- 

## 5.7 Collapse threshold (specification gate)

Given  $\hat{\alpha}$ , compute residuals  $\tilde{y}_i = y_i^{obs} - \hat{\alpha}x_i^{obs} - \hat{c}$  and regress  $\tilde{y}$  on  $x$  (with the same weights used in estimation). Define

$$R_{\text{collapse}}^2 := R^2(\tilde{y} \sim x).$$

### Decision rule (default):

- Accept the bin if **all** hold:
    1.  $R_{\text{collapse}}^2 < 0.05$  (or the 95% CI of the slope in  $\tilde{y} \sim x$  contains 0),
    2. LOESS smooth shows no trend,
    3. **Clock placebo:** scaling  $T \mapsto cT$  leaves  $\hat{\alpha}$  and  $R_{\text{collapse}}^2$  unchanged,
    4. Changepoint scan (PELT/BIC) finds none inside the bin.
  - Otherwise flag (NO\_COLLAPSE or REGIME\_MIX) and **do not** publish a slope or include it in fusion.
- 

### 5.8 Coverage and leverage gates

To avoid brittle fits:

- **Distinct Lpoints**  $\geq 6$  and  $\log L$  span  $\geq 0.6$ .
- **Balanced leverage:** the largest leverage point contributes  $\leq 25\%$  of the ODR slope influence.
- **Replicates:** if  $> 3$  replicates per  $L$ , either summarize to a robust mean/SE or pass replicate weights to ODR.

Bins failing these gates are flagged THIN\_COVERAGE or LEVERAGE\_RISK.

---

### 5.9 Putting it together (per-bin algorithm)

1. **Prep:** build the tidy table (Sec. 4.7); run changepoint scan; window if needed.
2. **Init:** compute Theil–Sen slope/intercept; remove obvious catastrophics (keep both versions for sensitivity).
3. **Fit ODR/TLS:** weighted by replicate SEs; obtain  $\hat{\alpha}, \hat{c}$ , bootstrap CIs.
4. **SIMEX (optional):** if  $\sigma_\xi^2$  is available, compute  $\hat{\alpha}_{\text{SIMEX}}$ .

5. **Collapse gate:** compute  $R_{\text{collapse}}^2$ , smooth check, placebo clock.
  6. **Decision:** if all gates pass, **accept**  $\hat{\alpha}$  with uncertainty; else **reject/split**.
  7. **Report:** slope, CI, diagnostics (collapse  $R^2$ , leverage plot, changepoints). Store flags.
- 

## 5.10 What we publish per accepted family

- $\hat{\alpha}_f \pm 50/95\%$  CI (ODR); Theil–Sen as robustness; SIMEX band if applicable.
- Collapse diagnostics:  $R_{\text{collapse}}^2$ , placebo check, window width  $h$ .
- Coverage: # distinct  $L, \log L$  span, leverage summary.
- Notes: any exclusions (catastrophics), changepoint status.

Only accepted families enter **fusion** (Sec. 6). If  $\geq 2$  families pass, we apply random-effects with  $Q, I^2$  and heterogeneity gates; otherwise we report family-wise slopes without fusion.

---

## 5.11 Summary

- Use **ODR/TLS** as the primary EIV estimator; **Theil–Sen** for robust init/check; **SIMEX** when  $\sigma_\xi^2$  is estimable.
- Enforce **collapse** as a **specification test** ( $R_{\text{collapse}}^2 < 0.05$  + placebo + no changepoints).
- Control **finite-window bias** by choosing  $h$  small enough (adiabatic regime) and splitting bins when needed.
- Publish complete **diagnostics** and **flags**; only clean families proceed to fusion and to the real-time  $\text{ECI}_{\text{QC}}(t)$ .

# 6. Building the Real-Time Indicator $\text{ECI}_{\text{QC}}(t)$

We now construct a **single, real-time** coherence indicator for a platform by fusing the **accepted** family-wise slopes  $\{\hat{\alpha}_{f,t}\}$  from Section 5. The fusion is **random-effects** (to acknowledge between-family heterogeneity), runs on a rolling clock, and drives **QA gates** and **decoherence alerts**.

---

## 6.1 Inputs and preconditions (per time $t$ )

For each family  $f \in \mathcal{F}_t$  (Physical, QEC, Compiler/Runtime, I/O-Cryo):

- A binwise estimate  $\hat{\alpha}_{f,t}$  with variance  $\hat{\sigma}_{f,t}^2$  (bootstrap or replicate-weighted),
- Collapse passed (Section 5.7), coverage/leverage gates satisfied (Section 5.8),
- Environment tags (BIN) unchanged within the window that produced  $\hat{\alpha}_{f,t}$ .

A fusion at time  $t$  proceeds **only if**  $|\mathcal{F}_t| \geq 2$ .

---

## 6.2 Random-effects fusion

We estimate the between-family variance  $\hat{\tau}_t^2$  (default **REML**; DerSimonian–Laird as sensitivity). Define weights

$$w_{f,t} = \frac{1}{\hat{\sigma}_{f,t}^2 + \hat{\tau}_t^2}.$$

Then the fused slope and its variance are

$$\hat{\alpha}_{QC}(t) = \frac{\sum_{f \in \mathcal{F}_t} w_{f,t} \hat{\alpha}_{f,t}}{\sum_{f \in \mathcal{F}_t} w_{f,t}}, \quad \text{Var}(\hat{\alpha}_{QC}(t)) = \frac{1}{\sum_{f \in \mathcal{F}_t} w_{f,t}}.$$

Report 50% and 95% intervals via normal approximation or by a **bootstrap-over-families** (resample families with replacement, recompute  $\hat{\tau}_t^2$  and the fused mean).

---

## 6.3 Heterogeneity diagnostics and gates

Compute the fixed-effect baseline

$$w_{f,t}^{FE} = \frac{1}{\hat{\sigma}_{f,t}^2}, \quad \hat{\alpha}_{FE}(t) = \frac{\sum_f w_{f,t}^{FE} \hat{\alpha}_{f,t}}{\sum_f w_{f,t}^{FE}}.$$

**Cochran's  $Q$  and  $I^2$ :**

$$Q_t = \sum_f w_{f,t}^{FE} (\hat{\alpha}_{f,t} - \hat{\alpha}_{FE}(t))^2, \quad I_t^2 = \max \{0, \frac{Q_t - (|\mathcal{F}_t| - 1)}{Q_t}\} \times 100\%.$$

### Fusion gates (pre-registered):

- Proceed with a single number **only if**
    - (i)  $|\mathcal{F}_t| \geq 2$ ,
    - (ii)  $I_t^2 < 50\%$  (*moderate or lower heterogeneity*), and
    - (iii) REML converges with finite  $\hat{\tau}_t^2$  not exceeding a historical cap (e.g.,  $\leq 90$ th percentile over past clean windows).
  - If any fails, **withhold fusion** and publish family-wise  $\hat{\alpha}_{f,t}$ + diagnostics; flag **FAMILY\_DIVERGENCE**.
- 

### 6.4 Real-time operation (rolling windows)

- **Cadence.** Recompute each family's  $\hat{\alpha}_{f,t}$  on a **rolling window** in  $x = \log L$  of width  $h$  (chosen by the adiabatic rule; Sec. 5.5) and a **wall-clock horizon** (e.g., last 7–28 days of data).
  - **Backfill and missingness.** If a family is missing at  $t$ , fuse over the available  $\mathcal{F}_t$  provided  $|\mathcal{F}_t| \geq 2$ ; otherwise **suspend**  $\text{ECI}_{\text{QC}}(t)$  and publish a **THIN\_FAMILIES** flag.
  - **Clock placebo.** Once per day, multiply all contributing  $T$  by a constant and verify  $\hat{\alpha}_{\text{QC}}(t)$  and  $I_t^2$  are unchanged (stored as a QA artifact).
- 

### 6.5 Decoherence events (alerting logic)

We define a **decoherence event** as a significant, QA-clean **drop** in  $\text{ECI}_{\text{QC}}(t)$ , robust to smoothing and not explained by heterogeneity spikes.

#### Filters:

1. **Smoothing:** maintain a 3-point median  $\tilde{\alpha}(t)$  of  $\hat{\alpha}_{\text{QC}}(t)$ .
2. **Z-score:**  $Z(t) = \frac{\tilde{\alpha}(t) - \text{EWMA}_{30}[\tilde{\alpha}]}{\sigma_{\text{EWMA}}(t)}$ .

#### Alert tiers (default):

- **Advisory:**  $Z(t) \leq -1.5$  for  $\geq 2$  consecutive ticks **and**  $I_t^2 < 50\%$ .
- **Watch:**  $Z(t) \leq -2.0$  once **or** persistent  $Z(t) \leq -1.5$  for  $\geq 4$  ticks,  $I_t^2 < 40\%$ .

- **Warning:**  $Z(t) \leq -2.5$  and a coincident family-wise drop ( $\geq 2$  families with  $Z_f \leq -2$ ).

**Playbooks triggered:** throttle scheduling (reduce concurrency/multiplexing), run segmented recalibration, or switch to RTM-aware routing until  $\tilde{\alpha}(t)$  normalizes.

---

## 6.6 Reporting and visualization

- **Primary panel:**  $\hat{\alpha}_{QC}(t)$  with 50/95% bands, heterogeneity ribbon colored by  $I_t^2$  (green <25%, amber 25–50%, red  $\geq 50\%$ ).
  - **Forest plot:** per-family  $\hat{\alpha}_{f,t}$ , weights  $w_{f,t}$ , and CIs; show  $Q_t, I_t^2, \hat{t}_t^2$ .
  - **Collapse dashboard:** per family, show  $R_{collapse}^2$ , LOESS residuals, window width  $h$ , coverage and leverage metrics.
  - **Flags legend:** NO\_COLLAPSE, REGIME\_MIX, LEVERAGE\_RISK, THIN\_COVERAGE, FAMILY\_DIVERGENCE, THIN\_FAMILIES.
- 

## 6.7 Sensitivity and ablation

- Publish the **fixed-effect** summary  $\hat{\alpha}_{FE}(t)$  alongside random-effects.
  - Report DL-based  $\hat{t}_{DL}^2$  as a sensitivity.
  - **Leave-one-family-out:** recompute  $\hat{\alpha}_{QC}^{(-f)}(t)$  to expose dominance.
  - **Clock placebos and shuffle nulls** (shuffle  $L$  within family) must not produce tiered alerts; if they do, review gates.
- 

## 6.8 Governance and provenance

Every fused point stores:

- Source families and their BIN tags,
- Estimator settings (ODR init, bootstrap seeds,  $h$ ),
- Collapse metrics,  $Q_t, I_t^2, \hat{t}_t^2$ ,
- Placebo outcome hashes,
- Versioned code/config (methods YAML).

This ensures **reproducibility** and enables post-mortems when alerts fire.

---

## 6.9 Summary

$\text{ECI}_{\text{QC}}(t)$  is a **random-effects fusion** of QA-clean, binwise slopes. Heterogeneity gates ( $I_t^2 < 50\%$ ,  $|\mathcal{F}_t| \geq 2$ ) prevent misleading single numbers when proxies disagree. Real-time smoothing and Z-scores turn slope dynamics into **actionable alerts for decoherence events**, while dashboards and provenance keep the system auditable.

# 7. RTM-Aware Design: Engineering $\alpha$ without Sacrificing Throughput

This section turns RTM into **design levers**. Goal: increase the **coherence exponent  $\alpha$**  (stronger tempo stratification across scale) while keeping or improving throughput. We give layer-specific controls, optimization targets, and guardrails.

---

## 7.1 Design objective and guardrails

We treat  $\alpha$  as an **operational objective** within a bin:

$$\max_{\text{controls } \theta} \alpha(\theta) \text{ s.t. } \text{throughput} \geq B, \text{ fidelity} \geq F, \text{ collapse passes.}$$

- **Controls**  $\theta$ : scheduler parameters, QEC cadence/jitter, routing constraints, multiplexing limits, module sizes.
- **Constraints**: a throughput floor  $B$  (e.g., jobs/hour), fidelity floor  $F$ , and **collapse gates** (Sec. 5.7).
- **Monitor**: track per-family  $\hat{\alpha}_f$  and the fused  $\hat{\alpha}_{\text{QC}}(t)$  with QA (Sec. 6).

---

## 7.2 Scheduler: batching & variance-aware routing

**Problem.** Long, tightly coupled operations launched in parallel **flatten  $\alpha$**  (fast cascades across scale).

### Controls.

1. **Wavefront batching (readout & long ops).** Partition time into short waves; pack readouts into waves instead of free-running concurrency.

- 2. Staggered resets.** Add small offsets  $\delta \in [-\epsilon, \epsilon]$  to reset times to avoid synch peaks.
  - 3. Low-variance routing.** Prefer routes with **low path-time variance** even if path length increases slightly.

**Objective.** For a job DAG with ops  $o$  having nominal durations  $\tau_o$  and routes  $p(o)$ :

subject to makespan budget. This reduces temporal “pile-ups,” lifting  $\alpha$ .

**Heuristic (greedy, practical).**

- Sort ops by duration desc; assign start times into **waves** so that each wave's total long-op load is balanced.
  - For each route candidate, penalize time-variance and crosstalk score; pick minimum penalized cost.

### 7.3 QEC cadence: avoid phase-lock (jitter/desynchronization)

**Problem.** A fixed syndrome cadence can **phase-lock** with physical noise rhythms, creating cross-layer synchronization  $\rightarrow \alpha_{\text{QEC}}$  falls.

## Controls.

- **Micro-jitter** the cycle period:  $P_k = P(1 + \eta_k)$  with  $\eta_k \sim \mathcal{U}[-\rho, \rho]$ ,  $\rho \ll 1$  (e.g., 1–3%).
  - **Multi-phase extraction:** split the code into sublattices whose cycles are offset by small phases  $\phi_j$ .

**Design rule.** Choose  $\rho$  so that the **main lobe** of the syndrome cycle's line spectrum moves **off** strong peaks of the error PSD while keeping decoder timing valid. Validate by: (i) increased  $\hat{\alpha}_{\text{QEC}}$  vs.  $d$ , (ii) stable logical error at fixed  $d$ .

## 7.4 Gradients and wells of $\alpha$

## Two architectural motifs to steer flows:

- **Gradient:** arrange resources so  $\alpha$  **increases** towards critical compute regions. Small disturbances decay as they travel inward.
- **Well:** create a **high- $\alpha$  basin** around sensitive qubits (e.g., clocking and buffering that slow large-scale cascades).

**Implementation cues.** Increase temporal buffering (queues, damped scheduling) and reduce crosstalk fanout as you approach the “core,” but cap buffering (Sec. 7.1 guardrails) so throughput doesn’t suffer.

---

## 7.5 Modular sizing: pick a sweet spot by balancing intra vs. inter latency

Let total qubits  $Q$  be partitioned into  $Q/m$  modules of size  $m$ . Approximate **characteristic time**:

$$T(m) = A m^a + B \left(\frac{Q}{m}\right)^b \quad (\text{intra-module cost} + \text{interconnect cost}).$$

**Optimal module size** (minimizes  $T$ ):

$$m^* = \left(\frac{B b}{A a}\right)^{\frac{1}{a+b}} Q^{\frac{b}{a+b}}.$$

- $a > 0$ : intra-module scaling (e.g., calibration, routing within module).
- $b > 0$ : inter-module scaling (e.g., photonic/ion link latency).

**Design use.** Measure  $a, b$  empirically (RTM per mechanism), estimate  $A, B$ , compute  $m^*$ . Operate near  $m^*$  and verify that  $\hat{\alpha}$  **does not collapse** (still power-like) in that neighborhood.

---

## 7.6 Multiplexing & I/O: hold tails in check

**Problem.** Aggressive multiplexing reduces per-shot time, but can synchronize queue tails  $\rightarrow \alpha_{\text{IO}} \downarrow$ .

### Controls.

- Cap multiplexing such that the **tail ratio**  $p95/p50$  of readout latency stays below a threshold (e.g.,  $\leq 1.6$ ).
- Use **phase-offset readout windows** across channels to avoid coherent tail growth.

- Buffer sizing: maintain buffer utilization < 70% to avoid tail amplification.

**Signal.** If  $p95/p50$  grows and  $\hat{\alpha}_{IO}$  drops with clean collapse, back off multiplexing and introduce offsets.

---

## 7.7 Online control loop (closed-loop $\alpha$ engineering)

A simple controller to keep  $\alpha$  high under constraints:

```
every Δt:
    estimate {α_f(t), σ_f(t)} per accepted family (Sec. 5)
    if |F_t| ≥ 2 and I^2_t < 50%:
        compute α_QC(t) (Sec. 6)
        if α_QC(t) < α_floor and constraints met:
            apply actions A = {↑wave size, ↑reset jitter ρ, ↑routing penalty on variance,
                                ↓multiplex cap, move toward m*}
        else if throughput < B:
            relax A minimally (keep collapse passing)
    log QA: collapse R^2, I^2_t, flags; revert actions if flags trip
```

- $\alpha_{floor}$ : pre-registered minimal acceptable fused slope.
  - **Revert** any action that causes **NO\_COLLAPSE** or  $I_t^2 \geq 50\%$ .
- 

## 7.8 Safety and validation

- Any intervention must **re-pass collapse** in the affected families.
  - Run A/B windows ( $\geq 2-4$  weeks) with **pre-registered** KPIs: throughput, makespan, logical error, uptime,  $p95/p50$ , and  $\hat{\alpha}_f$ .
  - If  $\alpha$  arises but KPIs worsen beyond budgets, you are **over-layering** (too much buffering). Roll back to the Pareto frontier.
- 

## 7.9 Quick-start playbooks

- If  $\alpha_{QEC} \downarrow$ : add 1–3% cadence jitter; introduce 2–3 phase groups for syndrome; re-measure collapse.

- **If**  $\alpha_{\text{IO}} \downarrow$ : reduce multiplex cap 10–20%; add 1–2 cycle offsets; keep  $p95/p50 \leq 1.6$ .
  - **If**  $\alpha_{\text{runtime}} \downarrow$ : enable readout batching; penalize high-variance routes; cap concurrent long ops per wave.
  - **Architectural planning:** estimate  $a, b, A, B$  and set module size near  $m^*$ ; confirm power-like scaling around that point.
- 

## 7.10 Summary

- **Scheduler** (waves, staggered resets, low-variance routing) and **QEC cadence** (micro-jitter, multi-phase) are first-line levers to **raise**  $\alpha$ .
- **Modular sizing** admits a closed-form optimum  $m^*$  balancing intra/inter costs; operate near it while watching collapse.
- **I/O controls** keep latency tails from synchronizing.
- A **closed-loop controller** maintains  $\alpha$  above a floor under throughput/fidelity budgets.

# 8. Falsifiable Experimental Protocols (Superconducting & Trapped-Ion)

This section specifies **testable** RTM-QC experiments with concrete  $(L, T)$  choices, data collection, analysis plans, and success criteria. Each protocol is binwise (fixed environment) and includes **placebos**, **changepoint guards**, and a **pre-registered** decision table.

---

## 8.1 Common scaffolding (applies to all protocols)

### **BIN (environment) lock.**

{platform; temperature band; firmware hash (FPGA/DSP); topology ID; routing policy; syndrome cadence; utilization band}. Any change  $\Rightarrow$  new bin.

**Data schema (tidy).** For each record:

$$x = \log L, y = \log T, \text{family}, \text{BIN tags}, \text{replicate ID}, \text{timestamp}, \text{weights}]$$

### **QA gates (must pass):**

- Coverage:  $\geq 6$  distinct  $L$ , span  $\geq 0.6$  in  $\log L$ .

- EIV fit converged (ODR), leverage <25%, robust init (Theil-Sen).
- Collapse:  $R_{\text{collapse}}^2 < 0.05$ , no LOESS trend, clock placebo holds.
- Changepoints: none inside bin (else split).

### Outcomes (primary, per family):

- Slope  $\hat{\alpha}_f$  with 50/95% CI; collapse diagnostics.
- For fused results,  $\hat{\alpha}_{\text{QC}}(t)$ ,  $Q$ ,  $I^2$ ,  $\hat{\tau}^2$  (Sec. 6).

### Statistical plan.

Bootstrap CIs (pairs/cluster). Predefine **minimal detectable effect** (MDE) on  $\alpha$  (e.g.,  $\Delta\alpha = 0.15$ ) and **operational KPIs** (throughput, makespan, logical error rate, uptime, p95/p50). Thresholds below.

---

## 8.2 Protocol A — Physical layer (Superconducting)

**Hypothesis (H1-Phys).** Increasing **cluster desynchronization** (staggered resets + readout waves) **raises**  $\alpha_{\text{phys}}$  without exceeding throughput budget.

### Design.

- $L$ : active-qubit **cluster size** (simultaneously engaged).
- $T$ : **stable calibration interval** (time to first out-of-tolerance flag in cluster).
- Arms: **Control** (baseline scheduler) vs. **RTM-aware** (readout batching + staggered resets,  $\pm 2\text{--}4\%$  offsets).
- Duration: 2–4 weeks; interleave arms daily to balance drift.

### Analysis.

- Fit ODR per arm, pass collapse.
- Primary effect:  $\Delta\hat{\alpha}_{\text{phys}} = \hat{\alpha}_{\text{RTM}} - \hat{\alpha}_{\text{CTRL}}$ .
- KPI guardrails: throughput drop  $\leq 5\%$ , no increase in gate/RO error  $> 0.2\sigma$ .

### Success criteria.

- $\Delta\hat{\alpha}_{\text{phys}} \geq 0.15$  and CI excludes 0, **and** guardrails satisfied.
- If collapse fails in any arm, declare **inconclusive** and rebin.

**Placebos.** Multiply  $T$  by a constant;  $\hat{\alpha}$  unchanged. Shuffle  $L$  within day; no significant slope.

---

### 8.3 Protocol B — QEC cadence (Superconducting or Ions)

**Hypothesis (H1-QEC).** Introducing **micro-jitter** (1–3%) in syndrome period and/or **multi-phase extraction** increases  $\alpha_{\text{QEC}}$  vs. code distance  $d$  at fixed decoder.

#### Design.

- $L$ : **code distance**  $d$  (e.g.,  $d \in \{3,5,7,9\}$ ).
- $T$ : **cycles to logical failure** (median or survival quantile at fixed target error).
- Arms: Control (fixed period  $P$ ) vs. Jitter ( $P_k = P(1 + \eta_k)$ ,  $\eta_k \sim \mathcal{U}[-0.02,0.02]$ ) and/or **2–3 phase groups**.
- Keep decoder parameters fixed; no change in noise bias mitigation.

#### Analysis.

- ODR per arm; collapse gate.
- Effect:  $\Delta\hat{\alpha}_{\text{QEC}}$ .
- KPI guardrails: logical error at fixed  $d$  not worse by >5% relative.

#### Success criteria.

- $\Delta\hat{\alpha}_{\text{QEC}} \geq 0.15$  with CI excluding 0 and guardrails pass.

**Diagnostics.** Check PSD of error processes; confirm jitter moves cadence lines off dominant peaks.

---

### 8.4 Protocol C — compiler/runtime scheduling

**Hypothesis (H2-Run).** Wavefront batching of readout and **low-variance routing** reduce synchronization cascades, increasing  $\alpha_{\text{runtime}}$  and lowering makespan tails.

#### Design.

- $L$ : **post-mapping circuit width** (or active layers).
- $T$ : **makespan** (submit→complete).
- Arms: Baseline policy vs. RTM-aware (waves + variance-penalized routing).
- Control utilization band; same job mix across arms.

#### Analysis.

- ODR slope per arm; collapse.
- KPIs: median makespan ( $\leq$  baseline), p95/p50 latency  $\downarrow \geq 10\%$ .

#### **Success criteria.**

- $\Delta\hat{\alpha}_{\text{runtime}} \geq 0.10$  (CI excludes 0) and p95/p50 improves  $\geq 10\%$ .
- 

## **8.5 Protocol D — I/O-Cryo multiplexing**

**Hypothesis (H2-IO).** Phase-offset readout windows across channels maintain or raise  $\alpha_{\text{IO}}$  while reducing p95 tails at a given multiplexing degree.

#### **Design.**

- $L$ : multiplexing degree (channels/line).
- $T$ : readout latency p95 (and p50).
- Arms: Synchronous windows vs. offset windows (phase pattern  $\phi_j$ ).
- Sweep  $L$  across operational range.

#### **Analysis & success.**

- $\Delta\hat{\alpha}_{\text{IO}} \geq 0.10$ ; p95/p50  $\leq 1.6$  in RTM arm over majority of  $L$ ; collapse passes.
- 

## **8.6 Protocol E — Modular sizing (planning study)**

**Hypothesis (H3-Mod).** There exists a module size  $m^*$  that minimizes  $T(m) = Am^a + B(Q/m)^b$  with empirically measured  $a, b > 0$ , and operating near  $m^*$  preserves power-like scaling (collapse holds).

#### **Design.**

- Platforms with photonic/ion links between modules.
- Measure  $T(m)$  by varying module size (or emulating interconnect cost) at fixed total  $Q$ .
- Fit  $a, b, A, B$  via ODR on each term's dataset; compute  $m^*$ .

#### **Success criteria.**

- Observed  $T(m)$  minimized near  $m^*$  (within CI), and log-log fits around  $m^*$  retain collapse (no curvature).

---

## 8.7 Fusion and alerting (cross-protocol)

Across A–D, if  $\geq 2$  families pass gates at overlapping times, compute  $\hat{\alpha}_{\text{QC}}(t)$  (Sec. 6).

**H2 (anticipation):** declare a **decoherence event** if Z-score tiers (Sec. 6.5) are met; test **lead-lag** vs. spikes in logical error/makespan/queues. Additive predictive value is assessed against baselines (fidelity, utilization, temperature) using time series regression with HAC errors; pre-register horizons (e.g., 7–30–90 days).

---

## 8.8 Placebos, shuffles, and robustness

- **Clock placebos:** multiply all  $T$  by constants;  $\hat{\alpha}$  and  $R_{\text{collapse}}^2$  invariant.
  - **Shuffle nulls:** permute  $L$  within day; slopes collapse to  $\sim 0$  (within CI).
  - **Leave-one-family-out** fusion to reveal dominance.
  - **Changepoints:** automatic split if detected; re-estimate on both sides.
- 

## 8.9 Power and duration (rules of thumb)

- With span  $\geq 0.8$  in  $\log L$ , 8–12 distinct  $L$  points, and moderate noise ( $\text{SNR} \approx 5–10$ ), ODR detects  $\Delta\hat{\alpha} \approx 0.10–0.15$  at 95% with  $\approx 200–400$  total observations per arm.
  - If noise is higher or drift suspected, shrink windows (Sec. 5.5) and extend duration.
- 

## 8.10 Decision table (pre-registered)

Outcome	Action
$\Delta\hat{\alpha} \geq \text{MDE and guardrails pass}$	Promote intervention to production in that bin; monitor with $\text{ECI}_{\text{QC}}(t)$ .
$\Delta\hat{\alpha}$ significant but KPI guardrail violated	Tune intensity (e.g., reduce buffering/jitter) and retest.
Collapse fails or heterogeneity high ( $I^2 \geq 50\%$ )	Do not fuse; report family-wise; revisit binning or mechanisms.
No effect ( $(\Delta\hat{\alpha} \approx 0)$ )	Document as <i>scope boundary</i> ; keep as negative control.

## 8.11 Ethics, safety, and reproducibility

- **Safety:** no unsafe RF power increase; jitter bounds keep decoders valid; rollback on `NO_COLLAPSE` or KPI breach.
  - **Reproducibility:** versioned methods YAML (BIN, estimator settings, seeds), public plots (collapse panels, forest plots), and stored placebo/shuffle artifacts.
  - **Transparency:** publish both successes and failures (negative results define scope).
- 

## 8.12 Summary

These protocols make RTM-QC **falsifiable**: each claims a directional change in  $\alpha$  from a specific control, under binwise constancy, with collapse as a specification test and operational guardrails. Success improves not only the slope but also **run-time stability** (tails, recalibrations) without sacrificing throughput.

# 9. Results Templates and Reporting Standards

This section defines **what to publish** once the protocols (Sec. 8) are run. It standardizes figures, tables, robustness panels, and a one-page checklist so that results are interpretable, reproducible, and directly comparable across labs and platforms.

---

## 9.1 Figure set (minimum)

### Fig. 1 — Collapse panels (per accepted family).

Four small multiples per family  $f$  within a bin:

1. **Log-log fit:**  $y = \log T$  vs.  $x = \log L$  with ODR line and 95% band.
2. **Residual vs.  $x$ :**  $\tilde{y} = y - \hat{\alpha}x - \hat{c}$  with LOESS; show  $R^2_{\text{collapse}}$ .
3. **Coverage/leverage:** scatter highlighting leverage points; annotate span in  $\log L$ , # distinct  $L$ .
4. **Placebo check:** overlay of fits before/after  $T \mapsto cT$  (curves coincide).

### Fig. 2 — Forest plot & heterogeneity.

Per time slice (or per experiment arm), show  $\hat{\alpha}_f \pm \text{CI}$ , weights  $w_f$ , the fused  $\hat{\alpha}_{\text{QC}}$  (diamond), and heterogeneity stats:  $Q$ ,  $I^2$ ,  $\hat{\tau}^2$ .

### Fig. 3 — $ECI_{\text{QC}}(t)$ time series.

Rolling fused slope with 50/95% bands; background ribbon colored by  $I^2$  (green <25%,

amber 25–50%, red  $\geq 50\%$ ). Mark **decoherence events** (advisory/watch/warning) and platform events (recalibrations, firmware changes).

#### **Fig. 4 — KPI panel (paired with Fig. 3).**

Aligned time axes for: logical error rate (at fixed  $d$ ), makespan median and p95, queue p95, uptime between recalibrations. Overlay shaded regions for alert tiers from Fig. 3.

#### **Fig. 5 — A/B outcomes (per protocol).**

For each arm: distribution plots (violin/box) of  $\hat{\alpha}_f$ , makespan p95/p50, logical error; include  $\Delta\hat{\alpha}$  with CI and guardrails.

#### **Optional Fig. 6 — Spectral diagnostics (QEC).**

PSD of error processes showing how cadence jitter/multi-phase moves line spectra off dominant peaks.

---

## 9.2 Core tables

**Table 1 — Accepted families (per bin/arm).**

Family	#L pts	$\log L$ span	$\hat{\alpha}_f$ (ODR, 50/95% CI)	Theil-Sen	SIMEX band ( $R^2_{coll}$ )	Leverage max	Flags
Physical	9	1.05	0.62 [0.55, 0.70]	0.60	0.58–0.66	0.02	0.18
QEC	8	0.82	0.74 [0.66, 0.82]	0.71	—	0.03	0.22
...	...	...	...	...	...	...	...

**Table 2 — Fusion & heterogeneity (per time slice or arm).**

Time/Arm	Families	$\hat{\alpha}_{QC} \pm SE$	(Q) (df)	$I^2$	$\hat{\tau}^2$	Fusion?
RTM-aware	3	$0.69 \pm 0.04$	3.2 (2)	37%	0.005	Yes
Control	3	$0.54 \pm 0.05$	6.8 (2)	71%	0.018	No (report family-wise)

**Table 3 — Protocol outcomes (A/B).**

Protocol	Metric	Control	RTM-aware	Effect ( $\Delta$ )	95% CI	Pass guardrail?
A (Phys)	$\hat{\alpha}_{phys}$	0.48	0.64	+0.16	[0.07, 0.25]	✓

Protocol	Metric	Control	RTM-aware	Effect ( $\Delta$ )	95% CI	Pass guardrail?
A (Phys)	Throughput	100%	97%	-3%	[-6, 0]%	✓
B (QEC)	$\hat{\alpha}_{QEC}$	0.68	0.83	+0.15	[0.06, 0.24]	✓
C (Run)	p95/p50	1.85	1.60	-0.25	[-0.35, -0.15]	✓

**Table 4 — Pre-registered thresholds & flags.**

Gate	Threshold	Status
Collapse $R^2$	< 0.05	Pass
Heterogeneity $I^2$	< 50% for fusion	Pass
MDE on $\Delta\alpha$	$\geq 0.10\text{--}0.15$	Pass
KPI guardrails	$\leq 5\%$ throughput loss; $\leq +5\%$ logical error	Pass

### 9.3 Robustness and sensitivity panel

- **Estimators:** ODR (primary), Theil–Sen, SIMEX ( $\pm$  bands for  $\sigma_{\xi}^2$ ).
- **Windows:** repeat with  $h \pm 25\%$ ;  $\hat{\alpha}$  stable and collapse still passing.
- **Placebos:** clock rescaling; **Shuffles:** permute  $L$  within-day—slope  $\rightarrow \sim 0$ .
- **Leave-one-family-out fusion:** report  $\hat{\alpha}_{QC}^{(-f)}$ .
- **Catastrophics:** re-estimate excluding flagged events; show  $\Delta$ .
- **Fixed-effect vs. random-effects:** publish both; divergence implies genuine heterogeneity.

### 9.4 Negative results & scope boundaries

Publish bins/arms that failed:

- **NO\_COLLAPSE** (curvature), **REGIME\_MIX** (kinks), **THIN\_COVERAGE**, **LEVERAGE\_RISK**, **FAMILY\_DIVERGENCE** (high  $I^2$ ).
- Include a short note: suspected mechanism and next steps (rebin, instrumentation change, mechanism isolation). Negative results define **where RTM does not apply**.

---

## 9.5 One-page checklist (for each figure/table set)

- ✓ BIN keys listed and unchanged.
  - ✓ # distinct  $L \geq 6$  and span  $\geq 0.6$ .
  - ✓ ODR converged; Theil–Sen reported; SIMEX (if  $\sigma_\xi^2$  known).
  - ✓ Collapse:  $R^2 < 0.05$ ; placebo OK; no changepoints.
  - ✓ Fusion:  $|\mathcal{F}_t| \geq 2$ ;  $I^2 < 50\%$ ; REML converged.
  - ✓ KPIs: throughput, makespan p95/p50, logical error, uptime—guardrails applied.
  - ✓ Robustness panel completed (windows, shuffles, LOO).
  - ✓ Provenance hashes (methods YAML, seeds, code version) included.
- 

## 9.6 Narrative template (short “Results” text)

*Physical layer.* Across 9 cluster sizes (span 1.05 in  $\log L$ ), RTM-aware scheduling increased the slope from 0.48 to 0.64 ( $\Delta = 0.16$ , 95% CI [0.07,0.25]); residuals showed  $R^2_{\text{collapse}} = 0.02$ . Throughput remained within the 5% guardrail.

*QEC.* With 1–3% cadence jitter,  $\alpha_{\text{QEC}}$  rose from 0.68 to 0.83 ( $\Delta = 0.15$ , CI [0.06,0.24]), logical error at fixed  $d$  did not worsen.

*Runtime.* Wavefront batching and variance-aware routing reduced p95/p50 from 1.85 to 1.60;  $\alpha_{\text{runtime}}$  increased by 0.12.

*Fusion.* Three families passed gates;  $I^2 = 37\%$ . The fused  $\hat{\alpha}_{\text{QC}} = 0.69 \pm 0.04$ . A **watch**-level decoherence alert fired on day 17; it preceded a makespan spike by 3 days.

---

## 9.7 Summary

The templates above ensure every claim is backed by: (i) **collapse** visual and numeric proof, (ii) EIV-aware estimation, (iii) **heterogeneity** accounting for fusion, (iv) KPI guardrails, and (v) complete **robustness** evidence.

# 10. Discussion

This section interprets RTM-QC results, clarifies how a **slope-first** view complements fidelity/QEC paradigms, and lays out trade-offs, risks, and adoption paths.

---

### 10.1 What does a higher $\alpha$ actually buy?

A larger binwise slope  $\alpha$  means **time stretches more steeply with scale**, i.e., larger aggregates slow down *relative* to smaller ones within a stable environment. Operationally:

- **Shock damping:** disturbances at small scale are less likely to synchronize larger layers (runtime → QEC → I/O), reducing cascades that inflate tails (p95/p50), queues, and forced recalibrations.
- **Predictability:** higher  $\alpha$  typically reduces **run-to-run variance** (narrower KPI distributions) because the stack’s “tempo gradient” prevents alignment of rare long events.
- **Control leverage:**  $\alpha$  is unit-agnostic; we can optimize it with scheduler/QEC/interconnect knobs without conflating unit changes (clocks) with structural change.

**Not a substitute for fidelity.** RTM improves **how** timing behaves across scale; it does not increase single/two-qubit fidelities by itself. Gains arrive through fewer cascades and better use of existing fidelity.

---

### 10.2 Complementarity with QEC and compilation

- **QEC:** Traditional design picks code distance  $d$  from error rates. RTM adds a second axis: **cadence geometry**. Slight **desynchronization** (jitter/multi-phase) can raise  $\alpha_{\text{QEC}}$  at fixed  $d$  and decoder, often improving stability without extra overhead.
- **Compilation/runtime:** State-of-the-art routing minimizes depth/length. RTM asks also to minimize **time-variance** and **coincidence of long ops**, which can improve tails even if mean depth changes marginally.

---

### 10.3 Trade-offs and Pareto front

- **Throughput vs. layering:** Raising  $\alpha$  by adding buffers/batching can reduce raw concurrency. We therefore optimize on a **Pareto front** (Sec. 7.1): increase  $\alpha$  *subject to* throughput/fidelity floors.

- **Jitter vs. decoder timing:** Micro-jitter must stay within decoder validity; otherwise you trade higher  $\alpha$  for logical failures.
  - **Modular size:** Operating near  $m^*$  (Sec. 7.5) balances intra/inter costs, but drifting too far (bigger or smaller modules) can either flatten  $\alpha$  (synchronization) or throttle bandwidth.
- 

## 10.4 Failure modes (informative by design)

RTM's **collapse** gate turns failures into diagnostics:

- **NO\_COLLAPSE:** curved log-log  $\rightarrow$  missing mechanism (e.g., scale-dependent “clock” or nonlinear overhead).
- **REGIME\_MIX:** kinks  $\rightarrow$  hidden seams (firmware/scheduler swaps); rebin or split.
- **High  $I^2$ :** proxies disagree  $\rightarrow$  do **not** fuse; inspect per-family controls.

Publishing these cases maps **scope boundaries** (where RTM does *not* apply), which is scientifically useful and prevents overreach.

---

## 10.5 Why a single fused indicator—and when not to use it

**Pros:**  $\text{ECI}_{\text{QC}}(t)$  summarizes multiscale coherence, enabling **alerts** (Sec. 6.5) and trend tracking.

**Cons:** Fusion can hide heterogeneity. Hence the **gates** (at least two families,  $I^2 < 50\%$ , REML convergence). If they fail, publish **family-wise**  $\hat{\alpha}_f$  only; the lack of fusion is itself a result (“the stack is speaking with different slopes”).

---

## 10.6 Relation to time-changed diffusions and queueing

The PDE view (RTM as a **state-dependent clock**) explains why **tails** shrink when  $\alpha$  rises: the **effective dynamic exponent**  $z$  increases, and exit/first-passage times scale more steeply with “radius” (Sec. 6 of the math paper). In queueing terms, scheduling that raises  $\alpha$  **decorrelates** service bursts and dampens tail amplification.

---

## 10.7 External validity and portability

Because  $\alpha$  is **gauge-invariant**, comparisons hold across labs and generations when bins are matched (environment keys). The same pipeline ports to **trapped ions, superconducting,**

**neutral atoms**, and **annealers** with layer-appropriate  $(L, T)$ . What changes is instrumentation; the **collapse logic** and **EIV estimation** remain.

---

## 10.8 Adoption path (practical)

1. **Shadow mode:** compute per-family  $\hat{\alpha}_f$  and collapse panels without changing operations.
  2. **Low-risk knobs:** enable **readout batching**, **staggered resets**, and tiny **cadence jitter** ( $\leq 3\%$ ).
  3. **Close the loop:** bring  $\text{ECI}_{\text{QC}}(t)$  into on-call dashboards with alert tiers and playbooks.
  4. **Architectural planning:** measure  $a, b, A, B$  (Sec. 7.5) to choose module sizes; iterate quarterly.
- 

## 10.9 Open questions

- **Decoder co-design:** how to include  $\alpha$  directly in decoders' scheduling/graph updates?
  - **Learning controllers:** can RL tune  $\alpha$  subject to KPI floors without violating collapse?
  - **Holonomy tests:** practical statistics to distinguish curvature from topological obstructions (global collapse failure).
  - **Cross-layer causality:** when do  $\alpha$  changes at the physical layer *cause* changes at runtime vs. just correlate via utilization?
- 

## 10.10 Takeaway

RTM-QC adds a **third axis**—the **geometry of tempo**—to fidelity and scale. With strict gates (collapse, heterogeneity) and modest controls (batching, jitter, routing variance),  $\alpha$  becomes a reliable lever for **stability and throughput**, yielding early warnings and design guidance while respecting scientific falsifiability.

# 11. Limitations & Scope

**Bin dependence.** RTM is a **binwise** theory. If the environment (temperature, firmware, topology, decoder, utilization) drifts, the slope  $\alpha$  is undefined until the bin is split. Results are only valid within clearly documented BIN keys.

**Proxy choice sensitivity.**  $(L, T)$  proxies must reflect a **single dominant mechanism** per family. Mis-specified proxies (e.g., mixing readout and routing in the same  $T$ ) induce curvature and validly fail collapse.

**Finite-window bias.** When  $\alpha(u)$  drifts, any finite window of width  $h$  incurs  $O(\varepsilon h)$  bias. Our adiabatic guidance mitigates but does not eliminate it; reported  $\hat{\alpha}$  should be interpreted as **local**.

**EIV model assumptions.** ODR/TLS and SIMEX assume well-behaved errors (mean-zero, finite moments) and independence from  $x$ . Heavy-tailed or state-dependent errors require robustness checks (Theil-Sen, bootstrap, sensitivity bands).

**Fusion heterogeneity.** Random-effects fusion is appropriate only when families are **commensurate** and  $I^2 < 50\%$ . Otherwise the single-number indicator is withheld by design; RTM does not force agreement across mechanisms.

**Causality limits.**  $\alpha$  is **structural but not causal** by default. Design sections propose interventions and A/B protocols, yet causal claims require the pre-registered controls and guardrails we specify.

**Scope boundaries.** Systems with **non-power** timing (persistent curvature), **scale-dependent clocks** (overheads that grow with  $L$  inside a bin), or strong **holonomy** (global seams) lie **outside** RTM's applicability. In such domains, treat  $\alpha$  as undefined and publish negative results.

## 12. Methods & Reproducibility

### 12.1 Data schema and BINs

- **BIN key:** {platform, temperature band, firmware hash (FPGA/DSP), topology ID, routing policy, syndrome cadence, utilization band}.
- **Tidy table (per bin):** [x=log L, y=log T, family, BIN tags, replicate\_id, timestamp, weight].
- **Coverage gates:**  $\geq 6$  distinct  $L$ , span  $\geq 0.6$  in  $\log L$ .

### 12.2 Estimation pipeline (per family, per bin)

1. **Changepoint scan:** PELT/BIC on  $(x, y)$  and on residuals if available; split if detected.

2. **Init:** Theil–Sen slope/intercept; flag catastrophics; build replicate weights.
3. **Primary fit:** ODR/TLS (orthogonal residuals) with replicate or bootstrap SEs.
4. **SIMEX (optional):** when  $\sigma_\xi^2$  is estimable; extrapolate to  $\lambda = -1$ .
5. **Collapse test:** regress  $\tilde{y} = y - \hat{\alpha}x - \hat{c}$  on  $x$ ; require  $R_{\text{collapse}}^2 < 0.05$ , flat LOESS, clock placebo holds.
6. **Diagnostics:** leverage  $\leq 25\%$ ; residual plots; window width  $h$  logged.
7. **Accept/Reject:** accept if all gates pass; else flag (NO\_COLLAPSE, REGIME\_MIX, THIN\_COVERAGE, LEVERAGE\_RISK, EIV\_FAIL).

### 12.3 Fusion and heterogeneity (rolling)

- **Weights:**  $w_f = 1/(\hat{\sigma}_f^2 + \hat{\tau}^2)$  with  $\hat{\tau}^2$  via REML (DL as sensitivity).
- **Fused slope:**  $\hat{\alpha}_{\text{QC}} = \sum w_f \hat{\alpha}_f / \sum w_f$ ; **variance:**  $1/\sum w_f$ .
- **Diagnostics:** fixed-effect baseline, **Cochran's Q** and  $I^2$ .
- **Gates:** fuse only if  $|\mathcal{F}| \geq 2$  and  $I^2 < 50\%$ . Otherwise publish family-wise.

### 12.4 Real-time operation and alerts

- **Rolling windows:** sliding horizon in  $x$  (width  $h$ ) and wall-clock (7–28 days).
- **Smoothing:** 3-point median; **Z-score** against 30-day EWMA.
- **Alert tiers:** Advisory/Watch/Warning thresholds (Sec. 6.5).
- **Playbooks:** throttle concurrency, stagger resets, cadence jitter, variance-aware routing; all interventions must re-pass **collapse**.

### 12.5 Robustness & sensitivity

- **Estimators:** publish ODR (primary), Theil–Sen, SIMEX bands.
- **Windows:**  $\pm 25\% h$  sensitivity;  $\hat{\alpha}$  stability required.
- **Placebos & shuffles:** clock rescaling invariance;  $L$ -shuffles yield near-zero slopes.
- **Leave-one-family-out** fusion; **fixed-effect** vs **random-effects** comparison.

### 12.6 Provenance (methods YAML)

- BIN keys, estimator settings, bootstrap seeds, SIMEX  $\Lambda$ , window  $h$ , collapse thresholds, heterogeneity gates, versions of analysis code.

- All plots and numbers include hash of the methods YAML; re-runs with the same YAML reproduce numbers within bootstrap noise.

## 13. Conclusion & Outlook

We presented **RTM-aware quantum computing (RTM-QC)**: a **slope-first** framework that measures and **engineers** the geometry of time across scale. Inside stable bins, the characteristic time  $T$  scales with a size proxy  $L$  as  $T \propto L^\alpha$ ; the **coherence exponent**  $\alpha$  is invariant to clocks and thus comparable across devices, stacks, and labs. With **collapse** as a falsifiable gate and **errors-in-variables** estimation,  $\alpha$  becomes a reliable operational signal. Fusing clean, layer-wise slopes yields a real-time  $ECI_{QC}(t)$  that supports **early warnings** (decoherence events) and **design decisions** (scheduler, QEC cadence, modular sizing, I/O offsets).

**What this adds.** RTM-QC complements fidelity/QEC by introducing a third axis—**tempo geometry**—that explains and controls tails, queues, and synchronization cascades. Modest, reversible controls (batching, staggered resets, micro-jitter, low-variance routing) can **raise**  $\alpha$  without degrading throughput or fidelity when used with guardrails.

**What it does not do.** RTM-QC does not replace physical improvements (fidelities,  $T_1/T_2$ ), nor does it guarantee causality without the A/B protocols and guardrails we specify. Failures to collapse, high heterogeneity, or regime seams are **informative**, delineating scope boundaries rather than inviting post-hoc fixes.

### Near-term agenda.

1. **Run the protocols** (Sec. 8) on superconducting and ion platforms; publish both successes and negatives with full collapse/fusion diagnostics.
2. **Close the loop:** deploy  $ECI_{QC}(t)$  dashboards and alert playbooks in production; evaluate lead-lag vs. KPI spikes.
3. **Co-design with decoders** and compilers so cadence and routing optimize  $\alpha$  subject to throughput/fidelity floors.
4. **Standardize reporting:** figures/tables in Sec. 9, methods YAML, and open robustness artifacts.

**Longer-term questions.** Incorporate  $\alpha$  into **time-changed diffusion models** of queues; develop **holonomy tests** to distinguish curvature from seams; extend to **modular networks** and **neutral-atom** platforms; integrate learning-based controllers that respect collapse gates.

**Bottom line.** RTM-QC gives quantum teams a **unit-robust, falsifiable lever** over multiscale timing. Measure the slope, **validate by collapse**, fuse when families agree, and **engineer**  $\alpha$ —not as a slogan, but as a reproducible practice to deliver more stable and efficient quantum computation.

## Appendices

### Appendix A — Mathematical Background (RTM essentials for QC)

#### A.1 Semigroup → power law

Assume binwise scale semigroup  $T(bL) = f(b)T(L)$ ,  $f(1) = 1$ , and measurability near  $b = 1$ . Then  $f(b) = b^\alpha$  and

$$T(L) = \kappa L^\alpha, v(u) = \log T = \alpha u + \log \kappa, u = \log L.$$

$\alpha$  is **gauge-invariant**;  $\kappa$  is a **clock**.

#### A.2 1-form & collapse

Define the RTM 1-form  $\omega = dv - \alpha du$ . **Collapse** (residual independence of  $v - \alpha u$  from  $u$ ) is equivalent to **exactness** of  $\omega$  on a simply connected bin:

$$\omega = d\psi(x), d\omega = 0, \psi \text{ independent of } u.$$

If  $\alpha = \alpha(x, u)$ , then  $d\omega = -d\alpha \wedge du$ ; nonzero curvature breaks collapse.

#### A.3 Variable exponents (finite-window bias)

For slowly varying  $\alpha(u)$ :

$$v(u) = \int_{u_0}^u \alpha(s) ds + \log \kappa(u), \hat{\alpha}(u; h) = \alpha(u) + O(\varepsilon h),$$

and  $R_{\text{collapse}}^2 = O((\varepsilon h)^2)$  for window width  $h$ .

## Appendix B — Estimators & Algorithms

### B.1 Orthogonal Distance Regression (TLS/ODR)

Minimize orthogonal residuals:

$$\min_{\alpha, c} \sum_i \frac{(y_i - \alpha x_i - c)^2}{\sigma_{y,i}^2 + \alpha^2 \sigma_{x,i}^2}.$$

**Init:** Theil-Sen; **CIs:** bootstrap pairs/cluster; **checks:** condition number  $< 10^4$ ; max leverage  $< 25\%$ .

## B.2 Theil-Sen

Median of pairwise slopes  $\alpha_{ij} = (y_j - y_i)/(x_j - x_i)$ ; robust to outliers; mild EIV attenuation.

## B.3 SIMEX (optional)

If  $\sigma_\xi^2 = \text{Var}(\xi)$  is estimable, simulate  $x^{(\lambda)} = x^{obs} + \sqrt{\lambda}\tilde{\xi}$  and extrapolate  $\hat{\alpha}(\lambda)$  to  $\lambda = -1$ .

## B.4 Collapse gate

Regress residuals  $\tilde{y} = y - \hat{\alpha}x - \hat{c}$  on  $x$ ; require  $R^2_{\text{collapse}} < 0.05$  and flat LOESS; pass clock placebo.

---

## Appendix C — Protocol Cards (copy-paste templates)

### C.1 Physical (staggered resets + readout waves)

- **L/T:**  $L$  = active cluster size;  $T$  = stable calibration interval.
- **Arms:** Control vs RTM-aware (waves + 2–4% reset offsets).
- **Duration:** 2–4 weeks, interleaved.
- **Success:**  $\Delta\alpha_{\text{phys}} \geq 0.15$  (95% CI excludes 0), throughput loss  $\leq 5\%$ , collapse passes.

### C.2 QEC (micro-jitter / multi-phase)

- **L/T:**  $L = d$ ;  $T$  = cycles to logical failure.
- **Arms:** Fixed period vs  $Pk = P(1 + \eta k)$ ,  $|\eta k| \leq 0.02$  or 2–3 phase groups.
- **Success:**  $\Delta\alpha_{\text{QEC}} \geq 0.15$ , no logical-error regression ( $> 5\%$ ) at fixed  $d$ .

### C.3 Runtime (batching + low-variance routing)

- **L/T:**  $L$  = post-mapping width;  $T$  = makespan.
- **Arms:** Baseline vs waveform + variance-penalized routing.

- **Success:**  $\Delta\alpha_{\text{runtime}} \geq 0.10$  and p95/p50 latency  $\downarrow \geq 10\%$ .

#### C.4 I/O (phase-offset windows)

- **L/T:**  $L$  = multiplexing degree;  $T$  = readout latency p95 (and p50).
  - **Arms:** Synchronous vs phase-offset windows.
  - **Success:**  $\Delta\alpha_{\text{IO}} \geq 0.10$ , p95/p50  $\leq 1.6$  over majority of  $L$ .
- 

#### Appendix D — Methods YAML (skeleton)

```

bin:
  platform: "SC"          # or "IONS", "NA"
  temperature_band: "10-15mK"
  firmware_hash: "fpga_1.4.2_dsp_0.9.8"
  topology_id: "mesh-v3"
  routing_policy: "baseline" # or "rtm-aware"
  syndrome_cadence: "P=3.2us, jitter=0%"
  utilization_band: "30-60%"

estimation:
  min_L_points: 6
  min_logL_span: 0.6
  eiv: "odr"
  odr:
    init: "theil-sen"
    leverage_cap: 0.25
  bootstrap: {clusters: true, reps: 2000, seed: 123}

simex:
  enabled: false
  lambda: [0.5,1.0,1.5,2.0]

collapse:
  r2_threshold: 0.05
  placebo_clock: true
  changepoint_scan: {method: "PELT", penalty: "BIC"}

```

```

fusion:
  heterogeneity_gate_I2: 0.5
  tau2_method: "REML"
  min_families: 2

eci_rt:
  window_logL: 0.8
  horizon_days: 14
  smoothing: "median3"
  alert:
    z_advisory: -1.5
    z_watch: -2.0
    z_warning: -2.5

```

## Appendix E — Notation Glossary

- $L$ : scale proxy (layer-specific);  $u = \log L$ .
- $T$ : characteristic time;  $v = \log T$ .
- $\alpha$ : **coherence exponent** (slope; clock-invariant).
- **Bin**: environment slice with fixed {platform, temperature band, firmware hash, topology ID, routing policy, syndrome cadence, utilization band}.
- **Collapse**:  $R^2(\tilde{y} \sim x) < 0.05$  for  $\tilde{y} = y - \hat{\alpha}x$ ; residuals show no trend vs  $x$ .
- **ECI<sub>QC</sub>( $t$ )**: fused slope via random-effects at time  $t$ .
- $Q, I^2, \tau^2$ : heterogeneity statistics for fusion.
- ODR/TLS, Theil–Sen, SIMEX: slope estimators under EIV.
- **Adiabatic window**: width  $h$  in  $u$  where  $|\partial_u \alpha| h \ll 1$ .

---

## Appendix F — Reproducible Figure Recipes (minimal)

- **Collapse panel:**
  - Fit ODR; compute residuals  $\tilde{y}$ .
  - Plot  $y$  vs  $x$  + ODR band; residual vs  $x$  with LOESS.
  - Annotate  $R^2_{\text{collapse}}$ , # $L$ , span, leverage.
- **Forest plot:**
  - For accepted families, display  $\hat{\alpha}_f \pm \text{CI}$ ; compute  $w_f, Q, I^2, \hat{\tau}^2$ .
  - Overlay fused  $\hat{\alpha}_{\text{QC}}$ .
- **ECI<sub>QC</sub>( $t$ ):**
  - Rolling fusion; show 50/95% bands; background colored by  $I^2$  tiers; mark alert tiers.