



Rhythmic Economics

Measuring Systemic Resilience with the RTM Coherence Exponent

Álvaro Quiceno

Abstract

We propose a rhythmic view of economic dynamics grounded in the RTM (Relativity of Time in Multiscale systems) principle that characteristic times scale with system size as $\tau \propto L^\alpha$. Translating this into economics, we define an Economic Coherence Exponent α that captures how quickly processes at different scales, households, firms, sectors, markets, settle, propagate, or recover. From α we construct a real-time Economic Coherence Index (ECI): a slope-first, errors-in-variables meta-estimate obtained from multiple independent proxies of economic “length” L (network size, capitalization tier) and economic “time” τ (recovery half-lives, relaxation times). Because shifts of clocks or levels affect intercepts rather than slopes, α is designed to be robust to unit changes and regime-level confounders.

Computational validation. We implement and test the ECI framework through three simulation suites. S1 demonstrates α estimation from market cap tiers and recovery times across five market regimes (stable growth $\alpha \approx 0.45$, pre-crisis $\alpha \approx 0.35$, crisis $\alpha \approx 0.20$), recovering the true exponent within 0.6% error. Meta-analysis across four proxy families (recovery half-life, volatility persistence, autocorrelation decay, order flow relaxation) yields combined ECI estimates with quantified heterogeneity (I^2). S2 validates Hypothesis H2, that α decline anticipates recessions, through backtesting on three episodes (2001 Dot-Com, 2008 GFC, 2020 COVID), finding mean lead times of 9 months with larger α drawdowns ($\Delta\alpha = 0.14$ -0.27) preceding more severe crises. S3 demonstrates cross-country variation: α correlates strongly with crisis frequency ($r = -0.91$) and average drawdown ($r = -0.95$), with developed economies ($\alpha \approx 0.48$ -0.55) showing higher resilience than emerging markets ($\alpha \approx 0.25$ -0.35).

We articulate falsifiable hypotheses: H1 (Resilience), higher baseline α predicts smaller drawdowns; H2 (Anticipation), sharp α drops precede recessions by 6-18 months; H3 (Cascade), α is non-decreasing across aggregation layers. The framework offers a complementary early-warning signal distinct from volatility or leverage metrics, with implications for coherence-aware stress testing and macroprudential policy.

Empirical validation → (Chapters 11 & 12). Beyond simulation, we subject the framework to a forensic stress test using high-frequency Bitcoin market microstructure and historical crashes across the S&P 500 and Gold. An earlier version of this work reported a large in-sample separation between healthy-market and crash DFA- α values ($d = -1.45$) with a ~ 10 -day lead, derived from a 13-event cross-market table. That table (`crash_alpha_analysis.csv`) was subsequently found to contain **synthetically generated α values** (reproducible from a fixed random seed), and its statistics are therefore withdrawn as empirical claims. Independent validation on real market data (5,104 days of S&P 500 returns spanning the Dot-Com, GFC, and COVID crises) finds that the genuine crisis-vs-calm DFA- α separation is small and not significant (Cohen’s $d \approx -0.2$, $p \approx 0.22$), and that α does **not** reliably decline before crises. The framework’s surviving empirical contribution is the **Multi-Scale Coherence** metric (below).

Data-provenance correction (July 2026). The 13-event α table underlying the $d = -1.45$ in-sample result was found to be synthetically generated; those numbers are withdrawn (see Chapter 12.6). The findings below are computed from raw market data and are unaffected. **Red Team caveat (April 2026).** Independent adversarial validation (5 analytical flanks) found that out-of-sample crash prediction accuracy is **25%** (1 of 4 post-2022 events correctly classified when threshold trained on pre-2022 crashes). The α - drop patterns in Chapters 11 and 12 are in-sample forensic observations, not validated prospective predictions. The surviving novel finding is the

Multi-Scale Coherence metric: during BTC crash months (COVID March 2020, FTX November 2022), the cross-scale standard deviation of α across time windows is $\sigma = 0.031\text{-}0.034$; during the control month (September 2023), $\sigma = 0.310$, a 10x difference. This effect replicates in direction, out-of-sample, on daily S&P 500 data across three crises (crisis cross-scale σ below calm σ ; Mann-Whitney $p = 0.039$), though the equity separation is far smaller ($\sim 1.3\times$) than the intraday BTC separation ($\sim 10\times$). The metric is real but its magnitude is asset- and timescale-dependent. This metric is RTM-native and not measured by standard financial indicators. Details: Chapter 12.7.

Monte Carlo simulations across 16 global markets yield a tail exponent of $\alpha = 2.966 \pm 0.236$, consistent with the RTM “Inverse Cubic Law” and with the empirically established inverse cubic law of price fluctuations (Gabaix et al. 2003). This is a convergent result, correctly framed as confirming known financial scaling, not as rejecting Gaussian models categorically.

1. Introduction

1.1 Motivation

Standard macro-financial indicators (GDP growth, inflation, unemployment, volatility indices) summarize levels or dispersion but rarely measure how **tempo** and **organization across scales** change as systems evolve. Yet crises, supply-chain breakdowns, and sudden sentiment cascades are multiscale phenomena: the *time it takes* for a disruption to propagate or dissipate depends on the *size* and *connectivity* of the structures it traverses. RTM, the empirical observation that characteristic times scale with size via a power law, offers a compact way to model that dependency. Bringing RTM into economics suggests a single, interpretable quantity, the **Economic Coherence Exponent** α_{econ} , that quantifies how “fast” or “structured” an economy is at a given moment, across layers.

1.2 From RTM to α_{econ}

RTM’s core claim is a scaling symmetry: if two subsystems are geometrically similar but differ in scale, their characteristic times relate as $T \propto L^\alpha$. Estimating α relies on **slopes** in log–log space within fixed environments; shifts in clocks, units, or baselines alter **intercepts** but not **slopes**. In economics we cast “length” L as a scale proxy, firm size, capitalization, supply-chain path length, network degree, or jurisdictional scope, and “time” T as a persistence or relaxation metric, recovery half-lives, order-book resiliency windows, lead-time decay, sentiment relaxation. The **Economic Coherence Index (ECI)** is a rolling, errors-in-variables meta-estimate of α_{econ} that blends multiple (L, T) families with cross-validation and uncertainty quantification.

1.3 What α_{econ} means (intuition)

Higher α_{econ} implies that larger structures slow down *more than proportionally*, often reflecting **greater organization and controllability**: information is filtered, buffers exist, and flows are orchestrated. Lower α_{econ} implies a flatter time-scale gradient: shocks traverse layers rapidly, sometimes beneficial for throughput but hazardous for stability. Thus α_{econ} reframes the classic trade-off between raw speed and systemic resilience as a tunable **coherence** parameter.

1.4 How ECI differs from familiar metrics

Volatility (e.g., VIX) measures dispersion at a given scale; leverage measures balance-sheet sensitivity; liquidity measures transaction cost/market depth. **ECI measures the *slope of time-vs-scale***: a structural property that complements those signals. Because ECI is built on slopes, it is comparatively robust to unit choices, nominal drifts, and many level-shifting policy changes.

1.5 Empirical program

We will (i) construct (L, T) pairs across independent domains, market microstructure, logistics, credit rollover, information decay, (ii) estimate α_{econ} within environment-fixed bins via robust errors-in-variables regression, (iii) validate internal **collapse** (curves coincide when rescaled by L^α inside a bin), (iv) build a rolling ECI(t) with uncertainty, and (v) test H1–H3 on retrospective episodes and, prospectively, in live pilots. Failures of slope separation or collapse are recorded as **negative results** delimiting ECI's scope.

1.6 Systematic Empirical Validation: Phase Transitions and Market Microstructure (Chapters 11 & 12)

Within the analytical paradigm of RTM, a market crash is not a purely exogenous or random panic event, but rather the final result of a quantifiable topological phase transition. To bridge the gap between theoretical thermodynamics and applied risk management, we utilize historical crashes and the high-frequency Bitcoin market as a computational wind tunnel.

An initial analysis of 13 major historical crashes appeared to show strong correlations between DFA- α decorrelation and crash severity; that analysis was later found to rest on synthetically generated α values and is withdrawn (Chapter 12.6). On real market data the DFA- α crisis signal is weak and does not anticipate crashes (Chapter 12.7). To address ecological fallacies from static point-estimates, we deployed Continuous Errors-in-Variables (ODR) modeling and Monte Carlo noise injection across thousands of trading hours.

The previously reported shift from laminar baseline coherence ($\alpha \approx 0.55$) to anti-persistence during crashes ($\alpha \approx 0.46$, Cohen's $d = -1.45$, ~ 10 -day in-sample lead) derived from the synthetic α table and is withdrawn (Chapter 12.6). Recovery time scaling follows a punishing nonlinear slope (3.59 ± 0.70), consistent with the known asymmetry between crash speed and recovery duration. Return distributions across 16 global markets converge to $\alpha = 2.966 \pm 0.236$, consistent with the empirically established inverse cubic law (Gabaix et al. 2003).

Important limitation: Out-of-sample validation (Red Team, April 2026) found 25% accuracy (1/4 post-2022 events) when applying the trained α -drop threshold to new data. The forensic patterns (COVID, FTX, GFC) are in-sample results. Whether the structural signal generalizes prospectively requires further validation on held-out data. The Multi-Scale Coherence metric (σ of α across time scales) shows more promise as a novel indicator, see Chapter 12.7.

Empirical Validation of Phase Bifurcation in High-Frequency Markets (Chapters 11)

This chapter stress-tests the RTM Real-Time Monitor against the extreme variance of the Bitcoin microstructure. By abandoning static daily closes and injecting the full continuous noise profile of minute-by-minute OHLCV data, we track the exact moment of structural fracture. The continuous analysis isolates the Phase Bifurcation threshold ($\alpha < 0.5$), distinguishing mechanical liquidity failures (e.g., March 2020) from high-viscosity political stress events ($\alpha > 0.6$, e.g., May 2021). The October 2025 event showed an elevated α reading consistent with a

technical anomaly (subsequently attributed to a Binance platform glitch) rather than a fundamental crash. The 15-hour lead observation is a single in-sample case and should not be generalized as validated prospective prediction until replicated on out-of-sample events.

Empirical Analysis: The Collapse of α as a Predictive Signal (Chapters 12)

This chapter departs from the traditional Gaussian assumptions of market recoveries and crash predictions. Initial naive OLS models predicting crash recoveries suffered from massive attenuation bias due to the ambiguous boundaries of “market recovery.” By applying Orthogonal Distance Regression (ODR) to absorb a 20% measurement noise margin, we reveal that recovery time scaling is substantially more punishing (slope = 3.59 ± 0.70) than previously modeled.

An earlier version reported that a Monte Carlo analysis of 13 major crashes showed structural α -decorrelation preceding price troughs by ~ 9.75 days ($d = -1.45$). The underlying 13-event table was subsequently found to contain synthetically generated α values (Chapter 12.6), so these figures are **withdrawn**. Independent analysis on real S&P 500 daily returns (Chapter 12.7) finds no significant crisis-vs-calm α separation ($d \approx -0.2$, $p \approx 0.22$) and no pre-crash α decline. The recovery-time scaling result (above) and the multi-scale coherence metric (Chapter 12.7) are computed from independent data and are retained.

2. RTM Primer for Economists

This section distills the Multiscale Temporal Relativity (RTM) into tools you can use with economic data. We explain the master scaling, why **slopes** (not levels) are the robust signal, how to estimate the coherence exponent, and what would falsify the approach.

2.1 The master law and why a power law

Claim (RTM). In multiscale systems, characteristic times T scale with system size L via a power law:

$$T = \kappa L^\alpha,$$

where $\kappa > 0$ is a scale factor determined by the environment (units, baseline frictions, “clock”), and α is the **coherence exponent**. Taking logs:

$$\log T = \alpha \log L + \log \kappa.$$

Why a power law? If (i) rescaling the system by a factor $b > 0$ simply rescales its characteristic time by a deterministic function $f(b)$, and (ii) independent rescalings compose ($f(b_1 b_2) = f(b_1) f(b_2)$), then $f(b) = b^\alpha$. This Cauchy-type functional equation is the standard route from *scale symmetry* to *power laws*. In economics, the “system” might be a firm, a supply chain, or a market subgraph; “time” might be a relaxation, renewal, or recovery time.

2.2 Slope vs. intercept: the “clock-invariance” insight

Changes of units or uniform measurement clocks typically **multiply** durations by a constant c (i.e., $T' = cT$), which adds a constant to $\log T$ and shifts the intercept without changing the slope α , provided the environment is fixed.

However, regime shifts that add a roughly constant delay b to observed durations ($T_{\text{obs}} = T + b$) do **not** add a constant to $\log T$ and can bias slope estimates unless modeled or corrected (e.g., fit $\log(T_{\text{obs}} - b)$ after validating b , or restrict to $T \gg b$). Hence:

Intercept: level effects (unit changes, multiplicative clocks, baseline rescalings; additive offsets only after correction).

Slope: structural scaling (how time grows with scale within a fixed environment).

Practical upshot: to get robustness, estimate slopes in environment-fixed bins **and** audit/adjust for additive offsets before logging.

Practical upshot. If you want something robust to level shifts and many policy rebasings, estimate **slopes in environment-fixed bins** (e.g., country $\times$ policy regime $\times$ quarter).

2.3 What α is, and is not

- α is **not** a temperature, volatility, or a “speed dial.” It is a **gradient of tempo across scale**.
- α is **not** the dynamic exponent z from critical phenomena; here we use a direct operational slope between empirical L and T .
- Heuristically, **larger α implies greater organization/persistence**: bigger structures slow down more than proportionally, suggesting buffering, staged decision-making, and filtered information flow. **Smaller α implies flatter cross-scale timing**, fast propagation and potential exposure to cascades.

2.4 Mapping economics to L and T

RTM requires pairs (L, T) drawn from a **fixed environment**:

- **Candidate L (size/scale) proxies**
 - Micro: employees, balance-sheet size, capitalization tier, supplier count.
 - Meso: supply-chain path length, network degree/centrality, regional market size.
 - Macro: trade network breadth, interbank network size, jurisdictional scope.
- **Candidate T (time/persistence) proxies**
 - Micro: order-to-cash cycle, time-to-fill vacancy, microprice reversion half-life.
 - Meso: shipment lead-time persistence, inventory replenishment half-life.
 - Macro: output-gap recovery half-life after shock, credit rollover renewal time, sentiment/news decay time.

Compatibility rule. Within a bin, L must be monotone with scale and T a **characteristic time** tied to *the same process layer*. Mixing layers in one bin (e.g., firm-level L with sector-level T) violates comparability.

2.5 Estimating α : slope-first with measurement error

Real economic L and T are noisy. A vanilla OLS regression of $\log T$ on $\log L$ is biased when L has error. Use **errors-in-variables (EIV)** or robust alternatives:

- **Total least squares / orthogonal distance regression** for symmetric noise.
- **Theil-Sen** slope for outlier robustness (median of pairwise slopes).
- **SIMEX** (simulation-extrapolation) if you can approximate the noise level of L .

Bin-by-bin pipeline (sketch).

1. Fix an environment (e.g., US manufacturing, 2012–2019, stable policy).
2. Partition into size tiers (or sliding windows) such that the *ambient clock* is approximately constant.
3. In each bin, fit $\log T = \alpha \log L + c$ with EIV; report $\hat{\alpha}$ and CI.
4. **Collapse test:** rescale each curve by $L^{\hat{\alpha}}$ and verify that residual structure vanishes within that bin (curves “collapse”). Failure to collapse $\Rightarrow$ the bin mixes incompatible regimes or α is not well-defined there.
5. Combine multiple independent (L, T) families via **random-effects meta-analysis** to get $\hat{\alpha}_{\text{econ}}$ and uncertainty.

2.6 Interpreting levels vs. changes in α_{econ}

- **Level** $\bar{\alpha}_{\text{econ}}$: background coherence/resilience of a system over a period.
- **Change** $\Delta\alpha_{\text{econ}}$: early-warning signal; sudden drops indicate **decoherence** (cross-scale timings become too similar $\rightarrow$ shocks traverse quickly). Sudden rises can indicate re-organization, sometimes at the cost of raw throughput.

Design trade-off. Higher α often means slower raw tempo at the largest scales, but better control and stability (fewer catastrophic cascades). Lower α increases throughput but can raise systemic risk.

2.7 Universality bands (heuristic guidance)

While RTM does not *fix* a universal α , empirical bands help interpret ranges:

- **Flattened/co-moving** ($\alpha \approx 1$): times scale $\sim$ linearly with size, fast propagation, minimal buffering.
- **Diffusive/mediated** ($\alpha \approx 2$): larger structures slow down more; coordination layers are visible.
- **Hierarchically buffered** ($\alpha > 2$): deep staging, long planning cycles, substantial slack.

These are **interpretive heuristics**, not hard thresholds; the data and collapse tests arbitrate.

2.8 Falsifiability: where RTM should fail in economics

RTM makes claims strong enough to be **wrong**:

- **No slope separation:** If, within a fixed environment, $\partial \log T / \partial \log L$ is indistinguishable from zero (or wildly unstable) across multiple independent (L, T) families, RTM is not informative for that domain.
- **No collapse:** If rescaling by $L^{\hat{\alpha}}$ fails to collapse curves within a bin, α is not well-defined there.
- **Reverse cascade:** If aggregation layers show **decreasing** α (macro faster across scale than micro) systematically and robustly, the RTM cascade signature fails.
- **Directionality symmetry:** If information transfer (e.g., transfer entropy) is symmetric or backward-dominant across layers in steady state, the cascade claim fails.

These criteria act as **guardrails**, they delimit where ECI is valid and where classical models may suffice.

2.9 Worked micro-example (thought experiment)

Suppose we study manufacturing firms within a single country and stable policy period. Let:

- L = log employee count tier;
- T = median **order-to-cash** cycle per tier.

We fit $\log T = \alpha \log L + c$ using Theil–Sen within each year. Findings:

- 2014–2018: $\hat{\alpha} \in [1.8, 2.2]$ and clean collapses → **coherent, buffered regime**.
- 2019Q4–2020Q2: $\hat{\alpha}$ drops to 1.2 with poor collapse → **decoherence event** (shock transmission), consistent with supply-chain strain.
- 2021–2022: $\hat{\alpha}$ partially rebounds to 1.6 as re-shoring and inventory buffers increase.

Even without magnitudes of GDP or inflation, the **slope** narrates structure: whether scale differentials in timing are present (buffered) or flattened (exposed).

2.10 Implementation notes (for Section 5 reuse)

- **Binning.** Prefer small multiples of environment-fixed bins (country $\times$ sector $\times$ regime $\times$ quarter). Use changepoint detection to keep regimes stable inside bins.
- **Uncertainty.** Bootstrap firms/edges/shipments; report percentile CIs on α . Track drift in coverage (data availability) as a QA metric.
- **Placebos.** Rescale clocks (e.g., convert days $\leftrightarrow$ weeks) to verify slope invariance. Shuffle L within bins to estimate the bias you’d get by chance.
- **Ledger.** Maintain a “slope vs. intercept ledger”: every estimate of α must be accompanied by the intercept c and a note of known level shifts (policy, unit, inflation rebases). This documents that robustness genuinely lives in slopes.

Takeaways for economists

1. RTM provides a **single structural parameter**, the slope α , to summarize how timing stretches with scale.
2. Estimating α **bin-by-bin** and testing **collapse** shields you from many confounders that plague level-based indicators.
3. α_{econ} is **complementary**, not substitutive: it adds a coherence lens to volatility, liquidity, and leverage metrics.
4. RTM is **falsifiable** in this domain; clear failure modes prevent overreach.

3. Defining α_{econ} and Constructing the Economic Coherence Index (ECI)

This chapter formalizes the **Economic Coherence Exponent** α_{econ} and specifies the **Economic Coherence Index (ECI)**, a rolling, uncertainty-aware estimate of coherence derived from multiple (L, T) proxy families. We present (i) measurement definitions, (ii) a slope-first estimation pipeline with errors-in-variables, (iii) a collapse test to validate binwise scaling, (iv) a

random-effects meta-estimator that fuses proxies, and (v) real-time nowcasting and quality assurance.

3.1 Objects and environments

Let $\mathcal{U}$ be a **fixed environment** (e.g., country $\times$ policy regime $\times$ sector $\times$ quarter). Inside $\mathcal{U}$, we observe N units $u = 1, \dots, N$ (firms, edges, products, ports, tickers...) and build **paired measurements**

$$(L_u, T_u) \quad \text{with } L_u > 0, T_u > 0.$$

- L is a **scale proxy** (size, path length, capitalization tier, network degree, geographic scope).
- T is a **characteristic time** of a *compatible process layer* (relaxation, renewal, recovery, persistence).

RTM scaling within $\mathcal{U}$:

$$T_u = \kappa_u L_u^{\alpha_u} \varepsilon_u, \mathbb{E}[\log \varepsilon_u] = 0.$$

Taking logs:

$$y_u = \log T_u = \alpha_u x_u + c_u + \eta_u, \quad x_u = \log L_u, \quad c_u = \log \kappa_u.$$

We allow **measurement error** in both x and y :

$$x_u^{\text{obs}} = x_u + \xi_u, \quad y_u^{\text{obs}} = y_u + \zeta_u,$$

with ξ_u, ζ_u mean-zero, possibly heteroskedastic.

Goal. Estimate α_u robustly (slope-first), validate that a single α explains the bin via **collapse**, and combine across independent proxy families to obtain $\hat{\alpha}_{\text{econ}}(\mathcal{U})$.

3.2 Proxy families for L and T

We recommend using **at least two** independent families per environment; examples:

A. Market microstructure

- L : capitalization tier; median trade size tier; degree in a cross-impact network.
- T : microprice reversion half-life; order-book resiliency time (depth recovery); quote-stability persistence.

B. Logistics & supply chains

- L : path length (stages), multi-modal route size, port capacity tier.
- T : shipment lead-time persistence; dwell-time decay; inventory replenishment half-life.

C. Credit & financing

- L : maturity ladder span; interbank network degree; book size tier.
- T : roll-over renewal time; spread mean-reversion half-life after funding shocks.

D. Information flow

- L : audience/reach tier; outlet network centrality; jurisdictional scope.
- T : sentiment/news shock decay; disagreement dispersion half-life.

Compatibility rule. Within a family, L and T must describe the **same process layer**; do not mix micro L with macro T inside one regression.

3.3 Binwise slope estimation (EIV / robust)

Given $\mathcal{U}$ and a proxy family f , fit

$$y_u^{\text{obs}} = \alpha_{u,f} x_u^{\text{obs}} + c_{u,f} + \epsilon_u$$

with **errors-in-variables** to correct attenuation:

- **Orthogonal Distance Regression (ODR)** or **Total Least Squares (TLS)** when $\text{Var}(\xi) \approx \text{Var}(\zeta)$.
- **SIMEX** (simulation-extrapolation) if we can approximate σ_ξ^2 from repeated measures or instrument precision.
- **Theil-Sen** median slope (robust to outliers) as a sensitivity check.
- **Bootstrap** (clustered by unit/entity) for CIs and bias correction.

Deliverables per bin & family: $\hat{\alpha}_{u,f}$, 95% CI, intercept $\hat{c}_{u,f}$, fit diagnostics, and coverage diagnostics (how representative the sample is inside $\mathcal{U}$).

3.4 Collapse validation (binwise)

After estimating $\hat{\alpha}$, we test whether a **single scaling** explains the bin:

1. **Rescale** the data: $\tilde{y}_u = y_u^{\text{obs}} - \hat{\alpha} x_u^{\text{obs}}$.
2. **Null expectation:** within $\mathcal{U}$, $\tilde{y}_u \approx c_u + \text{noise}$ **independent of** x .
3. **Quantify collapse** with an ANOVA-style statistic:

$$\Delta_{\text{collapse}} = R^2 (\tilde{y} \sim x^{\text{obs}}).$$

We **pass** the collapse if Δ_{collapse} is below a small threshold (e.g., < 0.05) *and* residual diagnostics show no systematic trend vs. x . Failure indicates mixed regimes or a non-power relationship in that bin.

3.5 Multi-proxy fusion: random-effects meta-estimate

When at least two families pass **collapse** and QA, we fuse their binwise slopes $\{\hat{\alpha}_f\}_{f=1}^F$ into an **Economic Coherence Index** estimate for that window. We use a **random-effects** meta-analytic model that acknowledges between-family heterogeneity.

Estimator. Let $\hat{\sigma}_f^2$ be the (cluster/bootstrap) variance of $\hat{\alpha}_f$. Estimate the between-family variance $\hat{\tau}^2$ by **REML** (preferred; DerSimonian-Laird reported as sensitivity). Define random-effects weights

$$w_f = \frac{1}{\hat{\sigma}_f^2 + \hat{\tau}^2},$$

and compute the fused slope and its variance as

$$\hat{\alpha}_{\text{econ}} = \frac{\sum_{f=1}^F w_f \hat{\alpha}_f}{\sum_{f=1}^F w_f}, \quad \text{Var}(\hat{\alpha}_{\text{econ}}) = \frac{1}{\sum_{f=1}^F w_f}.$$

Report 50/95% intervals from the normal approximation (or bootstrap the fusion for robustness).

Heterogeneity diagnostics. We publish both the fixed-effect summary and the heterogeneity statistics:

- **Cochran's Q** (using *fixed-effect* weights $w_f^{FE} = 1/\hat{\sigma}_f^2$):

$$\hat{\alpha}_{FE} = \frac{\sum_f w_f^{FE} \hat{\alpha}_f}{\sum_f w_f^{FE}}, \quad Q = \sum_{f=1}^F w_f^{FE} (\hat{\alpha}_f - \hat{\alpha}_{FE})^2.$$

Under homogeneity, $Q \sim \chi_{F-1}^2$ approximately.

- I^2 (share of total variation due to heterogeneity):

$$I^2 = \max\{0, \frac{Q - (F - 1)}{Q}\} \times 100\%.$$

Gates and thresholds (pre-registered).

- Proceed with a single fused number only if:
 - at least **2 families** pass QA and collapse,
 - $I^2 < 50\%$ (*moderate or lower heterogeneity*), and
 - REML converges with $\hat{\tau}^2$ finite and $\hat{\tau}^2$ below a historical cap (e.g., **$\leq 90\text{th percentile}$** of past clean windows).
- If $I^2 \geq 50\%$ or the Q test rejects homogeneity at $p < 0.05$, we **do not publish a single ECI**. Instead we:
 - report the **family-wise** $\hat{\alpha}_f$ with CIs,
 - include **leave-one-family-out** influence diagnostics, and
 - annotate FAMILY_DIVERGENCE in QA.

Sensitivity panel. Alongside REML we report:

- **DL** estimate of τ^2 ,
- the fixed-effect summary $\hat{\alpha}_{FE}$,
- and a forest plot (per family $\hat{\alpha}_f$, weight w_f , CI), plus $Q, I^2, \hat{\tau}^2$.

Rationale. Random-effects down-weight families with large internal uncertainty ($\hat{\sigma}_f^2$) **and** windows where families disagree (large $\hat{\tau}^2$). The I^2 gate prevents a misleading single number when proxies tell materially different stories.

3.6 From bins to a real-time index: ECI(t)

To produce a time series, roll $\mathcal{U}$ across overlapping windows (e.g., monthly with 1-week stride; quarterly with 1-month stride).

Algorithm (high-level).

1. Define rolling environments $\mathcal{U}_t$ by time window and regime filters (changepoint detection to keep regimes stable within windows).
2. For each $\mathcal{U}_t$ and family f , compute $\hat{\alpha}_{\mathcal{U}_t, f}$ + collapse test.
3. Combine families via random effects $\rightarrow \hat{\alpha}_{\text{econ}}(t)$.
4. Apply **QA gates**: minimum sample size, proxy coverage, collapse pass-rate threshold (e.g., ≥ 2 families pass).
5. Smooth with a **causal filter** (e.g., EWMA with half-life 2–3 windows) to stabilize noise while preserving turning points.
6. Publish **ECI(t)** as $\hat{\alpha}_{\text{econ}}(t)$ with an uncertainty band and QA flags.

QA flags (examples).

LOW_COVERAGE, FAMILY_DIVERGENCE (high heterogeneity), NO_COLLAPSE, REGIME_MIX (changepoint inside window), CLOCK_SHIFT (unit rebasing detected).

3.7 Decoherence events and leading signals

Define **decoherence events** as large, significant downward moves:

$$\Delta \alpha_{\text{econ}}^-(t) = \hat{\alpha}_{\text{econ}}(t) - \hat{\alpha}_{\text{econ}}(t - h) \leq -\theta,$$

with horizon h (e.g., 3 months) and threshold θ chosen by pre-registered percentile (e.g., 10th percentile of historical changes) or by a multiple of the rolling standard error. Tag events only when QA flags are green (no regime mix; ≥ 2 families pass collapse). These events serve as **candidate early-warnings** for H2 (anticipation).

3.8 Reporting standards (bin and index level)

Per bin ($\mathcal{U}$, family f):

- $\hat{\alpha}_{\mathcal{U}, f}$, 95% CI; $\hat{\sigma}_{\mathcal{U}, f}$.
- Collapse statistic Δ_{collapse} and pass/fail.
- Sample size, coverage, leverage points, bootstrapping scheme.
- Known level shifts (units, policy rebase).

Per time t :

- $\hat{\alpha}_{\text{econ}}(t)$, 50/95% bands; heterogeneity $\hat{\tau}^2(t)$.

- Family contributions and leave-one-out influence.
- QA flags and notes on regime stability.

3.9 Robustness and ablations

- **Errors-in-variables sensitivity.** Compare ODR/TLS, Theil–Sen, and SIMEX-corrected slopes.
- **Alternative L, T choices.** Swap proxies within families (e.g., degree vs. path length) and check stability.
- **Placebo clocks.** Change time units (days $\leftrightarrow$ weeks) to verify slope invariance.
- **Shuffle tests.** Randomly permute L within bins to estimate chance slopes; reported as a null benchmark.
- **Subsample stability.** Jackknife entities/sectors/regions.
- **Non-power alternative.** Fit $\log T = g(\log L)$ with splines; a strong, consistent curvature across bins falsifies the power-law assumption for that domain.

3.10 Minimal pseudocode

for t in rolling_windows:

U_t = define_environment(t) # regime-stable window

family_estimates = []

for f in proxy_families:

data = load_pairs(U_t, f) # (L,T) with metadata

xobs, yobs = log(L), log(T)

alpha_hat, c_hat, se = EIV_fit(xobs, yobs) # ODR/SIMEX/Theil–Sen

collapse = R2_of_residual_vs_x(yobs - alpha_hat*xobs, xobs)

if collapse < threshold and coverage_ok(data):

family_estimates.append((alpha_hat, se))

if len(family_estimates) >= 2:

alpha_RE, se_RE, tau2 = random_effects(family_estimates)

if QA_ok(family_estimates, tau2):

ECI[t] = (alpha_RE, se_RE, flags=None)

else:

ECI[t] = (alpha_RE, se_RE, flags=QA_flags)

else:

ECI[t] = (nan, nan, flags={'LOW_COVERAGE'})

3.11 Interpretation guide (practical)

- **High ECI (larger α_{econ})**
Expect deeper staging and slower cross-scale propagation: improved shock absorption, potentially lower raw throughput at the largest scales; often preferable during fragile periods.
- **Low ECI (smaller α_{econ})**
Flatter time gradients: faster propagation, efficient in calm times but exposes the system to synchronized failures.
- **Rising ECI** may indicate post-shock re-organization (buffers rebuilding, governance improving).
- **Falling ECI**, especially with clean QA, warrants vigilance: decoherence that can precede stress episodes.

3.12 Limitations specific to ECI

- **Proxy fragility.** Some (L, T) pairs are cyclical or policy-sensitive; rotating families and documenting coverage is essential.
- **Regime detection.** Mis-specified environments mix clocks and bias slopes; changepoint detection is not foolproof.
- **Endogeneity.** Coherence can be the result of policy actions rather than an exogenous cause; causal interpretation requires additional designs (instruments, differences-in-differences).
- **Data latency.** Some T proxies update slowly; the index must disclose latency and use nowcasting judiciously.

4. Falsifiable Hypotheses & Study Design

This chapter turns the construct into tests that can *pass or fail*. We (i) state hypotheses, (ii) pre-register identification choices, (iii) define outcome metrics and evaluation windows, (iv) specify statistical models, (v) detail validation logic (collapse tests, QA), and (vi) enumerate failure modes that would falsify the approach.

4.1 Hypotheses

H1 — Resilience (cross-sectional and time-varying).

Within comparable environments, higher baseline coherence $\bar{\alpha}_{\text{econ}}$ is associated with (a) smaller peak-to-trough drawdowns during shocks and (b) faster recoveries (shorter half-lives back to trend).

H2 — Anticipation (leading indicator).

Large, QA-clean negative moves in coherence, **decoherence events** $\Delta\alpha_{\text{econ}}^-(t)$, predict subsequent macro-financial stress (recessions, crisis indicators, liquidity squeezes) with horizons of 6–18 months, out-of-sample.

H3 — Cascade signature (multilayer).

Across aggregation layers (micro $\rightarrow$ meso $\rightarrow$ macro) inside a fixed regime, (a) α is **non-decreasing** with the layer index, and (b) directed information flow is forward-biased (micro $\rightarrow$ meso $\rightarrow$ macro) as assessed by transfer entropy/Granger causality.

4.2 Pre-registration: estimands, windows, and QA gates

We will pre-register the following before any outcome peeking:

- **Estimands.**
 - $\alpha_{u,f}$: binwise slopes per proxy family f .
 - $\alpha_{\text{econ}}(t)$: random-effects fusion (Section 3) with uncertainty.
 - **Decoherence event:** $\Delta\alpha_{\text{econ}}(t) \leq -\theta_h$ over horizon $h \in \{1,3,6\}$ months, with θ_h set to the historical 10th percentile of changes or k times the rolling standard error (pre-registered k).
- **Windows & sampling.**
 - Rolling quarterly windows (primary), monthly stride; semi-annual windows (sensitivity).
 - Country-sector panels for cross-section; market/credit/logistics/information families for multi-proxy redundancy.
- **QA gates.**
 - Minimum of **two** proxy families passing **collapse** within a bin.
 - Coverage thresholds (min N per bin, min share of panel present).
 - Regime stability within a window (changepoint tests).
 - Clock invariance checks (placebo rescaling).

Observations failing QA gates will be flagged and excluded from hypothesis tests (kept for descriptive plots only).

4.3 Outcomes and ground-truth labels

- **Shock drawdown:** percentage peak-to-trough decline in target variable Y (e.g., industrial production, real sales, market index) during a shock window identified by external chronology or rule-based thresholds.
- **Recovery half-life:** time to recover 50% of the drawdown (or to return within a band of pre-shock trend).
- **Stress labels** (for H2): binary weekly/monthly markers for recessions (official business cycle dating), financial stress indices, liquidity crises, or rule-based market stress (e.g., top decile drawdown or spread blowout).
- **Layer definitions** (for H3):
 - **Micro:** firm/port/ticker-level;
 - **Meso:** sector/route/cluster aggregates;
 - **Macro:** country/market system aggregates.

4.4 Identification: conditioning sets and controls

To reduce confounding:

- **Fixed effects:** environment FE (country × sector × regime), time FE (calendar quarter), and where appropriate entity FE (firm/port/ticker).
- **Controls** (do **not** collide with time scaling): volatility level, leverage proxies, liquidity depth, credit spreads, and global factors (commodity indices, policy-rate changes). These enter as *covariates*, while α remains the **slope** estimand derived upstream; controls are not allowed to alter slope construction.

4.5 Statistical tests

H1 — Resilience

(a) Drawdown size (cross-section, panel).

$$\text{Drawdown}_{i,s} = \beta_0 + \beta_1 \bar{\alpha}_{\text{econ},i,s}^{(\text{pre})} + \gamma' X_{i,s} + \text{FE} + \varepsilon_{i,s}$$

where i indexes country (or sector), s the shock episode, X controls, and $\bar{\alpha}^{(\text{pre})}$ is the average ECI in the pre-shock baseline. **Prediction:** $\beta_1 < 0$. Cluster SEs by i and s .

(b) Recovery half-life (AFT/parametric survival).

Accelerated failure-time model:

$$\log(\text{HalfLife}_{i,s}) = \delta_0 + \delta_1 \bar{\alpha}_{\text{econ},i,s}^{(\text{pre})} + \phi' X_{i,s} + \text{FE} + \eta_{i,s}.$$

Prediction: $\delta_1 < 0$. Robustness: Cox model with ECI as covariate and shared frailty.

H2 — Anticipation

Event-study and classification.

1. Binary prediction.

Logit/Probit:

$$\Pr(\text{Stress}_{t+h} = 1) = \sigma(\theta_0 + \theta_1 \Delta^- \alpha_{\text{econ}}(t) + \psi' Z_t + \text{FE}),$$

with horizons $h \in \{6, 12, 18\}$ months and Z_t standard leading indicators. **Prediction:** $\theta_1 > 0$.

2. Scoring rules.

Out-of-sample backtests with rolling origin; evaluate **AUC**, **Brier score**, **PR-AUC**; benchmark vs. canonical indicators (volatility, term spread, credit spreads). Require statistically significant improvement (DeLong test for AUC; Diebold–Mariano for scores), controlling for multiple horizons.

3. Change-point alignment.

Kaplan–Meier plots of time-to-stress after decoherence events vs. after matched placebo windows; log-rank tests.

H3 — Cascade signature

(a) Monotonicity across layers.

Within regime-stable windows, compute $\hat{\alpha}_\ell$ for $\ell \in \{\text{micro}, \text{meso}, \text{macro}\}$ using layer-compatible (L, T) pairs. Test:

$$H_0: \alpha_{\text{micro}} \geq \alpha_{\text{meso}} \text{ or } \alpha_{\text{meso}} \geq \alpha_{\text{macro}} \text{ vs } H_A: \alpha_{\text{micro}} \leq \alpha_{\text{meso}} \leq \alpha_{\text{macro}},$$

using paired comparisons with bootstrap CIs and multiple-testing control across windows.

(b) Directionality (TE/Granger).

Compute **transfer entropy** $TE_{\ell \rightarrow \ell'}$ and **Granger causality** tests between ECI-compatible layer signals (e.g., meso activity vs. macro aggregates). **Prediction:** $TE_{\text{micro} \rightarrow \text{meso}} > TE_{\text{meso} \rightarrow \text{micro}}$ and similarly meso $\rightarrow$ macro. Pre-register embedding dimensions, lag orders, and surrogate tests for significance.

4.6 Validation and falsification logic

A hypothesis is **counted as passed** only if:

- Binwise **collapse tests** pass for the contributing proxy families;
- QA flags are clear;
- Effect signs match predictions with pre-registered significance levels;
- Out-of-sample performance exceeds baselines by pre-registered margins.

The approach is **falsified** for a domain if, repeatedly across regimes and datasets:

- **No slope separation** is detected (α indistinguishable from 0 or unstable) in well-formed bins;
- **No collapse** occurs after rescaling;
- **Reverse cascade** (α decreases with aggregation) is systematic;
- **Directionality symmetry** (no forward bias) persists after robustness checks;
- H2's predictive content vanishes relative to strong baselines under proper oos evaluation.

Negative results will be documented and published as scope boundaries.

4.7 Multiple comparisons, uncertainty, and robustness

- **Multiplicity.** Control FDR across horizons and layers (Benjamini–Hochberg).
- **Uncertainty propagation.** Carry $SE(\hat{\alpha})$ from binwise fits into downstream models via parametric bootstrap or Bayesian hierarchical variants.
- **Alternative constructions.** Replace ODR with Theil–Sen/SIMEX; swap L, T proxies; vary windows; test non-power alternatives (spline $g(\log L)$).
- **Placebos.** Clock rescalings, shuffled L within bins, pseudo-events matched on controls.
- **Stability.** Rolling leave-one-family-out ECI; heterogeneity thresholds ($\hat{\tau}^2$) for acceptance.

4.8 Data governance and reproducibility

- **Pre-registration** of hypotheses, windows, thresholds, and model classes.
- **Versioned artifacts:** raw inputs (where licensable), feature builders, bin definitions, α -estimates, ECI series, QA flags, and analysis notebooks.
- **Audits:** independent recomputation of α on a held-out subset by a separate team; red-team protocols to probe for leakage or circularity (e.g., using outcome variables in proxy construction).

4.9 Power considerations (back-of-envelope)

Given typical panel sizes (hundreds to thousands per bin) and moderate measurement noise, binwise α estimation via EIV yields SEs in the 0.05–0.15 range. Detecting $\Delta\alpha$ shifts of 0.2–0.3 (a practically meaningful change in coherence) at 5% significance with >80% power is feasible with monthly/quarterly windows over multi-year spans, provided at least two proxy families pass collapse.

4.10 What success and failure mean

- **Success:** α_{econ} adds distinct, robust information about the *structure of timing across scale*, improving resilience inference and early warnings beyond volatility/liquidity/leverage.
- **Failure:** α behaves as an unstable artifact of clocks, units, or regimes; collapse seldom passes; predictive value is nil out-of-sample. In such cases, the RTM lens is *not* informative for that economic domain, and we recommend sticking to classical tools.

5. Data & Methods

This section specifies how we turn raw economic streams into binwise slopes α and a real-time Economic Coherence Index **ECI(t)**. We detail (i) datasets, (ii) feature construction for (L, T) pairs, (iii) environment/bins and regime control, (iv) estimation algorithms (EIV/TLS/SIMEX/robust), (v) collapse testing, (vi) multi-proxy fusion, (vii) nowcasting with QA/latency handling, and (viii) reproducibility.

5.1 Datasets (families & cadence)

We organize inputs into **four proxy families**. Each family is optional on any given window; ECI requires ≥ 2 families passing QA.

A. Market microstructure (intra-day to daily)

Order book, trades/quotes (L1–L3), consolidated tape, index constituents, corporate actions.

B. Logistics & supply chains (daily to monthly)

Port calls & dwell times, freight bookings, shipment lead times, inventory levels, routing metadata.

C. Credit & financing (daily to monthly)

Interbank rates/volumes, funding spreads, roll-over/renewal activity, maturity ladders.

D. Information flow (hourly to daily)

Newswire timestamps, article/link graphs, audience/reach tiers, text embeddings/sentiment, social signals.

Data governance. For each stream we maintain a data sheet: provenance, cadence, coverage (entities×time), revisions policy, and legal/licensing constraints. All time stamps are normalized to UTC; trading-calendar effects are tracked.

5.2 Feature construction: mapping to (L, T)

Each family yields paired measurements inside a **fixed environment** $\mathcal{U}$ (Sec. 5.3). We compute:

A. Market microstructure

- **Scale L** : capitalization tier (quantiles), median trade size tier, degree/centrality in cross-impact or correlation networks.
- **Time T** :
 - *Microprice reversion half-life*: fit ARMA/ECM to mid-quote deviations; report $t_{1/2}$.
 - *Order-book resiliency*: time to replenish depth after a standardized shock.
 - *Quote stability persistence*: expected time above a spread threshold.

B. Logistics & supply chains

- **Scale L** : path length (stages in bill of materials), route size (TEU bands), port capacity tier.
- **Time T** :
 - *Lead-time persistence*: decay constant of shipment delays post-shock.
 - *Dwell-time decay*: exponential tail fit for yard/anchorage durations.
 - *Inventory replenishment half-life*: time to return to target stock bands.

C. Credit & financing

- **Scale L** : maturity ladder span, network degree (interbank exposures), book size tier.
- **Time T** :
 - *Roll-over renewal*: median time to refinance a maturing bucket.
 - *Spread mean-reversion*: $t_{1/2}$ of funding spread shocks.
 - *Queueing persistence*: time for primary issuance backlog to clear.

D. Information flow

- **Scale L** : outlet audience tier, graph centrality, jurisdiction scope.
- **Time T** :
 - *Sentiment decay*: relaxation time of topic polarity after a shock.
 - *Disagreement dispersion*: half-life of cross-source variance.
 - *Attention half-life*: time for article impressions to drop to 50%.

Measurement notes.

1. We compute L as monotone tiers or log-scale magnitudes; T is always a **characteristic time** (half-life, decay constant, return-to-target, resiliency).
2. Each (L, T) pair carries metadata: timestamp, entity ID, method recipe, SEs for T , and quality flags (fit R^2 , residual diagnostics).

5.3 Environments & binning (regime stability)

A **bin** $\mathcal{U}$ is defined by: $(country \mid currency) \times sector \text{ (or market)} \times policy \text{ regime} \times time \text{ window}$.

- **Time windows.** Primary: rolling **quarterly** windows with **monthly stride**; sensitivity: semi-annual windows.
- **Regime stability.** We apply univariate and multivariate **change point detection** (e.g., PELT, Bai–Perron) to guard that macro policy, reporting standards, or market microstructure *do not* shift inside $\mathcal{U}$. If they do, the window is split and flagged REGIME_MIX.
- **Coverage thresholds.** Min entities per family per bin (e.g., ≥ 50 for microstructure; ≥ 20 routes/ports; ≥ 10 banks) and min time stamps per entity to estimate T .

5.4 Estimation of binwise slopes α (errors-in-variables)

For each $\mathcal{U}$ and family f , we fit:

$$\log T_u^{\text{obs}} = \alpha_{u,f} \log L_u^{\text{obs}} + c_{u,f} + \epsilon_u,$$

allowing noise in both axes.

Estimators.

- **ODR/TLS (default):** minimizes orthogonal residuals; good when errors are comparable.
- **SIMEX (attenuation correction):** if we can estimate $\text{Var}(\xi)$ of $\log L$ from replicates/instruments.
- **Theil–Sen (robust check):** median of pairwise slopes; resists outliers/heavy tails.

Uncertainty. Block/bootstrap by entity (clustered resampling); 95% percentile CIs. We report **intercepts** $c_{u,f}$ in a “slope–intercept ledger” to document level shifts that should *not* affect α .

5.5 Collapse test (scaling validation)

After obtaining $\hat{\alpha}_{u,f}$, compute residualized outcomes:

$$\tilde{y}_u = \log T_u^{\text{obs}} - \hat{\alpha}_{u,f} \log L_u^{\text{obs}}.$$

We test for **independence** of $\tilde{y}$ from $\log L$ inside $\mathcal{U}$:

- **Statistic:** $\Delta_{\text{collapse}} = R^2(\tilde{y} \sim \log L)$.
- **Pass rule:** $\Delta_{\text{collapse}} < 0.05$ and no visible trend in residual plots (nonparametric smooth $<$ pre-registered bandwidth).
- **Fail actions:** mark family-bin as NO_COLLAPSE; exclude from fusion; note in QA.

5.6 Multi-proxy fusion (random effects)

Given family estimates $\{\hat{\alpha}_{u,f}\}$ that pass collapse:

$$\hat{\alpha}_{\text{econ}}(\mathcal{U}) = \frac{\sum_f w_f \hat{\alpha}_{u,f}}{\sum_f w_f}, \quad w_f = \frac{1}{\hat{\sigma}_{u,f}^2 + \hat{\tau}^2},$$

with $\hat{\sigma}_{u,f}^2$ the bootstrapped variance and $\hat{\tau}^2$ the between-family heterogeneity (REML default; DerSimonian–Laird as sensitivity). We publish $\hat{\alpha}_{\text{econ}}$, 50/95% bands, **Q-statistic**, $\hat{\tau}^2$, and a **leave-one-family-out** influence analysis.

5.7 From bins to ECI(t): nowcasting pipeline

Rolling construction.

1. **Define windows** $\mathcal{U}_t$ (quarterly, monthly stride), run regime checks.
2. **Per family**, compute $\hat{\alpha}_{u_{t,f}}$, CIs, collapse tests, coverage stats.
3. **Fuse** into $\hat{\alpha}_{\text{econ}}(t)$ with random effects.
4. **QA gates**: require ≥ 2 families passing collapse; cap heterogeneity ($\hat{t}^2$ below a pre-registered threshold); enforce coverage minima.
5. **Smoothing**: apply a **causal EWMA** (half-life 2–3 windows) to stabilize noise; smoothing is *never* used in hypothesis tests, only for publishing the headline ECI series.
6. **Flags**: attach LOW_COVERAGE, FAMILY_DIVERGENCE, NO_COLLAPSE, REGIME_MIX, CLOCK_SHIFT as relevant.

Latency handling. Each observation carries **as-of date** and **vintage**. We maintain a real-time archive and recompute ECI(t) on vintages to evaluate revision sensitivity (reliability plots).

5.8 Decoherence events (signal definition)

We define a **decoherence event** when all hold simultaneously:

- $\hat{\alpha}_{\text{econ}}(t) - \hat{\alpha}_{\text{econ}}(t - h) \leq -\theta_h$, with $h \in \{1, 3, 6\}$ months and θ_h pre-registered (percentile or $k \cdot SE$).
- QA gates pass at t and in the lookback window; no REGIME_MIX.
- The drop is **confirmable** by at least **two** families individually (sign-consistent, even if magnitudes differ).

Events are time-stamped and later aligned to stress outcomes in H2.

5.9 Controls and placebos

- **Clock placebos.** Convert time units (days $\leftrightarrow$ weeks $\leftrightarrow$ months) within a bin; slopes should remain invariant while intercepts change.
- **Shuffle placebos.** Permute L labels within bins to estimate null slope distributions (attenuation baselines).
- **Alternative forms.** Fit $\log T = g(\log L)$ with cubic splines; systematic curvature across bins falsifies the power-law form for that domain.

5.10 Robustness suite

- **Estimator swaps.** ODR $\leftrightarrow$ Theil–Sen $\leftrightarrow$ SIMEX; compare $\hat{\alpha}$ deltas and CI overlap.
- **Proxy swaps.** Replace degree with path length, median trade size with market-cap tier, etc.; recompute ECI.
- **Window sensitivity.** Semi-annual windows; alternative strides; overlapping vs. disjoint.
- **Coverage stress.** Downsample entities; check degradation of CI widths and collapse rates.

- **Heterogeneity thresholds.** Vary acceptable $\hat{\tau}^2$ and re-flag QA.

5.11 Software, computation, and artifacts

- **Stack.** Python/R for data engineering; ODR (scipy), SIMEX (custom or R package), Theil-Sen (statsmodels), changepoints (ruptures/strucchange), meta-analysis (metafor/py-meta).
- **Pipelines.** Reproducible DAGs (e.g., make, dvc, prefect) with deterministic seeds.
- **Artifacts.** Versioned parquet/feather tables for: features, bin definitions, $\hat{\alpha}_{u,f}$, collapse stats, fusion outputs, ECI(t) with QA flags, and all figures.
- **Documentation.** YAML spec for each proxy: formulae, filters, unit conventions, missing-data policy.
- **Testing.** CI checks for (i) invariance to unit changes (clock placebos), (ii) reproducibility of $\hat{\alpha}$ to 1e-6 tolerance on a fixed sample, (iii) collapse statistics bounds.

5.12 Privacy & ethics

- **Aggregation.** Publish only bin- and index-level outputs; suppress micro identifiers unless explicitly consented and anonymized.
- **Bias.** Monitor family participation by country/sector to avoid ECI reflecting data richness rather than structure; include a LOW_COVERAGE flag and abstain from inference when flagged.
- **Open science.** Pre-register hypotheses and thresholds; release code and synthetic replicas where licensing prevents raw-data sharing.

5.13 Summary

We transform heterogeneous economic streams into coherent (L, T) pairs, estimate **binwise slopes** with measurement-error correction, **validate scaling** via collapse tests, **fuse** multi-proxy estimates under random effects, and publish a QA-aware **ECI(t)** with uncertainty and latency tracking. The pipeline is expressly designed to be **falsifiable** (clear failure modes), **auditable** (versioned artifacts), and **complementary** to classical volatility/liquidity/leverage indicators.

6. Results — Retrospective Backtests

This section shows how the Economic Coherence Index (ECI) behaves on historical data. Because this is a methods paper, we emphasize transparent templates, pass/fail criteria, uncertainty, and negative findings. Concrete numbers are placeholders illustrating the reporting format; the preregistered replication package will replace them with actual estimates.

6.1 Setup and evaluation protocol

Windows. Rolling quarterly bins (monthly stride) for each country $\times$ sector $\times$ regime.

Families. At least two of: *Market microstructure, Logistics, Credit, Information*.

Acceptance gates. In each bin, families must pass **collapse** and coverage thresholds; heterogeneity ($\hat{\tau}^2$) within limits.

Uncertainty. 50/95% bands for $\hat{\alpha}_{\text{econ}}(t)$; cluster/bootstrap CIs flow into downstream tests.

Benchmarks. Volatility indices, term spread, credit spreads, composite financial stress indices (FSI).

Pre-registered metrics. AUC/PR-AUC, Brier score, Cox/AFT coefficients, Diebold–Mariano tests vs. benchmarks.

Figure 1 (template). $ECI(t)$ with 50/95% bands and QA flags, alongside benchmarks (scaled).

Table 1 (template). Collapse pass-rates, per family, per regime.

6.2 H1 — Resilience (drawdowns and recoveries)

6.2.1 Cross-sectional drawdowns

We regress peak-to-trough drawdowns during each shock episode on **pre-shock** $\bar{\alpha}_{econ}$ with controls and fixed effects.

Table 2 (template). Drawdown regressions

- β_1 on $\bar{\alpha}_{econ}$ (expect **negative**)
- Controls: volatility, leverage, liquidity; FE: country×episode
- Cluster SEs; R^2 ; N bins

Illustrative format:

$\beta_1 = -0.28[-0.41, -0.15]$, $p < 0.001$. Interpretation: +0.5 increase in baseline coherence associates with ~14% smaller drawdowns, ceteris paribus.

6.2.2 Recovery half-lives

Accelerated failure-time model (AFT) for time-to-50% recovery.

Table 3 (template). AFT estimates

δ_1 on $\bar{\alpha}_{econ}$ (expect **negative**); frailty terms; concordance.

Illustrative format:

$\delta_1 = -0.35[-0.52, -0.18]$. Interpretation: higher coherence predicts faster recoveries (shorter half-lives), beyond volatility/liquidity effects.

Figure 2 (template). Kaplan–Meier curves stratified by tertiles of $\bar{\alpha}_{econ}$.

Robustness. Results stable under: (i) semi-annual windows, (ii) leave-one-family-out ECI, (iii) alternative outcome definitions (return-to-trend vs. to pre-shock band).

6.3 H2 — Anticipation (leading stress signal)

We define **decoherence events** as QA-clean drops in ECI exceeding pre-registered thresholds over $h \in \{1, 3, 6\}$ months.

6.3.1 Classification performance

Logit/Probit predicting stress at horizons $h = 6, 12, 18$ months.

Table 4 (template). Out-of-sample performance

- AUC, PR-AUC, Brier score for: (i) ECI-only, (ii) Benchmarks, (iii) ECI + Benchmarks
- DeLong and Diebold–Mariano tests; multiplicity control (FDR).

Illustrative format:

At $h = 12$ m, ECI-only AUC 0.72 (0.68–0.76); Benchmarks 0.66 (0.62–0.70); **Combined** 0.77 (0.73–0.80), ECI adds significant incremental value ($p < 0.01$).

6.3.2 Event alignment

Survival analysis from decoherence events to first stress signal.

Figure 3 (template). Time-to-stress curves for **ECI events** vs. matched placebo windows; log-rank p -values.

Lead-time distribution. Median lead 9–14 months (illustrative), with interquartile range reported per regime.

6.3.3 False positives & negatives

- **False positives** (ECI drops without ensuing stress): catalogued with QA notes (e.g., REGIME_MIX close to the window, or idiosyncratic sector shocks that reverse).
- **Misses** (stress without prior ECI drop): analyzed for **coverage** issues (too few families) or **proxy brittleness**.

6.4 H3 — Cascade signature (layer monotonicity and direction)

6.4.1 Layer monotonicity

Compute $\hat{\alpha}_{\text{micro}}$, $\hat{\alpha}_{\text{meso}}$, $\hat{\alpha}_{\text{macro}}$ inside regime-stable windows. Test non-decreasing order with bootstrap CIs.

Table 5 (template). Layer-wise α with pairwise differences

- Share of windows where $\alpha_{\text{micro}} \leq \alpha_{\text{meso}} \leq \alpha_{\text{macro}}$ (expect high).
- Violations documented by regime.

6.4.2 Directionality tests

Transfer entropy (TE) and Granger causality between layer aggregates compatible with ECI construction.

Figure 4 (template). TE arrows (micro→meso→macro) with confidence bands; surrogate tests for significance.

Interpretation template: Forward bias holds in X% of windows (FDR-controlled), consistent with cascade; exceptions coincide with REGIME_MIX or structural breaks.

6.5 Robustness, ablations, and diagnostics

6.5.1 Estimator swaps

ECI recomputed with Theil–Sen and SIMEX corrections.

Table 6 (template). $\Delta\hat{\alpha}$ vs. ODR; CI overlap rates; heterogeneity changes ($\hat{\tau}^2$).

Result format: median absolute change ≤ 0.06 ; no material effect on H1–H3 decisions.

6.5.2 Proxy swaps

Within families, alternative L, T proxies (e.g., degree↔path length, market-cap↔trade-size).

Figure 5 (template). Spider chart of family contributions to ECI under proxy swaps; stability lanes.

6.5.3 Window/coverage sensitivity

- Semi-annual vs. quarterly windows; different strides.
- Downsampling entities to stress coverage.
- QA flag rates (LOW_COVERAGE, FAMILY_DIVERGENCE, NO_COLLAPSE) reported.

Table 7 (template). QA flag incidence and impact on hypothesis tests.

6.5.4 Placebo & null diagnostics

- **Clock placebos:** unit rescalings change intercepts but preserve slopes (pass rates reported).
- **Shuffle nulls:** slope distributions under permuted L labels (should center near 0-attenuation baseline).
- **Non-power alternative:** spline $g(\log L)$ curvature tests, fraction of bins where power law is rejected.

6.6 Negative results and scope conditions

We document domains where ECI **fails** (by design):

- **No slope separation:** α unstable or indistinguishable from 0 despite good coverage → RTM uninformative at that layer (recorded).
- **No collapse:** persistent residual trends after rescaling → mixed regimes or wrong functional form (exclude).
- **Reverse cascade:** systematic $\alpha_{\text{macro}} < \alpha_{\text{micro}}$ in steady state → consider alternative architectures; RTM may not apply.
- **Directionality symmetry:** TE shows no forward bias after surrogates → cascade claim fails for that regime.

Table 8 (template). Negative result registry

- Domain, regime, reason, diagnostics, action (exclude / revise proxies / alternative model).

6.7 What the results imply (synthesis)

1. **Structure, not levels.** Where collapse passes, α captures an economy's **time-scale gradient** beyond volatility/liquidity.
2. **Resilience lens.** Higher pre-shock coherence aligns with shallower drawdowns and faster recoveries.
3. **Early warnings.** Decoherence events often **lead** stress by quarters, and add value beyond familiar benchmarks.

4. **Cascades are layered.** Forward-biased information flow and non-decreasing α across layers appear in stable regimes, precisely where policy design can influence buffers and transparency.

6.8 Replication checklist (what a reader should be able to redo)

- Recompute binwise $\hat{\alpha}$ for each accepted family with the posted code and data vintages.
- Verify collapse pass/fail and QA flags match.
- Recreate ECI(t), uncertainty bands, and decoherence event timestamps.
- Re-run H1–H3 tests with our seeds to reproduce tables/figures within tolerance.
- Swap estimators/proxies and see stability envelopes similar to ours.

7. Discussion

This section interprets what **ECI** measures, how it differs from familiar indicators, where it is most informative, when it should *not* be used, and how to read successes and failures. We also consider alternative explanations, causal identification limits, and implications for design and policy (expanded in Section 8).

7.1 What ECI actually measures

ECI is a structural slope: the gradient of characteristic **time** with respect to **scale** within a fixed environment. Where collapse passes, α_{econ} summarizes *how quickly timing stretches as you move up the size/aggregation ladder*. It is neither a volatility index nor a speedometer of the whole economy; it is a *geometry-of-tempo* statistic:

- **High α** $\rightarrow$ timing increases steeply with scale: large units are slower relative to small ones. This usually indicates *layering, buffers, and filtered information flow*, traits correlated with resilience but potentially lowering raw throughput at the largest scales.
- **Low α** $\rightarrow$ timing increases weakly with scale: propagation is fast across layers. This boosts short-run throughput but increases the chance of synchronized failures.

Because **clock/level shifts** live in the intercept, ECI is comparatively robust to rebasings, unit changes, and some regime-wide level shifts, *provided* the environment is correctly held fixed.

7.2 How ECI complements familiar signals

- **Volatility (e.g., VIX):** dispersion at a given scale; can be high in both coherent and incoherent regimes. ECI captures *cross-scale timing structure* that volatility cannot see.
- **Liquidity depth/spreads:** transactional frictions; they can improve as α rises (staged flow) or fall if buffering clogs execution. No fixed sign relation.
- **Leverage/credit spreads:** balance-sheet pressure; may co-move with ECI but conceptually distinct. A highly levered system can remain coherent, until leverage forces de-layering and α drops.
- **Business-cycle indicators (PMIs, unemployment):** level dynamics; ECI can lead or lag depending on whether coherence reorganizes *before* levels move.

Net: Treat ECI as a *third axis*, structure of time across scale, orthogonal to level and dispersion.

7.3 Mechanisms: why coherence tends to help resilience

Three generic channels explain the H1/H2 patterns we observe when bins pass collapse:

1. **Buffering & staging.** Larger units keep inventories, capital buffers, and decision checkpoints. As α rises, disturbances are dissipated at each stage, lengthening macro timing but reducing peak stress.
2. **Filtering of information.** Coherent systems slow rumor cascades and algorithmic reflex loops, reducing feedback overshoot.
3. **Heterogeneous clocks.** When layers run at differentiated tempos (large slow, small fast), cross-layer synchronization is harder; shocks struggle to lock all scales into the same phase.

These mechanisms can be engineered (governance, disclosure, circuit breakers, redundancy) and, crucially, *measured* with (L, T) proxies. They also clarify the trade-off: high α can reduce raw throughput or headline “speed,” which is sometimes misread as inefficiency.

7.4 Alternative explanations (and how we guard against them)

1. **Volatility regimes masquerading as coherence.**
If high volatility lengthens observed T uniformly, slopes could steepen mechanically. We mitigate by (i) estimating **within** environment-fixed bins, (ii) including volatility as a **control** in H1/H2 models, and (iii) requiring **collapse** to pass (uniform level shifts alone won’t pass).
2. **Measurement clocks and unit artifacts.**
We run **clock placebos** (days↔weeks↔months) to ensure slopes are invariant while intercepts move; we keep a **slope-intercept ledger**. Failures here invalidate bins.
3. **Endogeneity/selection.**
Coherence may be *chosen* in anticipation of shocks (reverse causality). We therefore (i) pre-register windows/thresholds, (ii) use out-of-sample evaluation for H2, and (iii) in extensions, leverage instruments or differences-in-differences where policy creates exogenous coherence shifts (e.g., disclosure mandates, circuit-breaker rules).
4. **Confounded layers.**
If L and T do not belong to the same process layer, spurious slopes arise. Our **compatibility rule** and bin-level **collapse** are designed to fail in that case, by design, an informative failure.

7.5 Scope conditions: where ECI is (and isn’t) useful

Works best when:

- Structure is quasi-stationary within bins (policy/microstructure stable).
- Multiple independent (L, T) families are available (≥ 2) with acceptable coverage.
- Timing processes are *internally generated* (renewal/relaxation) rather than entirely policy-clocked.

Should be avoided or flagged when:

- REGIME_MIX: fast-moving structural breaks inside windows.
- LOW_COVERAGE: too few entities per family; slopes unstable.

- **Single-family ECI:** no redundancy, report but abstain from inference.
- **Non-power form:** spline tests reveal consistent curvature (RTM form rejected).

7.6 Interpreting levels and changes across regimes

- **Cross-country comparison.** Compare only when *bin definitions match* (e.g., similar reporting standards and market microstructure). ECI is not a universal ranking; it is *contextual*.
- **Sector vs. market.** Sectors with engineered staging (utilities, pharmaceuticals) often exhibit higher α than hyper-competitive, just-in-time sectors. Policy that forces transparency and buffering can shift α up.
- **Trend breaks.** A *persistent* rise in ECI after a shock often reflects deliberate re-organization (inventory strategies, redundancy, governance). A transient spike with high heterogeneity may be noise or measurement artifacts.

7.7 Causality: what we can and cannot claim

ECI is observational and **structural-descriptive**. H1/H2/H3 provide *predictive* and *associational* evidence. To argue **causality**, we need:

- Exogenous or quasi-exogenous shocks to coherence (policy natural experiments).
- Instrumental variables that shift α but not outcomes except through α .
- Randomized pilot interventions (e.g., mandated disclosure cadence, circuit-breaker designs) with pre/post α measurement.

Until such designs are executed, we recommend phrasing causal claims cautiously (“associated with”, “predictive of”).

7.8 Model risk and overfitting

- **Proxy proliferation.** More proxies increase coverage but raise multiple-testing risk; we curb this by pre-registering families, using random-effects fusion, and publishing **negative results**.
- **EIV mis-specification.** If L - errors are misestimated, SIMEX corrections can bias slopes; we therefore publish **estimator-swap** results (ODR vs. Theil–Sen vs. SIMEX).
- **Look-ahead bias.** All ECI(t) series are computed by **vintage**, and hypothesis tests use only information available as of that date.

7.9 Relation to the broader RTM corpus

Economics inherits the same **slope-first** discipline seen in physical and biological RTM domains: **binwise scaling**, **collapse validation**, and **cascade signatures**. Conceptually, α_{econ} plays the role of a *coherence exponent* akin to chemistry’s environment-controlled kinetics or meteorology’s persistence gradients. Failures (no slope, no collapse) are not bugs; they are **scope boundaries**, signals that, in that domain or regime, RTM’s simple scaling does not describe timing.

7.10 Practical reading guide for practitioners

- **If ECI is falling** with clean QA: prepare for *faster* cross-scale propagation, tighten liquidity buffers, rehearse contingency playbooks, and re-check correlated exposures.
- **If ECI is rising** steadily: explore throughput trade-offs, can some buffers be streamlined without eroding resilience?
- **If families diverge** (high $\hat{\tau}^2$): investigate measurement breaks or sectoral idiosyncrasies before acting.
- **If QA flags trigger**: treat ECI as informative *context*, not a decision trigger.

7.11 Ethical and equity considerations

Coherence can be *engineered* in ways that unintentionally disadvantage smaller entities (e.g., disclosure burdens). Any policy use of ECI should:

- Publish methodology and QA transparently.
- Include **fairness impact assessments** (are small firms or low-income regions systematically penalized?).
- Prefer **carrots** (standards, tooling) over **sticks** that entrench incumbency.
- Respect data privacy and licensing; release synthetic replicas when raw sharing is restricted.

7.12 Takeaways

1. **ECI is a structural lens**: it measures the *shape* of timing across scales, not levels or instantaneous noise.
2. **Resilience ↔ coherence**: higher α often aligns with smaller drawdowns and faster recoveries, but with a throughput trade-off.
3. **Early warning**: clean ECI drops frequently precede stress, valuable alongside, not instead of, classical indicators.
4. **Bounded applicability**: where collapse fails or regimes mix, do not force RTM, log the negative result and fall back to domain-specific tools.
5. **Actionability** comes from **QA-aware interpretation**, redundancy across families, and, eventually, causal designs that move from prediction to policy.

8. Policy & Design Implications

This section turns **ECI** from a diagnostic into **design guidance**. We outline (i) coherence-aware stress testing, (ii) disclosure standards that make α_{econ} measurable, (iii) market-structure and supply-chain design patterns, (iv) macroprudential uses, (v) operational playbooks for public and private institutions, and (vi) governance guardrails. The theme is simple: **engineer tempo across scales** so that shocks dissipate rather than amplify, *without* freezing productive flow.

8.1 Coherence-aware stress testing

Objective. Move beyond level shocks (GDP, capital ratios) to **cross-scale timing shocks**: “what happens if the time gradient flattens (ECI↓) or steepens (ECI↑)?”

Test block A — Slope shocks.

- **Scenario A1 (Decoherence):** impose $\Delta\alpha_{\text{econ}} = -0.3$ for 2–3 windows with QA-clean conditions; propagate through sectoral input–output timing: shorter inventory half-lives, faster rumor cascades, reduced order-book resiliency.
- **Scenario A2 (Over-layering):** impose $\Delta\alpha_{\text{econ}} = +0.3$; propagate longer settlement and replenishment times; evaluate throughput loss vs. drawdown mitigation.

Metrics. Peak-to-trough outcomes, recovery half-life, synchronization indices (phase-lock across layers), and spillover multipliers.

Pass criteria. (i) Critical services remain above pre-registered continuity thresholds; (ii) no phase-lock across >2 layers in A1; (iii) in A2, throughput loss $\leq$ policy tolerance.

8.2 Disclosure standards that make α visible

Problem. Many jurisdictions collect levels and balance-sheet data but not **characteristic times**.

Minimum viable disclosure (per sector).

- **Logistics:** lead-time distributions, dwell-time tails, re-order cadence (anonymized).
- **Credit:** maturity ladders, roll-over windows, renewal rates by tenor.
- **Markets:** order-book resiliency metrics, standardized microprice reversion half-lives.
- **Information:** latency of correction/errata, editorial cadence, API timestamps.

Standard. Publish **quantiles of characteristic times** and the **bin definitions** (environment metadata). This enables third parties to compute α without exposing micro identifiers.

8.3 Market-structure patterns that raise coherence (without killing throughput)

M1 — Layered circuit breakers (time-aware).

Staggered pauses linked to *cross-scale* conditions (e.g., microstructure resiliency failing across capitalization tiers), rather than single-threshold halts. **Effect:** increases α transiently to prevent phase-lock.

M2 — Depth replenishment auctions.

Micro-auctions triggered when order-book depth drops below tiered thresholds; restore staging without long halts.

M3 — Clock desynchronization.

Randomized micro-offsets in batch auctions or reporting, small but sufficient to prevent algorithmic herding.

M4 — Transparency on timing rather than raw volume.

Mandate publishing resiliency/half-life metrics alongside liquidity stats; markets compete on *recovery speed* quality, not only spread.

8.4 Supply-chain & operations patterns

S1 — Buffer targeting by α .

Tie safety stocks and reorder points to sectoral $\hat{\alpha}$: when ECI drops, automatically widen buffers for critical inputs; when ECI rises, allow staged normalization.

S2 — Multipath routing by decoherence flags.

When ECI_EVENT triggers, switch to route sets that reduce **path-length variance** (not necessarily shortest), stabilizing T .

S3 — Cadenced procurement.

Avoid synchronized mega-orders; enforce **phase offsets** across suppliers to maintain timing heterogeneity.

S4 — Containment drills.

Treat decoherence like a cyber incident: playbooks to slow cross-layer propagation (temporary quotas, staggered schedules, alternate depots).

8.5 Macroprudential uses

P1 — Countercyclical ECI buffer.

Analogous to CCyB: when ECI falls below a percentile threshold (QA-clean), raise countercyclical capital/liquidity buffers; relax as ECI normalizes.

P2 — Maturity harmonics.

Discourage excessive bunching of corporate or sovereign maturities (reduce phase-lock); offer incentives for **staggered ladders**.

P3 — Disclosure cadence governance.

Stabilize α by setting **predictable announcement windows** (macro prints, policy updates) to avoid cascaded surprises.

P4 — Interbank contingency by timing.

Stress test *roll-over times* rather than only levels; pre-arrange facilities keyed to rollover half-lives, not just spreads.

8.6 Public-sector operationalization (roadmap)

Phase 0 — Baseline. Build an **ECI lab**: compute retrospective ECI on public data; publish methodology, collapse pass rates, QA.

Phase 1 — Pilot.

- Select 2–3 sectors with good coverage; run **live ECI(t)** for 12 months.
- Integrate with existing stress dashboards; define **ECI event** response playbooks.

Phase 2 — Standardization.

- Issue **timing disclosure templates**; onboard regulated entities.
- Convene an **ECI working group** (stat offices, central bank, market operators, supply-chain agencies).

Phase 3 — Policy integration.

- Tie **countercyclical tools** (buffers, facilities) to ECI triggers;
- Publish **transparency reports** (how often ECI flags were used; outcomes).

8.7 Private-sector operationalization

Firms & funds.

- Add ECI to risk dashboards; **red team** decoherence scenarios.
- Embed **ECI-gated rules**: e.g., leverage caps, VaR scaling, inventory buffers.
- Procurement and treasury coordinate on **maturity/ordering phase offsets**.
- Investor relations publish **timing KPIs** (recovery half-life, replenishment half-life).

8.8 Governance, fairness, and misuse risks

Guardrails.

- **No single-number tyranny.** Publish **QA flags** and **uncertainty bands**; never mandate actions on ECI alone.
- **Equal access.** Timing disclosures must be **public** (or symmetrically licensed) to avoid insider advantages.
- **SME/EM bias.** Provide tooling/support so small firms and emerging markets can meet timing disclosures without undue burden.
- **Privacy.** Release **aggregated quantiles** and synthetic replicas; independent audits for re-identification risk.
- **Auditability.** Keep a **slope-intercept ledger** and vintage archives; enable third-party recomputation.

8.9 Implementation templates (ready-to-use)

Template A — ECI policy trigger (public).

- **Trigger:** $ECI(t) - ECI(t - 3m) \leq -\theta$, QA clean, two families confirm.
- **Actions:** raise CCyB by X bps; activate liquidity facilities tied to *roll-over half-lives*; instruct market operators to enable **M1/M2**.
- **Sunset:** auto-review at +6 months; revert if ECI normalizes and stress absent.

Template B — Corporate playbook (private).

- **Trigger:** *ECI_EVENT* in sector/region.
- **Actions:** widen safety stocks by y%; enforce **S3** cadencing; diversify maturities; tighten algo throttles; send supplier phase-offset advisories.
- **KPIs:** drawdown containment, half-life to service-level recovery, exposure to phase-lock (share of synchronized suppliers).

8.10 Limits and unintended consequences

- **Over-layering risk.** Blindly “raising α ” can bloat bureaucracy; apply **throughput caps** and *sunset clauses*.
- **Gaming.** Entities might cosmetically stagger reports; require **ex post** collapse validation and random audits.
- **Coordination failures.** If only part of a network shifts cadence, temporary frictions may rise; use **pilot corridors** before national rollouts.

8.11 Summary

Policy and design can **shape tempo across scales**. ECI provides a **measurable, falsifiable** handle on that structure, enabling **coherence-aware** stress tests, disclosures, and interventions. The guiding principle is **differentiated clocks**: enough staging to prevent cascades, not so much that we strangle throughput. With transparent QA, fairness safeguards, and open replication, ECI can sit alongside volatility, liquidity, and leverage as a **third axis** for resilient economies.

9. Limitations

This section makes explicit where **RTM-ECI** can mislead, fail, or be outperformed by classical approaches. We group limitations into **data, measurement, identification, model form, operationalization**, and **external validity**, and state what evidence would change our minds.

9.1 Data limitations

- **Coverage heterogeneity.** Some proxy families are rich for developed markets (microstructure) but sparse for logistics/credit in smaller economies. **Risk:** ECI reflects *where data exist*, not coherence. **Mitigation:** LOW_COVERAGE flags, minimum-coverage gates, publish participation maps, abstain from inference when flagged.
- **Latency and revisions.** Logistics/credit series may arrive late or be revised. **Risk:** spurious ECI swings and hindsight bias. **Mitigation:** vintage accounting, latency disclosure, real-time backtests on frozen vintages.
- **Breaks in reporting standards.** Regulatory or vendor changes can shift measured T without structural change. **Risk:** step shifts in intercept that leak into slope if bins mix regimes. **Mitigation:** changepoint filters; slope–intercept ledger; exclude REGIME_MIX windows.

9.2 Measurement limitations

- **Proxy fragility.** Some (L, T) pairs depend on modeling choices (e.g., how “half-life” is fit). **Risk:** estimator-induced slopes. **Mitigation:** recipe sheets, estimator swaps (ODR/Theil–Sen/SIMEX), robustness envelopes.
- **Layer incompatibility.** Misaligned L and T (micro vs. macro) create spurious relations. **Mitigation:** compatibility rule; collapse test designed to **fail** such bins.
- **Errors-in-variables mis-specification.** If we under/overstate noise in L , SIMEX/ODR corrections can bias $\hat{\alpha}$. **Mitigation:** replicate measures where possible; sensitivity bounds; report null (shuffle) distributions.

9.3 Identification and causality limits

- **Associational nature.** ECI is structural-descriptive; H1–H2 give *predictive* content, not causal proof. **Risk:** policy overreach from correlation. **Mitigation:** reserve causal claims for settings with instruments, natural experiments, or randomized cadence interventions.
- **Confounding by policy clocks.** Uniform changes in timing (e.g., mandated settlement delays) may alter intercepts and sometimes slopes if heterogeneously adopted. **Mitigation:** bin by regime; test for slope invariance pre/post; document in ledger.

- **Reverse timing choice.** Agents may *increase layering* in anticipation of shocks, making ECI look prescient. **Mitigation:** preregistration, out-of-sample evaluation, difference-in-differences where cadence policies vary exogenously.

9.4 Model-form limitations

- **Power-law misspecification.** Some domains may follow $\log T = g(\log L)$ with curvature. **Risk:** biased α , false collapses. **Mitigation:** spline alternatives; declare *non-power* bins; treat as negative results (scope boundary).
- **Single- α assumption within bins.** Heterogeneous sub-regimes can require mixture models. **Risk:** average α hides opposed structures. **Mitigation:** stratify; finite-mixture fits; raise heterogeneity thresholds ($\hat{\tau}^2$) for fusion.
- **Temporal nonstationarity inside windows.** Rapid structural shifts violate the “fixed environment” premise. **Mitigation:** shorten windows; increase changepoint sensitivity; drop windows with REGIME_MIX.

9.5 Operational and governance limitations

- **One-number overreach.** Treating ECI as a master control risks **bureaucratic over-layering** ($\alpha \uparrow$ with throughput loss). **Mitigation:** multi-metric dashboards; sunset clauses; throughput caps; never trigger policy on ECI alone.
- **Gaming and Goodhart’s law.** If timing metrics become targets, agents may cosmetically stagger reports. **Mitigation:** random audits; ex-post collapse validation; cross-family consistency checks.
- **Equity and access.** Timing disclosures can burden SMEs/EMs. **Mitigation:** publish templates, subsidize tooling, allow aggregated quantiles, monitor fairness impacts.

9.6 External validity and transferability

- **Cross-country comparability.** ECI is contextual to bin definitions and data conventions. **Risk:** pseudo-rankings across incomparable regimes. **Mitigation:** harmonize bins before comparison; report comparability scores.
- **Sector heterogeneity.** High- α sectors (utilities, pharma) and low- α sectors (fast retail) differ structurally; uniform policy recipes are inappropriate. **Mitigation:** sector-specific playbooks; avoid universal targets.
- **Shock typology.** Some shocks are **clock exogenous** (policy moratoria) or purely level shocks; ECI may add little. **Mitigation:** declare “low-yield” shock classes a priori; use classical tools instead.

9.7 What would change our minds

We would regard RTM–ECI as **not useful** for a domain if, across multiple datasets and regimes:

1. **No slope separation** is detectable under robust EIV fits;
2. **Collapse routinely fails** in well-specified bins;
3. **Reverse cascade** (α decreasing with aggregation) appears persistently in steady state;
4. **H2 adds no predictive value** out-of-sample beyond strong baselines;
5. Results **invert** under reasonable proxy/estimator swaps.

Publishing such outcomes is part of the program: they define the method's **scope boundary**.

9.8 Roadmap to reduce limitations

- **Data:** expand timing disclosures; standardize recipe sheets; build open synthetic replicas.
- **Measurement:** invest in repeat measurements to calibrate EIV; broaden proxy families.
- **Identification:** seek natural experiments (cadence mandates, circuit-breaker reforms); pilot randomized cadence offsets.
- **Model form:** add mixture/spline diagnostics; automate non-power flags.
- **Governance:** codify QA gates, uncertainty bands, fairness audits; maintain public vintage archives.

10. Ethics & Governance

This chapter sets guardrails for **using, publishing, and acting on ECI**. Because ECI can influence capital allocation, regulation, and public narratives, the governance goal is twofold: (i) **prevent misuse** (one-number tyranny, gaming, unequal access), and (ii) **institutionalize good practice** (transparency, fairness, reproducibility). We structure guidance across (A) transparency & accountability, (B) fairness & access, (C) privacy, (D) decision protocols, (E) audits & red-teaming, and (F) stewardship & open science.

10.1 Transparency & accountability

10.1.1 Public method cards.

Every published ECI series must ship with a “method card” that states: bin definitions, proxy families, estimator choices (ODR/Theil–Sen/SIMEX), collapse pass rates, heterogeneity $\hat{\tau}^2$, QA flags, and vintage policy. Provide a human-readable summary and a machine-readable spec (YAML/JSON).

10.1.2 Slope–intercept ledger.

Maintain a ledger of known **level/clock shifts** (units, policy rebasings) alongside α . This clarifies why intercepts moved and defends slope robustness.

10.1.3 Negative results.

Register and publish bins where **collapse fails** or **α is unstable**. Non-publication of failures biases incentives and invites Goodhart's law.

10.2 Fairness & access

10.2.1 Equal-timing disclosure.

Timing metrics (lead-time distributions, resiliency half-lives) should be **public or symmetrically licensed**, not paywalled to selected actors. If regulators rely on ECI, they must ensure equal access to inputs.

10.2.2 SME/EM burden.

Small firms and emerging markets can't carry heavy reporting loads. Provide **templates, open-source tooling, and grants** so timing disclosures don't entrench incumbency.

10.2.3 Impact assessments.

Before ECI-guided policy (e.g., buffers keyed to ECI), run a **fairness impact assessment**: who bears costs/benefits across size, sector, and region? Publish mitigations (phase-in schedules, exemptions).

10.3 Privacy & confidentiality

10.3.1 Aggregate by design.

Publish **quantiles** of timing variables and ECI at bin level; avoid micro-identifiers. When micro data are necessary for research, use **secure enclaves** and audited access.

10.3.2 Re-identification audits.

Run periodic **linkage tests** against external registries to assess re-identification risk; rotate or coarsen binning if risk rises.

10.3.3 Synthetic replicas.

Release **synthetic datasets** that preserve distributional properties and collapse behavior, enabling independent verification without exposing raw micro data.

10.4 Decision protocols (how to act on ECI)

10.4.1 No single-number triggers.

ECI should **not** be a lone decision rule. Combine with volatility/liquidity/leverage metrics and qualitative intelligence. Document when ECI informed but did not decide.

10.4.2 QA-gated usage.

Actions keyed to ECI require **QA clean** status: ≥ 2 proxy families, collapse passed, heterogeneity below threshold, no REGIME_MIX. If QA fails, ECI can inform *monitoring*, not *action*.

10.4.3 Sunset clauses & throughput caps.

Policies that “raise α ” (more staging) must include **sunsets** and **throughput caps** to avoid bureaucratic over-layering.

10.4.4 Escalation ladders.

Tie **graduated responses** to the *magnitude and persistence* of ECI moves (e.g., advisory → targeted buffers → system-wide measures), with automatic de-escalation when ECI normalizes.

10.5 Audits, red-teaming, and model risk

10.5.1 Independent recomputation.

At least annually, a third party recomputes binwise $\hat{\alpha}$, collapse stats, and ECI(t) from posted artifacts. Differences beyond tolerance trigger a public postmortem.

10.5.2 Red-team scenarios.

Commission adversarial reviews probing: proxy leakage (outcomes feeding inputs), estimator brittleness, “clock hacks”, and data vendor exclusivity. Publish findings and fixes.

10.5.3 Stressing assumptions.

Run **non-power** alternatives (spline $g(\log L)$), mixture models, and regime-mixing simulations. Where power-law fails persistently, flag ECI as **not applicable**.

10.5.4 Governance overweights.

If a single proxy family repeatedly dominates fusion weights, require a **diversification plan** (add complementary proxies or cap weights) to reduce model monoculture risk.

10.6 Communications ethics

10.6.1 Avoid deterministic language.

Use “associated with”, “predictive of”, not causal claims, unless supported by explicit designs (instruments, natural experiments, RCTs).

10.6.2 Contextualize uncertainty.

Always show **bands and flags**. Provide plain-language explanations of what failure modes mean (“we could not validate scaling this quarter”).

10.6.3 Historical responsibility.

When ECI informs policy that affects livelihoods, publish **after-action reports**: what signals we saw, choices made, and outcomes (including mistakes).

10.7 Institutional stewardship & open science

10.7.1 Public registries.

Host a **registry of ECI vintages**, bin specs, QA logs, and hypothesis pre-registrations. Time-stamp everything.

10.7.2 Working groups.

Create cross-institution **ECI working groups** (stat offices, central banks, exchanges, ports, academia) to harmonize bins and share negative results.

10.7.3 Education.

Publish primers for practitioners and civic readers explaining slopes vs. levels, collapse tests, and why **negative findings** are successes for science.

10.8 Ethical red lines

- **No surveillance creep.** Timing disclosures must not morph into individual-level behavioral monitoring.
- **No punitive use without due process.** ECI flags are not grounds for sanctions absent statutory frameworks and rights to contest.
- **No exclusionary licensing.** If public entities act on ECI, core inputs and methods must be accessible to those affected.

Summary. ECI becomes ethically usable when institutions **share methods and uncertainty, guard against inequity and gaming, avoid one-number rulemaking, and invite independent recomputation.** Governance should make it *easy* to do the right thing (transparent, QA-gated, audited) and *hard* to do the wrong thing (opaque, exclusive, overconfident).

Chapter 11: Empirical Validation of Phase Bifurcation in High-Frequency Markets

This chapter stress-tests the RTM Real-Time Monitor against the extreme variance of the Bitcoin microstructure. By abandoning static daily closes and injecting the full continuous noise profile of minute-by-minute OHLCV data, we track the exact moment of structural fracture. The

continuous analysis isolates the Phase Bifurcation threshold ($\alpha < 0.5$), distinguishing mechanical liquidity failures (e.g., March 2020) from high-viscosity political stress events ($\alpha > 0.6$, e.g., May 2021). During the October 2025 event the noise-corrected metric showed an elevated α reading roughly 15 hours before the price move. This is a **single, in-sample observation**, and the event was subsequently attributed to a Binance platform technical glitch rather than a fundamental structural crash. It is reported as a case study only and must not be read as validated prospective prediction; replication on held-out events would be required to make any anticipatory claim.

Chapter 12: Empirical Analysis: The Collapse of α as a Predictive Signal

This chapter departs from the traditional Gaussian assumptions of market recoveries and crash predictions. Initial naive OLS models predicting crash recoveries suffered from massive attenuation bias due to the ambiguous boundaries of “market recovery.” By applying Orthogonal Distance Regression (ODR) to absorb a 20% measurement noise margin, we reveal that recovery time scaling is substantially more punishing (slope = 3.59 ± 0.70) than previously modeled. (This result derives from the historical crash-magnitude/duration record, not from the withdrawn α table; it was independently reproduced at slope ≈ 3.9 . It rests on 16 events with ambiguous recovery-date definitions and should be read as suggestive.)

An earlier version reported that a Monte Carlo analysis of 13 major crashes showed structural α -decorrelation preceding price troughs by ~ 9.75 days ($d = -1.45$). The underlying 13-event table was subsequently found to contain synthetically generated α values (Chapter 12.6), so these figures are **withdrawn**. Independent analysis on real S&P 500 daily returns (Chapter 12.7) finds no significant crisis-vs-calm α separation ($d \approx -0.2$, $p \approx 0.22$) and no pre-crash α decline. The recovery-time scaling result (above) and the multi-scale coherence metric (Chapter 12.7) are computed from independent data and are retained.

12. Conclusion

12.1 What This Paper Set Out to Do — and What It Found

This paper proposed that economic time is not an absolute background variable but a dynamic dimension that scales with the multiscale structural organization of the system. Through the derivation of the Economic Coherence Index (ECI) and its underlying cascade framework (Chapters 1–10), we built a complete theoretical and operational apparatus for this idea: proxy families, collapse tests, meta-analysis, alert logic, governance protocols. The theoretical architecture stands as a proposed framework.

The empirical results are more modest than early drafts of this work claimed, and intellectual honesty requires stating the revised record plainly.

A 13-event cross-market analysis originally reported a large in-sample separation between healthy and crash-state DFA exponents ($d = -1.45$, ~ 10 -day lead, severity correlation $r = 0.97$). That analysis rested on synthetically generated α values (Chapter 12.6) and is **withdrawn**. Independent validation on real S&P 500 daily returns (5,104 days spanning the Dot-Com, GFC, and COVID crises) finds the genuine crisis-vs-calm DFA- α separation to be small and not significant ($d \approx -0.2$, $p = 0.22$). The hypothesis that α declines before crises (H2) is contradicted: α rose before both testable onsets (GFC and COVID), and shows no pre-crash depression at monthly resolution across 155 years of data.

What survives is narrower but real. The **multi-scale coherence metric**, the cross-scale variance of α is lower during crisis windows than during calm windows on both BTC intraday data ($\sigma = 0.03$ vs 0.31) and daily S&P 500 ($p = 0.039$, $AUC = 0.60$). The direction replicates; the magnitude is asset- and timescale-dependent ($\sim 10\times$ on crypto, $\sim 1.3\times$ on equities). The **inverse-cubic tail** ($\alpha \approx 3$) is independently reproduced from raw daily returns (far-tail Hill estimator $\alpha = 2.9\text{--}3.3$), confirming a known result (Gabaix et al. 2003) rather than contributing a new one. The crisis-differentiation observation, that liquidity vacuums (COVID 2020, $\alpha < 0.5$) are structurally distinct from political shocks (China Ban 2021, $\alpha > 0.6$) remains a legitimate descriptive characterization from raw BTC microstructure data.

The Red Team's out-of-sample accuracy of 25% (1 of 4 post-2022 events) is the correct bottom line for prediction.

12.2 Why Economics Is Different: The Boundary Condition

Across the broader RTM corpus, the framework demonstrates genuine empirical traction in physics (quantum processor decoherence), geoscience (Gutenberg-Richter scaling, Omori aftershock decay, tornado discrimination), biology (mass-longevity allometry, flight-longevity bonus), and fluid dynamics (Richardson dispersion). In each of these domains, the scaling exponent α behaves as Doc 002 requires: a structural invariant, stable within measurement windows, set by the system's transport-network topology rather than by the observer.

Economics is the domain where this fails, and the failure is not accidental. It reveals a **boundary condition** of the RTM framework that, stated precisely, strengthens rather than weakens the theory by delimiting where it applies and where it does not.

Four properties jointly distinguish financial markets from every other system in the corpus:

Reflexivity. An earthquake does not read its own Gutenberg-Richter b -value and adjust its magnitude in response. A bat does not observe its longevity quotient and modify its metabolic rate. But a market participant who discovers that an α -drop predicts crashes will *act on that signal*, and by acting, will either eliminate or transform it. Markets are the only system in the corpus where the agents inside the system observe, interpret, and reconfigure the system in real time. Any scaling regularity that becomes known and actionable tends to be arbitrated toward extinction. In physics and biology, no equivalent mechanism exists: the laws of scale are not subject to strategic manipulation by the entities they govern.

Endogenous topology. RTM works when there is a transport network whose structure is stable over the measurement timescale. A mammal's vascular network, a superconducting chip's qubit lattice, a plate-boundary fault system, these topologies change on timescales (evolution, fabrication, geology) far longer than the phenomena RTM measures. In financial markets, the network, who trades with whom, which algorithms are active, which cross-asset correlations hold, reconfigures endogenously in milliseconds. The topology is itself a fast variable, not a slow parameter. When the network rewires faster than the exponent can equilibrate, α cannot function as a structural invariant. It becomes, at best, a snapshot of a transient state.

Non-conservation of the transported quantity. The semigroup derivation (Doc 001, Appendix J) models the transport of a conserved quantity through a fixed structure. In physics, what flows is conserved: energy, mass, charge, momentum. In markets, the relevant quantities, credit, information, sentiment, liquidity, are not conserved in the physical sense. Credit is created ex

nihilo by lending. Information is amplified, distorted, and destroyed. Panic self-reinforces without external input. When the substance flowing through the network can be endogenously created and destroyed, the transport laws that assume conservation hold less rigidly. They may still apply locally and transiently, during a crash, when the market briefly behaves like a physical cascade, but not as stable structural invariants.

Structural non-stationarity. In natural systems, scaling laws are approximately stationary over observational windows. The Gutenberg-Richter b -value holds across decades. Kleiber's law holds across taxonomic orders. The tornado α replicates across 7 of 9 independent outbreaks within a single year. But a market's scaling behavior changes with regulation, monetary policy, technology, and the collective learning of its participants. The pattern found in-sample in 2018–2022 may not exist in 2023–2025, not because the measurement was wrong, but because the system has genuinely changed. This is why the 25% out-of-sample accuracy is not merely “the effect is weaker”; it reflects the possibility that the specific scaling relationship discovered in-sample *ceased to exist* in the test period.

Taken together, these four properties explain why the RTM exponent functions as a structural invariant in physics, biology, and geoscience, but degrades into a transient descriptor in economics. The framework's mathematical machinery is not wrong when applied to markets; it is applied to a system that violates its foundational assumptions, stable topology, conserved transport, non-reflexive agents, and approximate stationarity.

12.3 The Crash as a Window of Physical Behavior

There is a revealing corollary. The one economic regime where RTM-like scaling *does* partially replicate — the multi-scale coherence signature during crashes, is precisely the regime where the four boundary violations are temporarily relaxed:

- During a panic, **reflexivity partially shuts down**: strategic calculation gives way to mechanical liquidation. Agents stop optimizing and start dumping. The market briefly acts more like a physical cascade than a strategic game.
- During a liquidity crisis, **the topology rigidizes**: counterparties withdraw, market-making ceases, and the network of active connections freezes into a simpler, more rigid structure. The fast-variable topology temporarily becomes a slow parameter.
- During a crash, **flows become more nearly conserved**: the dominant dynamic is the redistribution of a fixed (shrinking) pool of liquidity, not the creation of new credit or information.
- During a crash, the **structural non-stationarity narrows**: all participants face the same constraint (survive), and the diversity of strategies collapses into a near-uniform “sell everything.” The system becomes briefly stationary in its desperation.

This is why the multi-scale coherence metric, which measures whether α becomes scale-invariant — shows its strongest signal during crashes: the market is temporarily behaving like a physical system, and RTM's assumptions are temporarily satisfied. As soon as recovery begins and strategic diversity returns, the signal weakens. The 1.3× equity separation (vs 10× on BTC intraday) likely reflects the degree to which equities retain institutional structure even in panic, while crypto markets, less regulated, more retail-driven, thinner, approach the “pure physical cascade” limit more closely.

12.4 Diagnosis, Not Prediction

These considerations reframe what the Economic Coherence Index can and cannot aspire to:

What it cannot do. Reliably predict crashes out of sample with operational lead time. The 25% accuracy is not a starting point to be improved with more data; it is a consequence of the structural properties described above. A better RTM crash predictor would require a market whose topology is stable, whose participants do not learn, and whose transported quantities are conserved, a market that, by definition, no longer functions as a market. The document's earlier claims of "mathematical authority to predict structural failures" are withdrawn.

What it can do, and this is not nothing. Characterize the *type* of stress a market is under in near-real-time. The distinction between a liquidity vacuum ($\alpha < 0.5$, anti-persistent, memoryless) and a viscous political shock ($\alpha > 0.6$, persistent, structurally intact) is a genuine structural diagnostic that price and volatility alone do not capture. Whether the topology is *currently* coherent across scales (σ near zero) or fragmented (σ elevated) is a question RTM can answer, descriptively, not predictively. This positions the ECI as a **regime classifier**, not an early-warning system: it tells you *what kind of regime the market is in now*, not *what it will do next*.

For policymakers and risk managers, this is still useful. Knowing that a market is in a liquidity-vacuum state (requiring emergency provision of market-making) versus a viscous-stress state (requiring patient resolution of solvency uncertainty) is a diagnostic that informs *response*, even if it cannot inform *anticipation*. The analogy is medical: a thermometer tells you the patient has a fever and what kind (high, low, cycling), but it does not predict when the fever will break. That is still a useful instrument.

12.5 The Honest Position

RTM-Economics occupies a unique and instructive position in the broader corpus. It is the domain where the framework's predictions are weakest, and the domain that, by its very resistance, most precisely delineates the framework's boundary conditions. The four properties that make markets resistant to scaling analysis (reflexivity, endogenous topology, non-conservation, non-stationarity) are exactly the properties that *define* what a market is, as distinct from a physical or biological system. RTM does not fail in economics because its mathematics is wrong; it fails because markets are the one system in the corpus whose participants can observe the exponent and change it.

This boundary condition, **RTM functions where the topology is a slow parameter and transport obeys conservation; it degrades where the topology is endogenous and the agents are reflexive**, is, we believe, a genuine contribution to understanding the limits of scaling analysis in complex systems. It explains why the same mathematical framework that cleanly separates tornadic from non-tornadic storms ($d = 0.96$), quantifies the coherence cost of monolithic quantum scaling ($\alpha = -0.26$), and recovers the flight-longevity bonus across 2,400 endotherms ($d = 0.85$) cannot produce a reliable crash predictor in financial markets, and why the expectation that it should was always the wrong question to ask.

The right question was never "can RTM predict crashes?" It was: "what does it mean that RTM's scaling assumptions are violated in markets but not in physics?" The answer, that markets are reflexive, topologically unstable, and non-conservative in ways that no natural system is, is itself

a finding, and it belongs in the conclusion rather than being hidden behind an inflated claim of predictive power.

12.6 Data Provenance Correction (July 2026) A forensic audit established that `crash_alpha_analysis.csv`, the source of all 13-event α statistics in this chapter, was synthetically generated rather than measured from market data. Replaying `numpy`'s random generator seeded at 42 and consuming five uniform draws per event reproduces the file's columns to machine precision:

Column	Generator	Exact matches
Baseline_Alpha	$0.5 + 0.1 \cdot U(0,1)$	13/13
Pre_Alpha	Baseline + $(-0.03 + 0.06 \cdot U)$	13/13
Lead_Time_Hours	$\text{int}(72 + 432 \cdot U) \rightarrow U(3,21)$ days	13/13
Immediate_Alpha	function of Baseline and Drop_Pct	severity encoded by construction

Consequences: (i) the $d = -1.45$ separation is a property of the simulation design; (ii) the ~ 9.75 -day lead time is the mean of uniform 3–21-day draws on the significant subset; (iii) the $r = 0.97$ severity correlation was built into the generator (`Immediate_Alpha` depends on `Drop_Pct`). All three are withdrawn as empirical findings. The downstream ODR/Monte-Carlo pipeline (`analyze_financial_robust.py`) is not itself at fault, it faithfully processed a synthetic input that had entered the pipeline upstream and was thereafter treated as real. The April 2026 Red Team reproduced the *statistics* of the table but did not test *provenance*; provenance verification is now part of the audit standard. Reproduce with `forensic_audit_crash_alpha.py`.

12.7 Real-Data Replication (What Survives) To replace the withdrawn table, the core claims were re-tested on public, verifiable data: 5,104 days of S&P 500 OHLCV (2000–2020, covering the Dot-Com, GFC, and COVID crises), daily VIX (1990–2026), and Shiller monthly S&P (1871–2026). The DFA implementation was calibrated on known processes (white noise $\alpha = 0.51$; Brownian motion $\alpha = 1.50$).

Test	Prior (synthetic) claim	Real-data result	Status
Crisis vs calm DFA- α	$d = -1.45$	$d = -0.22$ (drawdown, $p = 0.22$); $d = +0.03$ (VIX, $p = 0.93$)	Not supported
H2: α drops before crises	6–18 mo lead	α rose before GFC (+0.05) and COVID (+0.07); no monthly-scale signal 1871–2026	Contradicted
Multi-scale coherence σ	10× (BTC, 3 months)	Direction replicates on equities: crisis $\sigma < \text{calm } \sigma$, $p = 0.039$, AUC 0.60, ratio 1.31×	Partially supported
Inverse cubic tail $\alpha \approx 3$	Literature citation	Hill estimator on raw returns: far-tail $\alpha = 2.9\text{--}3.3$	Independently reproduced

The multi-scale coherence metric is the only RTM-native signal that survives out-of-sample transfer to an independent asset class; it is real but weak on equities, and its magnitude is timescale-dependent. Hypothesis H2 (α -decline anticipation) is now contradicted by three

independent lines of evidence (real-data event study, century-scale monthly analysis, and the Red Team’s 25% out-of-sample accuracy) and should be retired or reformulated. Full pipeline: `rtm_econ_new_validations.py`.

Appendices

These appendices give implementers everything needed to reproduce, audit, and extend the **Economic Coherence Index (ECI)**: the math behind the scaling law and estimators, a full specification (data schemas, QA gates, defaults), evaluation metrics, and robustness/ablation checklists.

Appendix A — Mathematical Notes

A.1 From scale symmetry to a power law

Assume a characteristic time $T(L)$ depends on a size/scale proxy $L > 0$ and satisfies **scale symmetry**:

- For any $b > 0$, rescaling $L \mapsto bL$ rescales time by a factor $f(b)$: $T(bL) = f(b) T(L)$.
- Composition of rescalings is multiplicative: $f(b_1 b_2) = f(b_1) f(b_2)$.

Then f solves Cauchy’s exponential equation on $\mathbb{R}_{>0}$, yielding $f(b) = b^\alpha$ for some real α . Fixing any $L_0 > 0$ gives

$$T(L) = T(L_0) \left(\frac{L}{L_0} \right)^\alpha = \kappa L^\alpha, \quad \kappa = T(L_0) L_0^{-\alpha}.$$

Taking logs: $\log T = \alpha \log L + \log \kappa$.

Implication. Any uniform “clock” or level change multiplies κ (intercept), not α (slope).

A.2 Errors-in-variables (EIV) slope estimators

Let $x = \log L$, $y = \log T$, with noisy observations $x^{obs} = x + \xi$, $y^{obs} = y + \zeta$, $\mathbb{E}[\xi] = \mathbb{E}[\zeta] = 0$.

Orthogonal Distance Regression (ODR / TLS).

Estimate $(\hat{\alpha}, \hat{c})$ by minimizing the sum of **orthogonal squared residuals** to the line $y = \alpha x + c$:

$$\min_{\alpha, c} \sum_u \frac{(y_u^{obs} - \alpha x_u^{obs} - c)^2}{1 + \alpha^2}.$$

Closed forms exist (via SVD of centered design); most libraries implement iteratively.

SIMEX (simulation–extrapolation).

If σ_ξ^2 (or a bound) is known/estimable:

1. Simulate added noise: $x^{(\lambda)} = x^{obs} + \sqrt{\lambda} \xi$, $\lambda \in \Lambda \subset \mathbb{R}_{\geq 0}$.
2. Fit slopes $\hat{\alpha}(\lambda)$ for each λ .
3. Extrapolate $\lambda \rightarrow -1$ (zero measurement error) with a low-order polynomial to obtain $\hat{\alpha}_{\text{SIMEX}}$.

Theil-Sen (robust).

Median of pairwise slopes $\{(y_j - y_i)/(x_j - x_i)\}$ over all $i < j$. Resistant to outliers; use as a sensitivity check.

A.3 Collapse test statistic

Given $\hat{\alpha}$ for a bin, define residualized outcomes $\tilde{y}_u = y_u^{obs} - \hat{\alpha}x_u^{obs}$. In a valid power-law bin, $\tilde{y}$ should be **independent** of x (aside from noise). Use:

$$\Delta_{\text{collapse}} := R^2(\tilde{y} \sim x^{obs}).$$

Pass rule (default): $\Delta_{\text{collapse}} < 0.05$ and a nonparametric smooth (e.g., LOESS with pre-registered span) shows no visible trend.

A.4 Random-effects fusion (DerSimonian-Laird / REML)

With family-specific estimates $\hat{\alpha}_f$ and variances $\hat{\sigma}_f^2$,

$$\hat{\tau}^2 = \max\left\{\frac{Q - (F - 1)}{\sum w_f - \sum w_f^2 / \sum w_f}, 0\right\}, w_f = \frac{1}{\hat{\sigma}_f^2}, Q = \sum w_f (\hat{\alpha}_f - \bar{\alpha}_w)^2,$$

$\bar{\alpha}_w = \sum w_f \hat{\alpha}_f / \sum w_f$. The fused estimate:

$$\hat{\alpha}_{\text{econ}} = \frac{\sum_f w_f \hat{\alpha}_f}{\sum_f w_f}, \quad w_f = \frac{1}{\hat{\sigma}_f^2 + \hat{\tau}^2}.$$

REML may replace DL for $\hat{\tau}^2$ when F is small.

A.5 Directionality tests (TE/Granger) for H3

Granger causality. X “Granger-causes” Y if lagged X improves prediction of Y beyond lagged Y . Use VAR with pre-registered lag p ; Wald tests with multiplicity control.

Transfer entropy (TE). Information-theoretic asymmetry:

$$TE_{X \rightarrow Y} = \sum p(y_{t+1}, y_t^{(p)}, x_t^{(p)}) \log \frac{p(y_{t+1} | y_t^{(p)}, x_t^{(p)})}{p(y_{t+1} | y_t^{(p)})}.$$

Estimate via kNN or model-based surrogates; significance with block-shuffled surrogates.

Appendix B — ECI Specification (Data, QA, Defaults)**B.1 Data schema (entity-level features → bin table)**

Entity table (per family):

- entity_id
- timestamp (UTC)
- L_value (scale proxy, positive)
- T_value (characteristic time, positive)

- L_unit, T_unit (strings)
- L_method, T_method (recipe IDs)
- fit_r2, fit_se_T (optional)
- env_keys (country, sector, policy regime, window_id)
- quality_flags (bitfield)

Bin table (per family × environment):

- env_keys
- n_entities, coverage_share
- alpha_hat, alpha_ci_low, alpha_ci_high
- intercept_hat
- collapse_R2, collapse_pass (bool)
- clock_placebo_pass (bool)
- qa_flags (enum: LOW_COVERAGE, NO_COLLAPSE, REGIME_MIX, CLOCK_SHIFT)

ECI table (fused per environment/time):

- window_id, env_keys
- alpha_econ_hat, band50_low/high, band95_low/high
- tau2 (heterogeneity), Q_stat, families_used
- qa_flags (as above)
- vintage_asof

B.2 Environment definition and regime control

- **Windows:** quarterly (primary), monthly stride.
- **Regime checks:** univariate/multivariate changepoint tests on representative level series and timing proxies; split window on detected breakpoints; flag REGIME_MIX if unsplittable.

B.3 Estimation defaults

- **Estimator:** ODR/TLS (default), Theil–Sen (robust check), SIMEX (when σ_{ξ}^2 known).
- **Bootstrap:** clustered by entity ($\geq 1,000$ replicates); report percentile CIs.
- **Collapse threshold:** $\Delta_{\text{collapse}} < 0.05$ and smooth trend visually flat.
- **Fusion:** random-effects (REML), require ≥ 2 families with collapse_pass==True.
- **Heterogeneity cap:** refuse fusion if $\hat{\tau}^2$ exceeds pre-registered percentile (e.g., 90th of historical τ^2); flag FAMILY_DIVERGENCE.

B.4 QA & acceptance gates

A bin contributes to ECI only if:

1. coverage_share $\geq$ sector-specific minimum,
2. collapse_pass == True,
3. clock_placebo_pass == True,
4. no REGIME_MIX.

The fused ECI(t) at a window is **published** only if ≥ 2 bins (families) meet these gates; otherwise, populate qa_flags and withhold decision use.

B.5 Decoherence event definition (defaults)

- Horizons $h \in \{1, 3, 6\}$ months.
- Threshold θ_h : historical 10th percentile of $\Delta\alpha_{\text{econ}}$ at horizon h **or** $k \cdot SE_t$ with k pre-registered (e.g., 1.64).
- Confirmation: at least two families show sign-consistent drops; QA clean at t and over $[t - h, t]$.

B.6 Publishing & vintage policy

- Publish **vintage-stamped** ECI files (asof), with code to reconstruct any historical curve.
- Keep a **slope-intercept ledger** recording unit changes, policy rebasings, vendor switches.

Appendix C — Forecast & Inference Metrics

C.1 Classification (H2)

- **AUC / PR-AUC** with DeLong intervals and bootstrap for PR-AUC.
- **Brier score, log loss**; Diebold–Mariano tests for forecast-score differences.
- **Calibration**: reliability diagrams; Expected Calibration Error (ECE).
- **Confusion**: hit/false alarm rates at policy-relevant thresholds (pre-registered).

C.2 Survival / duration (H1b, H2)

- **AFT model** coefficients with robust SEs; concordance index.
- **Cox model** (sensitivity): hazard ratios for ECI tertiles; Schoenfeld tests for PH assumption.

C.3 Cross-sectional regression (H1a)

- Fixed effects; cluster-robust SEs.
- **Shapley / dominance** analysis for incremental explanatory power vs. volatility/liquidity/leverage.

C.4 Multilayer cascade (H3)

- **Paired-difference tests** for $\alpha_{\text{micro}} \leq \alpha_{\text{meso}} \leq \alpha_{\text{macro}}$ across windows; FDR control.
- **TE/Granger** asymmetry statistics with surrogate-based p -values.

C.5 Presentation templates

- **Figure A (mandatory):** ECI(t) + 50/95% bands + QA flags.
- **Table A:** collapse pass rates by family and regime.
- **Table B:** fusion weights, $\hat{\tau}^2$, influence (leave-one-family-out).
- **Figure B:** event-aligned plots (ECI drops vs. stress onsets).

Appendix D — Robustness & Ablations

D.1 Estimator & proxy ablations

- **Estimator swap:** ODR $\leftrightarrow$ Theil–Sen $\leftrightarrow$ SIMEX; report $\Delta\hat{\alpha}$, CI overlap, decision stability for H1–H3.
- **Proxy swap:** within each family, alternate L (degree $\leftrightarrow$ path length, cap $\leftrightarrow$ trade-size tier) and T (half-life variant, alternative fit).
- **Outcome swap:** alternative stress labels (e.g., local crisis chronologies).

D.2 Windowing & coverage

- Windows: quarterly vs. semi-annual; stride: monthly vs. biweekly (where feasible).
- Downsample entities; track CI widening and collapse failure rates.

D.3 Placebo & nulls

- **Clock placebos:** rescale time units; confirm slope invariance, intercept shift.
- **Shuffle null:** permute L within bins; store null slope distribution and ensure observed $\hat{\alpha}$ exceeds null by pre-registered margins.
- **Non-power alternative:** spline $g(\log L)$ fits; if curvature persists across bins, mark domain **non-power** and exclude.

D.4 Heterogeneity management

- Cap fusion when $\hat{\tau}^2$ is excessive; publish family-wise $\hat{\alpha}_f$ instead of a single ECI.
- Require a diversification plan if one family repeatedly dominates weights.

Appendix E — Reproducibility & Packaging

- **Repository layout.**
data_raw/, data_processed/, features/, bins/, alpha_estimates/, collapse/, fusion/, eci/, qa/, figures/.
- **Pipelines.** Directed acyclic graph with deterministic seeds; CI tests: (i) unit-invariance, (ii) exact reproduction to tolerance, (iii) collapse bounds.
- **Docs.** YAML specs for proxy recipes; CHANGELOG for estimator versions; reproducible environment files.

Appendix F — Glossary (minimal)

- L : scale proxy (size/path/degree/cap tier).
- T : characteristic time (half-life, decay, replenishment, resiliency).

- α : slope in $\log T = \alpha \log L + c$ within a fixed environment; coherence exponent.
- **ECI(t)**: fused, QA-gated time series of $\hat{\alpha}_{\text{econ}}$.
- **Collapse**: residual independence of $\tilde{y} = \log T - \hat{\alpha} \log L$ from $\log L$.
- **Decoherence event**: significant QA-clean ECI drop over a pre-registered horizon.
- **QA flags**: LOW_COVERAGE, NO_COLLAPSE, FAMILY_DIVERGENCE, REGIME_MIX, CLOCK_SHIFT.

APPENDIX G — Computational Validation of RTM-Econ Framework

G.1 Overview

This appendix presents computational validation of the Rhythmic Economics (RTM-Econ) framework. Three simulation suites demonstrate:

1. α can be reliably estimated from cross-sectional financial data (S1)
2. α decline provides early warning of recessions (S2)
3. α varies systematically across economies and predicts resilience (S3)

G.2 S1: α Estimation from Financial Data

G.2.1 Model

RTM-Econ Scaling:

$$\tau(L) = \tau_0 \times (L/L_{\text{ref}})^{\alpha}$$

where:

- τ = characteristic time (recovery, persistence)
- L = scale proxy (market cap, firm size)
- α = coherence exponent

G.2.2 Market Regime Parameters

Regime	Period	α	Interpretation
Stable Growth	2004-2006	0.45	Good coherence
Pre-Crisis	2007	0.35	Coherence declining
Crisis	2008-2009	0.20	Decoherence, cascade risk
Recovery	2010-2012	0.40	Rebuilding coherence
New Normal	2013-2019	0.42	Post-crisis stable

G.2.3 Estimation Results

Regime	True α	Estimated α	Error
Stable Growth	0.45	0.447	0.003
Pre-Crisis	0.35	0.346	0.004
Crisis	0.20	0.192	0.008

Regime	True α	Estimated α	Error
Recovery	0.40	0.396	0.004
New Normal	0.42	0.416	0.004

Mean absolute error: 0.0056 (1.3%)

G.2.4 Multi-Family Meta-Analysis

Combining four proxy families for Stable Growth regime:

Family	Estimated α	95% CI
Recovery Half-Life	0.44	[0.38, 0.50]
Volatility Persistence	0.46	[0.39, 0.53]
Autocorrelation Decay	0.43	[0.35, 0.51]
Order Flow Relaxation	0.47	[0.38, 0.56]

Combined ECI: 0.447 (True: 0.45)

Heterogeneity I^2 : 0.12 (low, families agree)

G.3 S2: Early Warning Backtesting

G.3.1 Hypothesis H2

Claim: Sharp drops in α precede recessions by 6-18 months.

G.3.2 Recession Analysis

Recession	$\alpha_{\text{pre}} \rightarrow \alpha_{\text{trough}}$	$\Delta\alpha$	Lead Time
2001 Dot-Com	0.42 $\rightarrow$ 0.28	0.14	9 months
2008 GFC	0.45 $\rightarrow$ 0.18	0.27	15 months
2020 COVID	0.40 $\rightarrow$ 0.22	0.18	3 months

Mean lead time: 9 months

Mean α drawdown: 0.20

G.3.3 Comparison to Other Indicators

Indicator	Type	Typical Lead Time
ECI (α)	Leading (structural)	6-15 months
Yield Curve	Leading (financial)	8-12 months
VIX	Concurrent	0-1 months
GDP Growth	Lagging	Negative

G.3.4 Detection Protocol

1. Monitor rolling ECI with 3-6 month window
2. Establish baseline α during expansion

3. Alert when α drops >15% below baseline
4. Confirm with other leading indicators
5. Expected lead time: 6-18 months

G.4 S3: Cross-Country Comparison

G.4.1 Country Classification

Type	Countries	Mean α	Resilience
Developed	Germany, Japan, Switzerland	0.52	Very High
Financial Hub	US, UK, Singapore	0.42	Moderate
Transition	China, India, S. Korea	0.39	Variable
Emerging	Brazil, Turkey, Argentina	0.28	Low

G.4.2 Correlation Results

Relationship	Correlation	p-value
α vs Crisis Frequency	$r = -0.91$	< 0.001
α vs Average Drawdown	$r = -0.95$	< 0.001
α vs GDP per capita	$r = +0.68$	< 0.05

G.4.3 Top Resilient Economies

Rank	Country	α	Resilience Score
1	Switzerland	0.55	0.87
2	Japan	0.52	0.80
3	Germany	0.48	0.74

G.5 Summary of Computational Validation

Test	Metric	Result
α estimation	Mean error	0.56%
Meta-analysis	Heterogeneity I^2	0.12
Early warning	Mean lead time	9 months
Cross-country	α -crisis correlation	$r = -0.91$

G.6 Falsifiable Predictions

RTM-Econ fails if:

1. **No scaling:** τ vs L shows no power-law within market regimes
2. **No anticipation:** α does not decline before recessions
3. **No cross-country pattern:** High- α economies have equal crisis rates
4. **High heterogeneity:** Proxy families disagree ($I^2 > 0.75$)

G.7 Policy Implications

1. **Stress Testing:** Include α monitoring in macroprudential surveillance
2. **Early Warning:** α decline signals building fragility
3. **Policy Design:** Interventions that increase α (buffers, staging) enhance resilience
4. **Cross-Country:** Low- α economies need stronger institutional buffers

G.8 Methodological Note on Recovery Scaling and Non-Gaussian Distributions

Standard Ordinary Least Squares (OLS) regressions severely underestimate the difficulty of market recovery due to an attenuation bias. Defining the exact day a market “recovers” involves immense noise (e.g., inflation boundaries, dividend reinvestment logic). By deploying Orthogonal Distance Regression (ODR) to absorb these massive boundary errors (10% variance in crash depth, 20% in recovery time), the recovery scaling slope steepens dramatically from a flawed 2.49 to a robust 3.59 ± 0.70 . This suggests that economic recovery is substantially more punishing and non-linear than classical models assume.

Furthermore, to rigorously map the shape of global financial distributions, we utilized a Monte Carlo simulation ($n = 16,000$) across 16 global markets. The universal mean tail exponent converges at $\alpha = 2.966 \pm 0.236$. This is consistent with the theoretical RTM “Inverse Cubic Law” limit ($\alpha \approx 3.0$) and departs from Gaussian assumptions. It is a convergent result, in line with the established inverse cubic law (Gabaix et al. 2003), in which extreme price moves are recurrent statistical features of markets rather than isolated anomalies.